

tim & koko

OpenTelemetry & Grafana Tempo

Distributed Tracing Made Easy

Thomas Philipona

Bern, 21.08.2025



Agenda

1. Why
2. OpenTelemetry Introduction
3. Grafana Tempo
4. Demo
5. Distributed Tracing Stack in Practice

Observability – See, Understand, Improve



Know what's happening

It allows you to understand a system from the outside.
It is the foundation to detect issues quickly and understand what is happening.



Understand why

Diagnose root causes with context, which is super important in increasingly complex environments.



Improve faster


The what and why helps to reduce the time to fix issues (Shorter Mean Time to Repair).
Creates business value and happy customers!

Why Distributed Tracing matters?

Distributed tracing is a method for tracking requests in **distributed systems** (e.g., microservices architectures) from **start** to **finish** in order to understand **how** and **where** they spend time and where **problems** occur.

Three components complete the picture:

- Logs: What happened?
- Metrics: How is it going?
- Traces: Where is it stuck?



Distributed Tracing

OpenTelemetry



OpenTelemetry

What is OTel?

An observability framework and toolkit designed to facilitate the

- Generation
- Export
- Collection

of telemetry data such as **traces**, **metrics**, and **logs**.

<https://opentelemetry.io>

A screenshot of the OpenTelemetry website homepage. The header features the OpenTelemetry logo (a yellow and blue icon) and the text "OpenTelemetry" and "Cloud Native Computing Foundation (C...)" in black. Below the header are buttons for "CNCF" (blue), "INCUBATING" (dark grey), and several icons. The main content area contains the text "OpenTelemetry community content" and "GitHub stars: 926".

OpenTelemetry

Cloud Native Computing Foundation (C...)

CNCF INCUBATING


OpenTelemetry community content

GitHub stars: 926




OpenTelemetry

Components / Architecture



OpenTelemetry

Components / Architecture for Traces



OpenTelemetry Instrumentation



Zero Code

Great for getting started! Packaged SDK, libraries and exporters are injected automatically into the application. You will get your code auto instrumented. Such as HTTP APIs, Backend Calls, DB Calls, Messagequeues and so on.

.NET, Go, Java, JavaScript, PHP, Python



Code-based

Import OTel API and SDKs which provide access to the OpenTelemetry functionality. Ideally for custom instrumentation, to implement custom traces and metrics.

.NET, Go, Java, JavaScript, PHP, Python, C++, C, Rust, Ruby, Swift, Erlang

Zero Code Instrumentation

Java Spring Boot

- Add dependency to your application
- configure an endpoint and export signals
- optionally configure the instrumentation
- Similarly done for other languages (.NET, Go, JavaScript, PHP, Python)

gradle.properties

```
import org.springframework.boot.gradle.plugin.SpringBootPlugin

plugins {
    id("java")
    id("org.springframework.boot") version "3.2.0"
}

dependencies {
    implementation(platform(SpringBootPlugin.BOM_COORDINATES))
    implementation(platform("io.opentelemetry.instrumentation:open
-instrumentation-bom:2.19.0"))
}
```

Auto Instrumentation Kubernetes Operator

```
apiVersion: opentelemetry.io/v1alpha1
kind: Instrumentation
metadata:
  name: demo-instrumentation
spec:
  exporter:
    endpoint: http://demo-collector:4318
  propagators:
    - tracecontext
    - baggage
  sampler:
    type: parentbased_traceidratio
    argument: '1'
  java:
    env:
      - name: OTEL_INSTRUMENTATION_KAFKA_ENABLED
        value: false
      - name: OTEL_INSTRUMENTATION_REDISCALA_ENABLED
        value: false
```

Annotation on Deployment

```
apiVersion: apps/v1
kind: Deployment
[...]
spec:
  template:
    metadata:
      annotations:
        instrumentation.opentelemetry.io/inject-java: "true"
      spec:
        containers:
          [...]
```

OpenTelemetry

Telemetry Data /Signals



Traces

Path of a request
through the entire
system



Metrics

A measurement
captured at runtime.
eg.

http.requests.total 42
db.queries.count 501



Logs

A timestamped
recording of an
event.
Can be structured or
unstructured.



Baggage

Contextual
information that can
be passed between
signals.

OpenTelemetry

Telemetry Data /Signals



Traces

Path of a request
through the entire
system



Metrics

A measurement
captured at runtime.
eg.

http.requests.total 42
db.queries.count 501



Logs

A timestamped
recording of an
event.
Can be structured or
unstructured.





Baggage

Contextual
information that can
be passed between
signals.

OpenTelemetry

What is a Trace?



OpenTelemetry


What is a Trace?

Root Span GET /api/persons

Person Services Span

Person DB

Address DB



```
{  
    "name": "GET /api/persons",  
    "context": {  
        "trace_id": "5b8aa5a2d2c872e8321cf37308d69df2",  
        "span_id": "051581bf3cb55c13"  
    },  
    "parent_id": null,  
    "start_time": "2025-04-29T18:52:58.114201Z",  
    "end_time": "2025-04-29T18:52:58.114687Z",  
    "attributes": {  
        "http.route": "/api/persons",  
        ...  
    },  
    "events": [  
        {  
            "name": "request.received",  
            "timestamp": "2025-04-29T18:52:58.114561Z",  
            "attributes": {  
                "event_attributes": 1  
            }  
        }  
    ]  
}
```

OpenTelemetry


What is a Trace?

Root Span GET /api/persons

Person Services Span

Person DB

Address DB



```
{  
    "name": "Call Person DB",  
    "context": {  
        "trace_id": "5b8aa5a2d2c872e8321cf37308d69df2",  
        "span_id": "5fb397be34d26b51"  
    },  
    "parent_id": "051581bf3cb55c13",  
    "start_time": "2025-04-29T18:52:58.114304Z",  
    "end_time": "2025-04-29T22:52:58.114561Z",  
    "attributes": {  
        "db.system": "mysql",  
        "db.statement": "SELECT * FROM person WHERE id=?",  
    }  
}
```


OpenTelemetry

Context Propagation


With context propagation, signals can be correlated with each other, regardless of where they are generated.

- Service A calls Service B includes a trace ID and a span ID
- Service B uses these values to create a new span that belongs to the same trace, setting the span from Service A as its parent.


Default propagator W3C TraceContext



OpenTelemetry Collector



OpenTelemetry Collector



OpenTelemetry Collector Configuration

- Configure the collector setup and chain is straight forward.
- It contains of reasonable defaults by default.
- Decoupling and offering of open source observability data formats
- It is not necessary to send telemetry data through a collector, it is recommended.



```
receivers:  
  otlp:  
    protocols:  
      grpc:  
        endpoint: 0.0.0.0:4317  
      http:  
        endpoint: 0.0.0.0:4318  
processors:  
  batch:  
  
exporters:  
  otlp:  
    endpoint: otelcol:4317  
  
extensions:  
  health_check:  
    endpoint: 0.0.0.0:13133  
  pprof:  
    endpoint: 0.0.0.0:1777  
  zpages:  
    endpoint: 0.0.0.0:55679  
  
service:  
  extensions: [health_check, pprof, zpages]  
  pipelines:  
    traces:  
      receivers: [otlp]  
      processors: [batch]  
      exporters: [otlp]  
    metrics:  
      receivers: [otlp]  
      processors: [batch]  
      exporters: [otlp]  
  logs:  
    receivers: [otlp]  
    processors: [batch]  
    exporters: [otlp]
```

OpenTelemetry

Additional Components and Concepts

Platforms

- Client-side Apps
- FaaS
- Kubernetes, Helm Charts and Operator

OTel Specification


- OTel
- OTelP

Sampling

Most of the requests are successful and finish with acceptable latency and no errors.

Those might not be worth the cost.

Sampling is a way to reduce the amount of traces collected.




Distributed Tracing

Grafana Tempo

Grafana Tempo

Backend for Traces



Grafana Tempo

Introduction

Tempo is cost-efficient, requiring only object storage to operate, and is deeply integrated with Grafana, Prometheus, and Loki. Tempo can ingest common open source tracing protocols, including Jaeger, Zipkin, and OpenTelemetry.

- Announced 2020
- GA 2021
- AGPLv3

<https://grafana.com/oss/tempo/>





«Grafana Tempo is an open source, easy-to-use, and high-scale distributed tracing backend.»

- Built for massing scale: affordable long term storage
- Cost-effective: Traces are not indexed




Grafana Tempo

Grafana Tempo Components / Architecture



Grafana Tempo Components / Architecture



Grafana Tempo – Deployment Variants



Kubernetes

Helm, Operator or Tanka
(Jsonnet) based deployments on
your kubernetes clusters.



Virtual Machine

Run a monolithic installation
based on docker-compose on a
linux virtual machine.



Grafana Cloud Traces


Grafana Tempo is part of the
Grafana Cloud Services.
Pay as you go.

Distributed Tracing Demo


<https://opentelemetry.io/docs/demo/>

<https://github.com/tim-koko/opentelemetry-demo>


Demo Application



Demo




Collector



Backends



GrafanaTempo



Visualization



Grafana

The best telescopes to see the world closer

[Go Shopping](#)

Hot Products



National Park Foundation Exploroscope
\$ 101.96



Starsense Explorer Refractor Telescope
\$ 349.95



Eclipsmart Travel Refractor Telescope
\$ 129.95



Lens Cleaning Kit
\$ 21.95



Roof Binoculars
\$ 209.95



Solar System Color Imager
\$ 175.00



Optical Tube Assembly

Capturing impressive deep-sky astroimages is easier than ever with Rowe-Ackermann Schmidt Astrograph (RASA) V2, the perfect companion to today's top DSLR or astronomical CCD cameras. This fast, wide-field f/2.2 system allows for shorter exposure times compared to traditional f/10 astrophotography, without sacrificing resolution. Because shorter sub-exposure times are possible, your equatorial mount won't need to accurately track over extended periods. The short focal length also lessens equatorial tracking demands. In many cases, autoguiding will not be required.

\$ 3599.00

Quantity

 Add To Cart

You May Also Like

frontend-web: HTTP POST


POST 200

[Explain in Assistant](#)
[Feedback](#)
[Share](#)


Trace ID 0ca9772eb7a5a55fb98a8cb7673be2f1 | Start time 2025-08-20 21:36:49.974 (12 hours ago) Duration 75ms Services 12 Route /oteldemo.CartService/GetCart

> Span Filters

46 spans | Prev | Next



Service & Operation



HTTP POST — Service: **checkout** Duration: 2.59ms Start Time: 55.59ms (21:36:50.029) Child Count: 1 Kind: client Status: unset
Library Name: go.opentelemetry.io/contrib/instrumentation/net/http/otelhttp Library Version: 0.62.0

[Logs for this span](#)

> Span attributes http.request.method POST http.response.status_code 200 network.protocol.version 1.1 server.address shipping server.port 50050 url.full http://shipping:50050/get-quote

> Resource attributes host.name b327c7bf3501 process.executable.name checkout process.owner nonroot service.name checkout service.namespace opentelemetry-demo service.version 2.0.2 telemetry.sdk.language go telemetry.sdk.name o...

> **shipping** /get-quote (1.65ms)

> POST quote (1.51ms)

1.65ms

1.51ms

 grafanacloud-timkoko-logs

| Split

d <

© 202

08-20

21:36:48 to 2025-08-20 21:36:52 ✓

>

Run query ▾

▶ Live

+ Add query ⓘ Query inspector

▼ Logs volume

Logs

-7

```
|> 2025-08-20 21:36:50.042 Successful to write message. offset: 0, duration: 112.747µs
|> 2025-08-20 21:36:50.042 sending to postProcessor
|> 2025-08-20 21:36:50.042 order confirmation email sent to "thomas@tim-koko.ch"
|> 2025-08-20 21:36:50.039 order placed
|> 2025-08-20 21:36:50.035 payment went through
|> 2025-08-20 21:36:50.026 [PlaceOrder]
```




[+ Add query](#) [Query inspector](#)

Table

Name	Rate	Error Rate	Duration (p90)	Links
ingress	0.01	0.005	0.00	24.0 ms Tempo
image-provider	0.00	0.002	0.00	4.50 ms Tempo
GET	0.00	0.002	0.00	33.8 ms Tempo
oteldemo.ProductCatalogService,	0.00	0.002	0.00	0 s Tempo
GET /api/products/{productId}	0.00	0.001	0.00	0 s Tempo

Node graph



Distributed Tracing

Distributed Tracing Stack in Practice

Tracing Platform PoC

based on OpenTelemetry and Grafana Tempo

Baloise faced the challenge that their existing solution was costly, had low adoption, and lacked self-service capabilities. Their objective was to assess a cloud-native alternative to the proprietary system in place.

- PoC was a success
- Familiar user experience
- Full ownership
- Open standards and vendor independence

<https://tim-koko.ch/en/references/baloise-distributed-tracing-opentelemetry/>




Distributed Tracing Made Easy

Learnings



Observability is key

Observability is key in distributed systems to understand what is going on. Distributed systems are complex, so is distributed tracing. It is a platform topic, but is much closer to the applications than you might guess.



Correlate Signals and Levels

The correlation of observability signals is extremely valuable and can be achieved on several levels. Signals have different goals and requirements, which support the approach of different backends.



Stick to standards

Decoupling your business from vendors using open standards and open source. Building your own open tracing and observability platform come with costs, though...

About us

tim&koko


We are hiring!

tim-koko.ch/jobs

Thomas Philipona

thomas@tim-koko.ch

+41 79 325 55 83



Thank you for your
attention

