



Dig Smart: Creating A Reliable Cloud-Native DNS Service

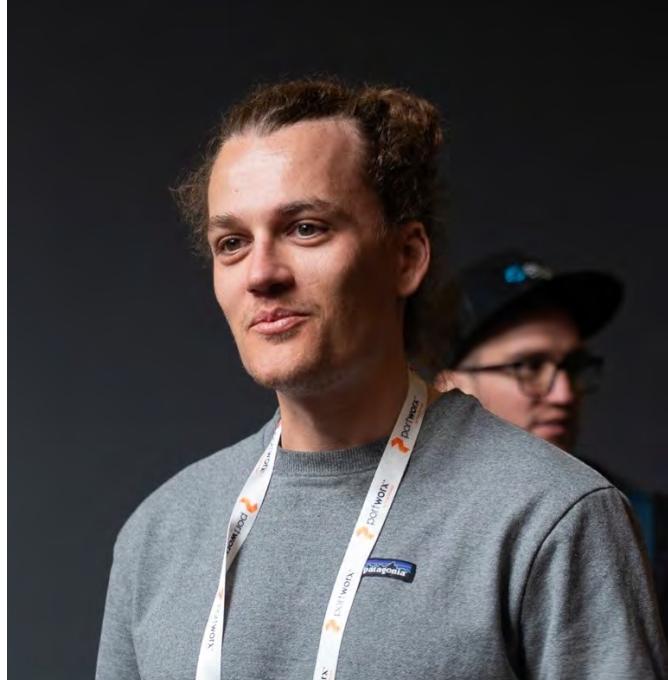
Joel Studler, Fabian Schulz & Georgios Daskalopoulos





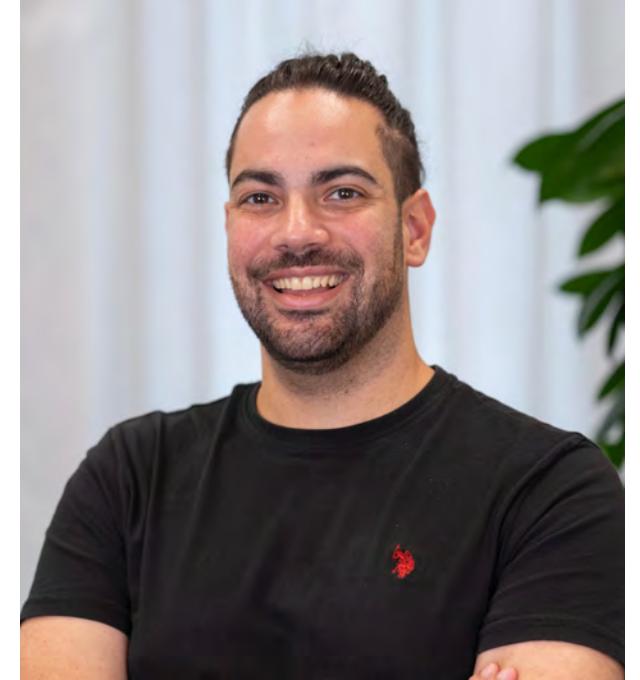
Joel Studler
DevOps Engineer

joel.studler@swisscom.com



Fabian Schulz
DevOps Engineer

fabian.schulz1@swisscom.com



Georgios Daskalopoulos
DevOps Engineer

gdasky@gmail.com



Part 1: Creating A Reliable Cloud Native DNS Service

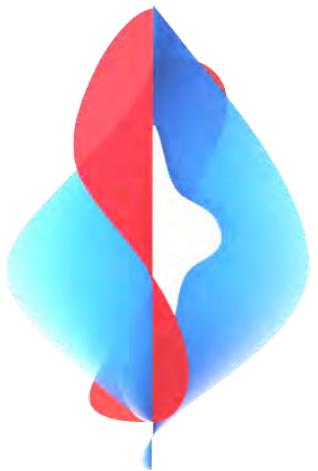
by Fabian Schulz & Joel Studler

Part 2: Resilience Testing on Cloud Native DNS

by Georgios Daskalopoulos



University of
Zurich ^{UZH}



swisscom



Context & Related Talks

5G - driving our journey from Telco to TechCo

by Swisscom CTIO Mark Düsener at Connect Conference 2022

<https://www.youtube.com/watch?v=hND7TiXJED8>

Evolving GitOps: Harnessing Kubernetes Resource Model for 5G

by Ashan Senevirathne and Joel Studler at Open Source Summit 2024

https://www.youtube.com/watch?v=35-fE_gHDjw

From Weeks to Seconds: Scaling Network Automation with Kubernetes Operators

by Alessio Diamanti and Adrian Kurt at ContainerDays 2025

<https://www.youtube.com/watch?v=RytDzZYDQTU>



DNS in 5G Core

5G

Specific Private Zones

Domains used in Mobile Network only such as 3gppnetwork.org



Moderate Throughput

10s to 100s of Requests/second



Low Latency

DNS is an important factor in the overall performance of the Mobile Network



Requirements for the 5G Core DNS Service



Proximity to Consumer

Minimal amount of hops between 5G Core and DNS

X No SaaS allowed



Fully Automated

GitOps driven and automated provisioning of DNS records

X No manual interaction allowed



Geo Redundant & HA

Spread across multiple K8s clusters and geo regions to increase reliability

X No singletons



Support of Advanced DNS features

Resource Records such as NAPTR and SRV supported for e.g. SIP Phone Calls

X Need to go beyond A and CNAME



K8s integration with ExternalDNS

The System leverages Kubernetes Patterns such as CRs and Operators

X No CRUD outside kube-api



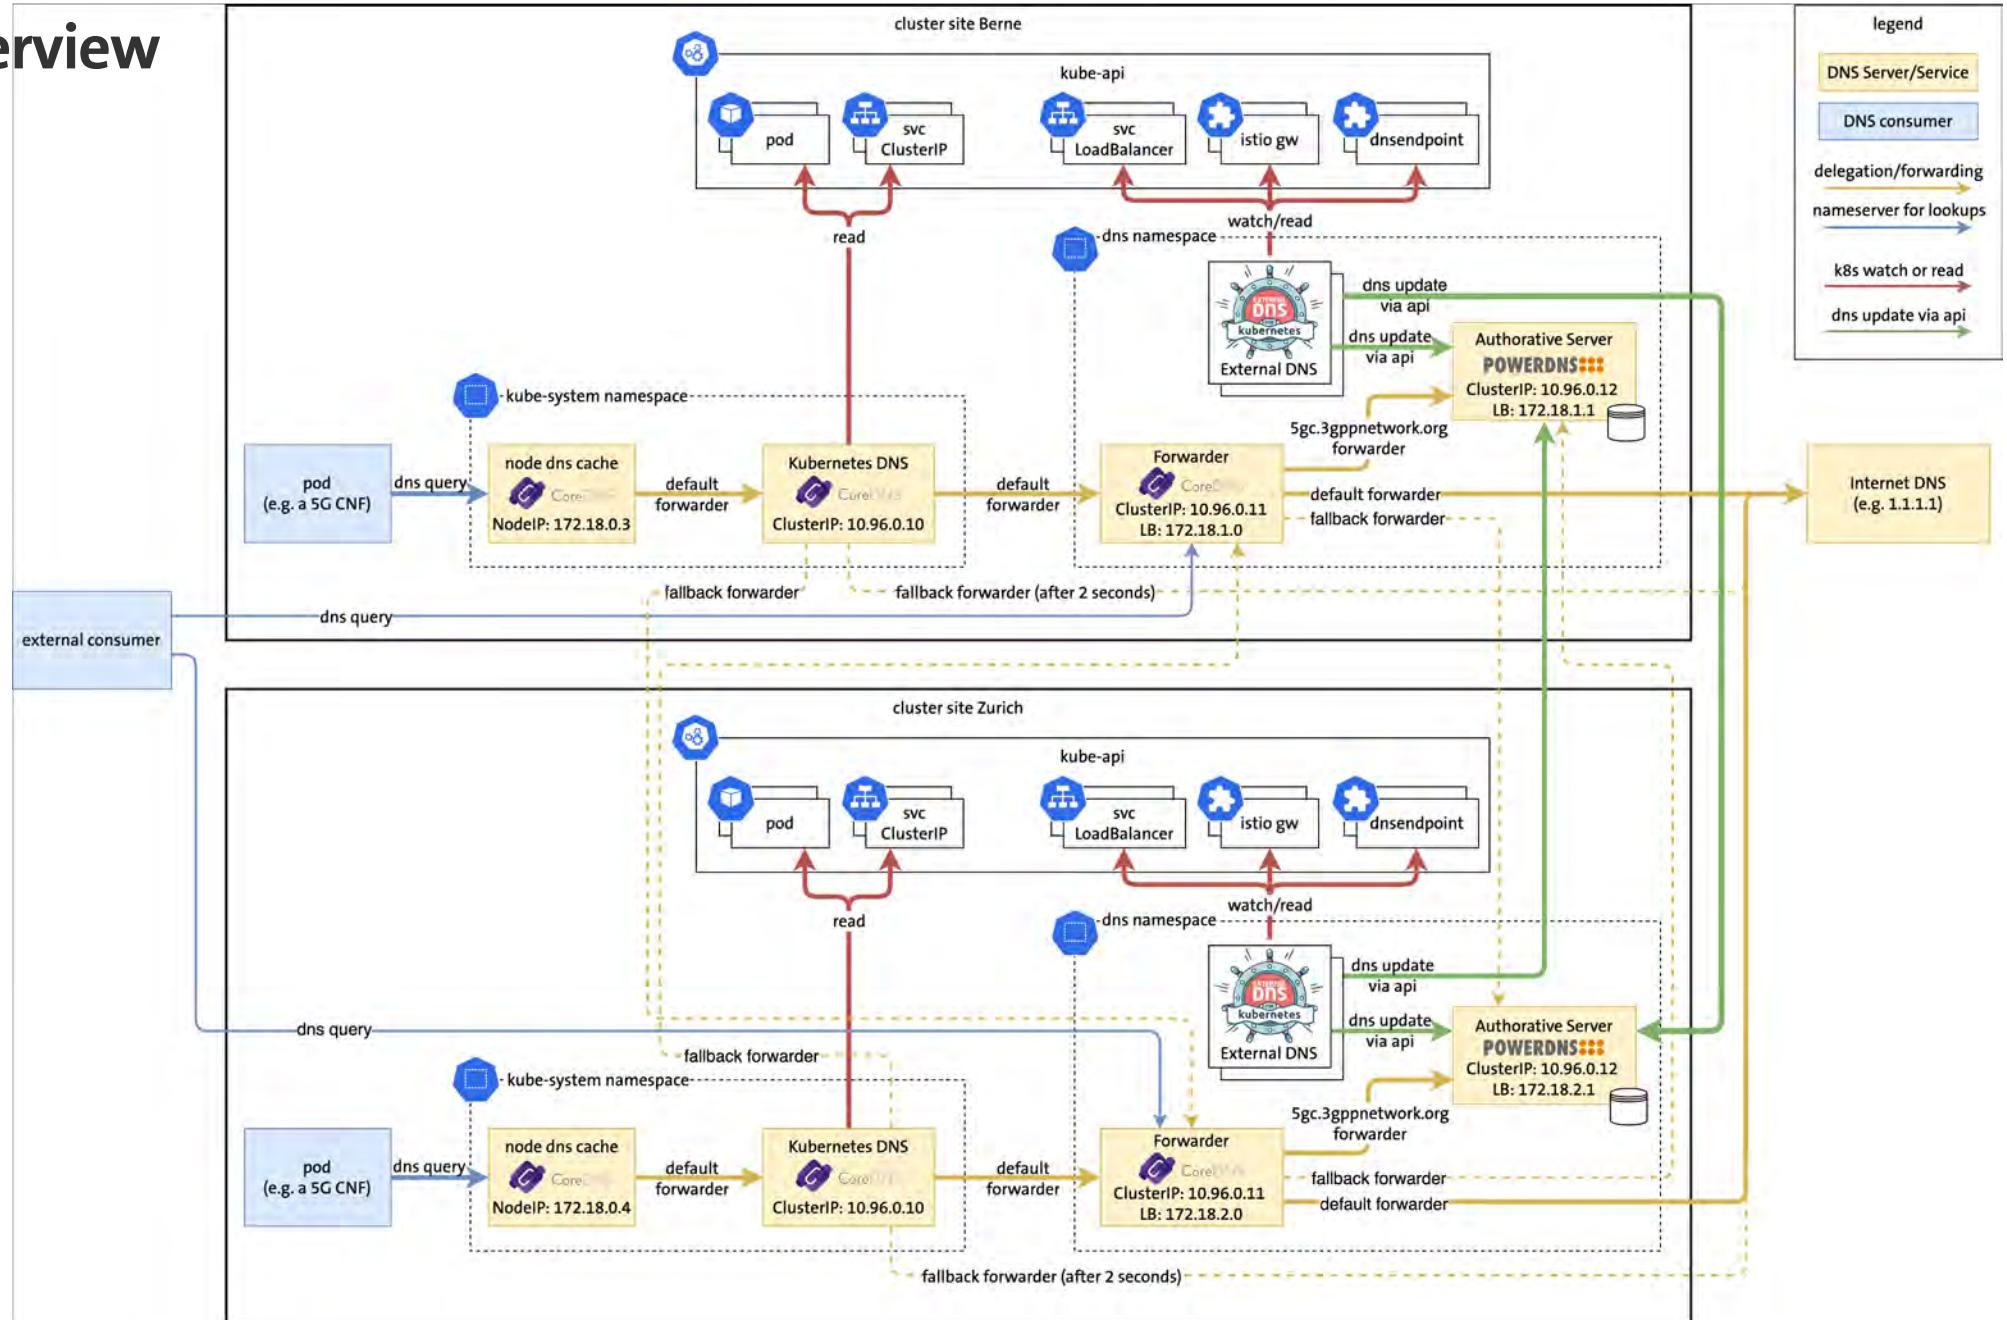
Minimal Amount of SPOFs

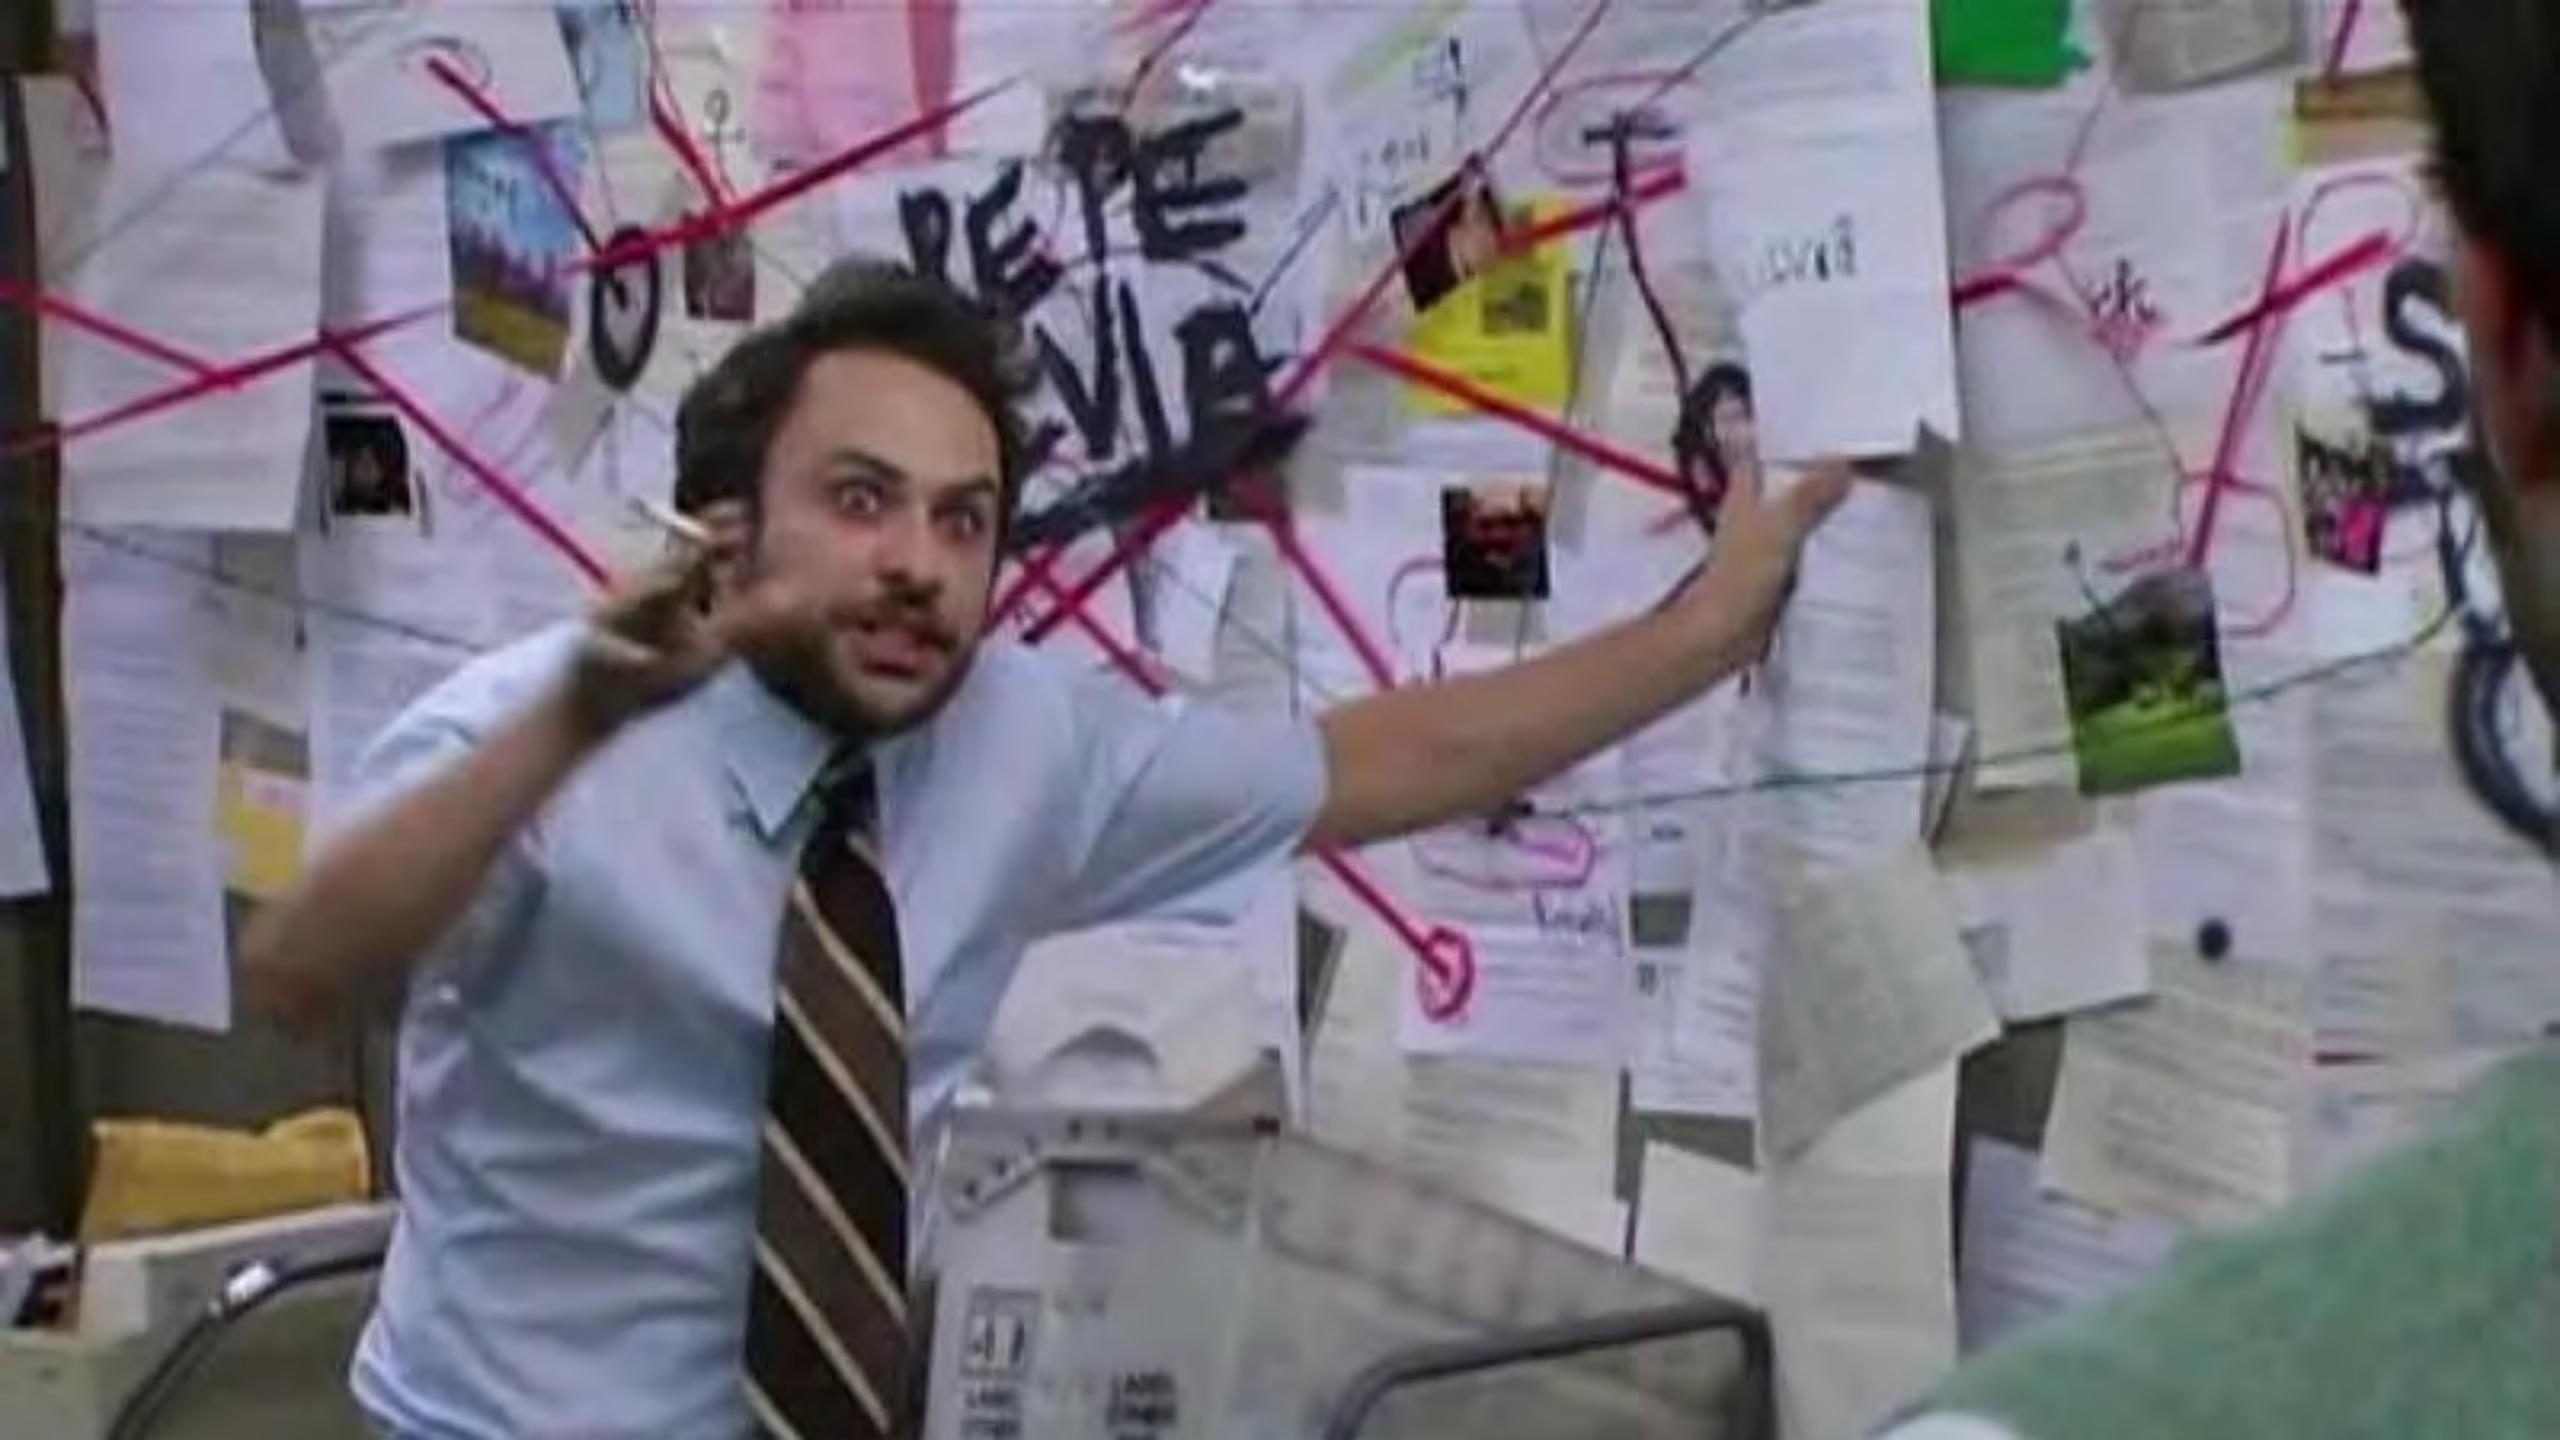
Share nothing by removing single points of failure from the System

X No shared mgmt system



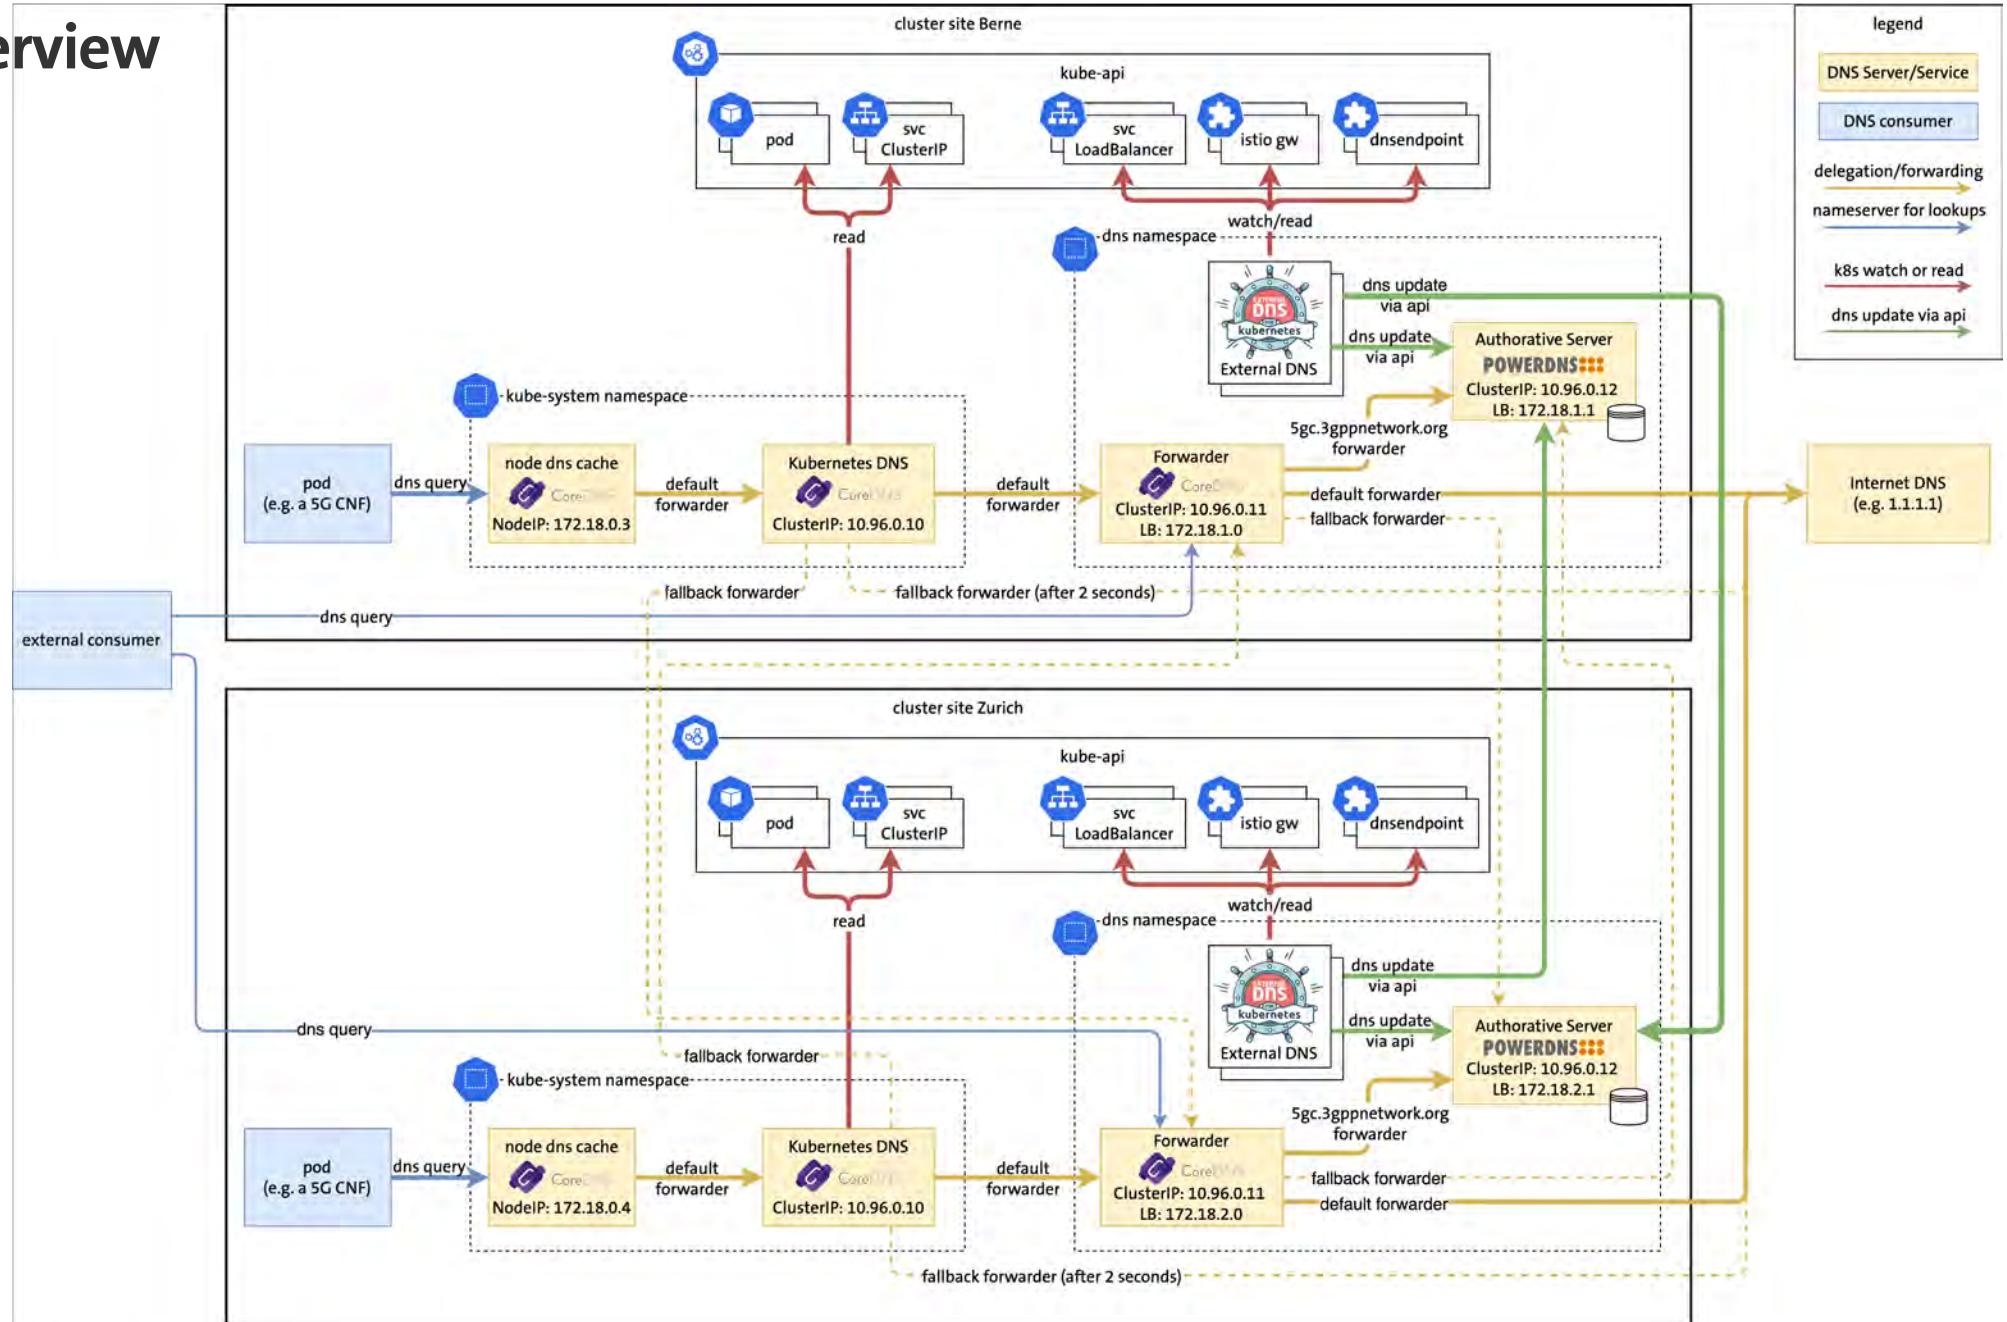
Overview





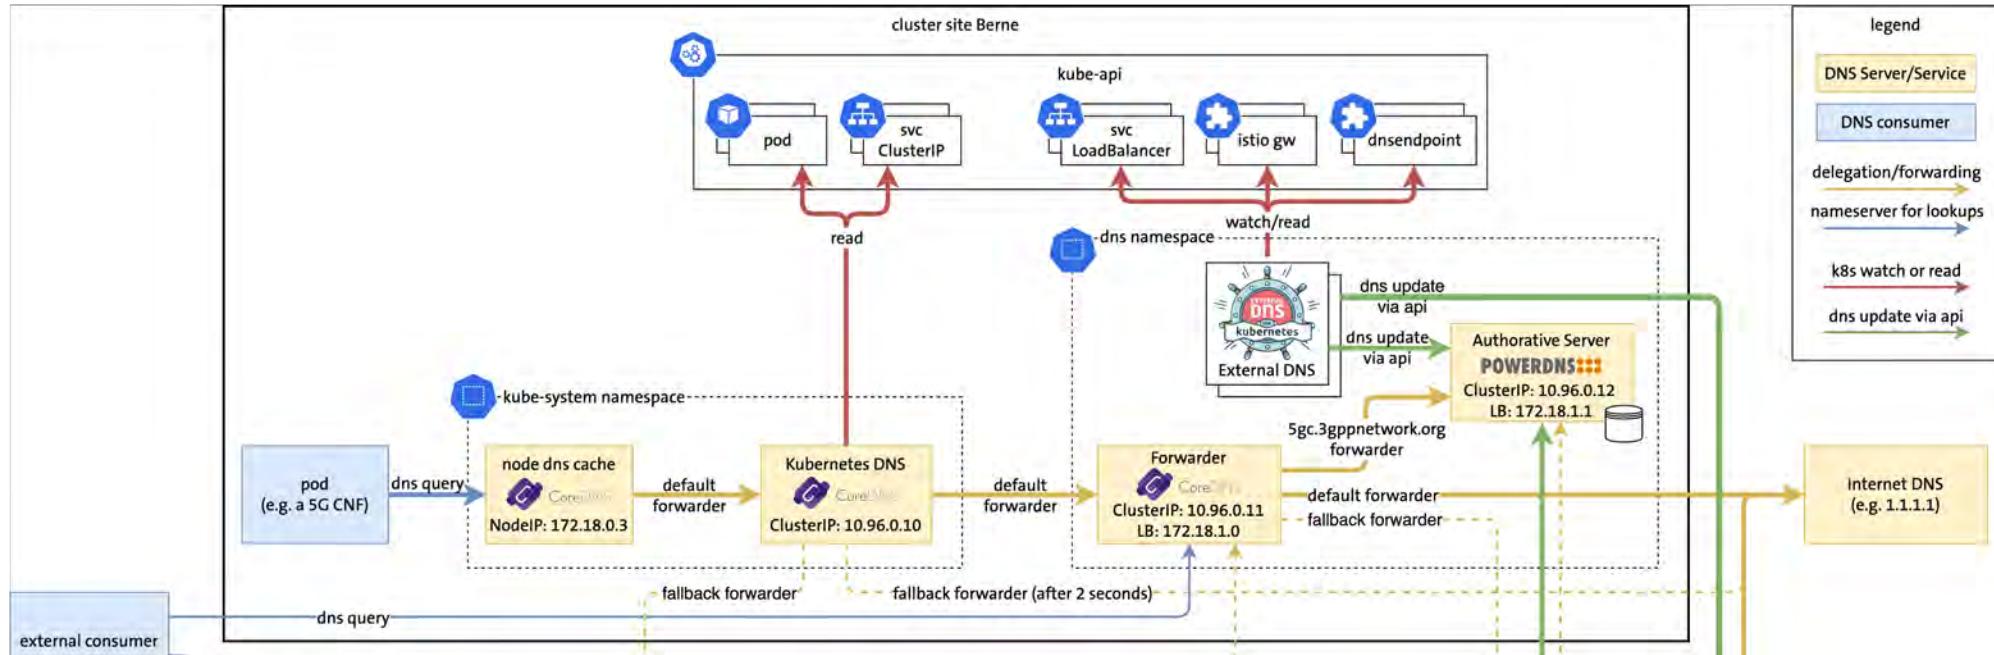


Overview





Overview Single Cluster





In-Cluster Service Discovery in Kubernetes

CoreDNS (<https://coredns.io>)

Reads from kube-api

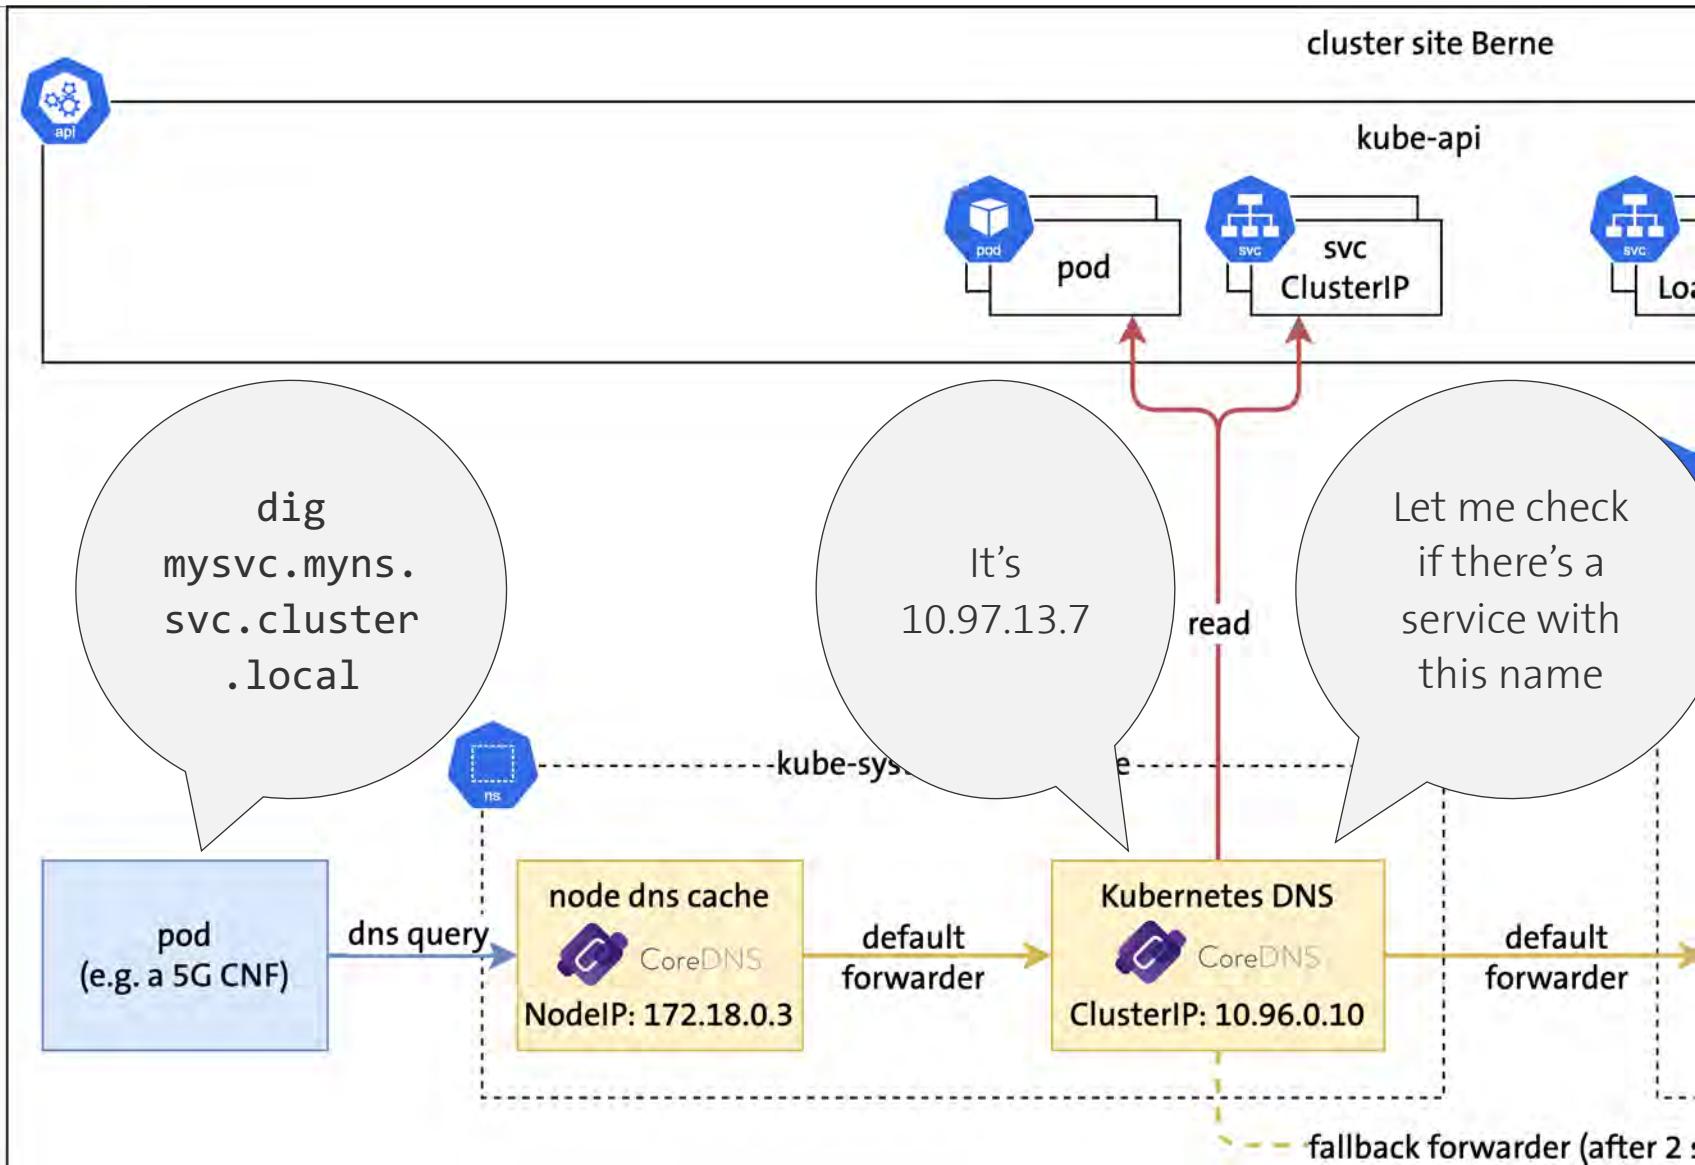
Features:

✓ In-Cluster Service Discovery

Missing:

X Not exposed outside of K8s

X No custom Resource Records





In-Cluster Service Discovery in Kubernetes: Resources

Kubernetes DNS: <https://kubernetes.io/docs/concepts/services-networking/dns-pod-service>

Reserved ClusterIP Address assignment: <https://kubernetes.io/docs/concepts/services-networking/cluster-ip-allocation/#why-do-you-need-to-reserve-service-cluster-ips>

Node Cache: <https://kubernetes.io/docs/tasks/administer-cluster/nodelocaldns>

Debugging Kubernetes DNS: <https://kubernetes.io/docs/tasks/administer-cluster/dns-debugging-resolution>

Customize DNS Service: <https://kubernetes.io/docs/tasks/administer-cluster/dns-custom-nameservers>



Requirements for Authoritative Server

Requirement

ExternalDNS* Support for K8s integration

A & CNAME Resource Records

NAPTR Resource Records (e.g. for SIP phone calls)

Proximity to Consumer



CoreDNS

CoreDNS



POWERDNS



PowerDNS



Authoritative

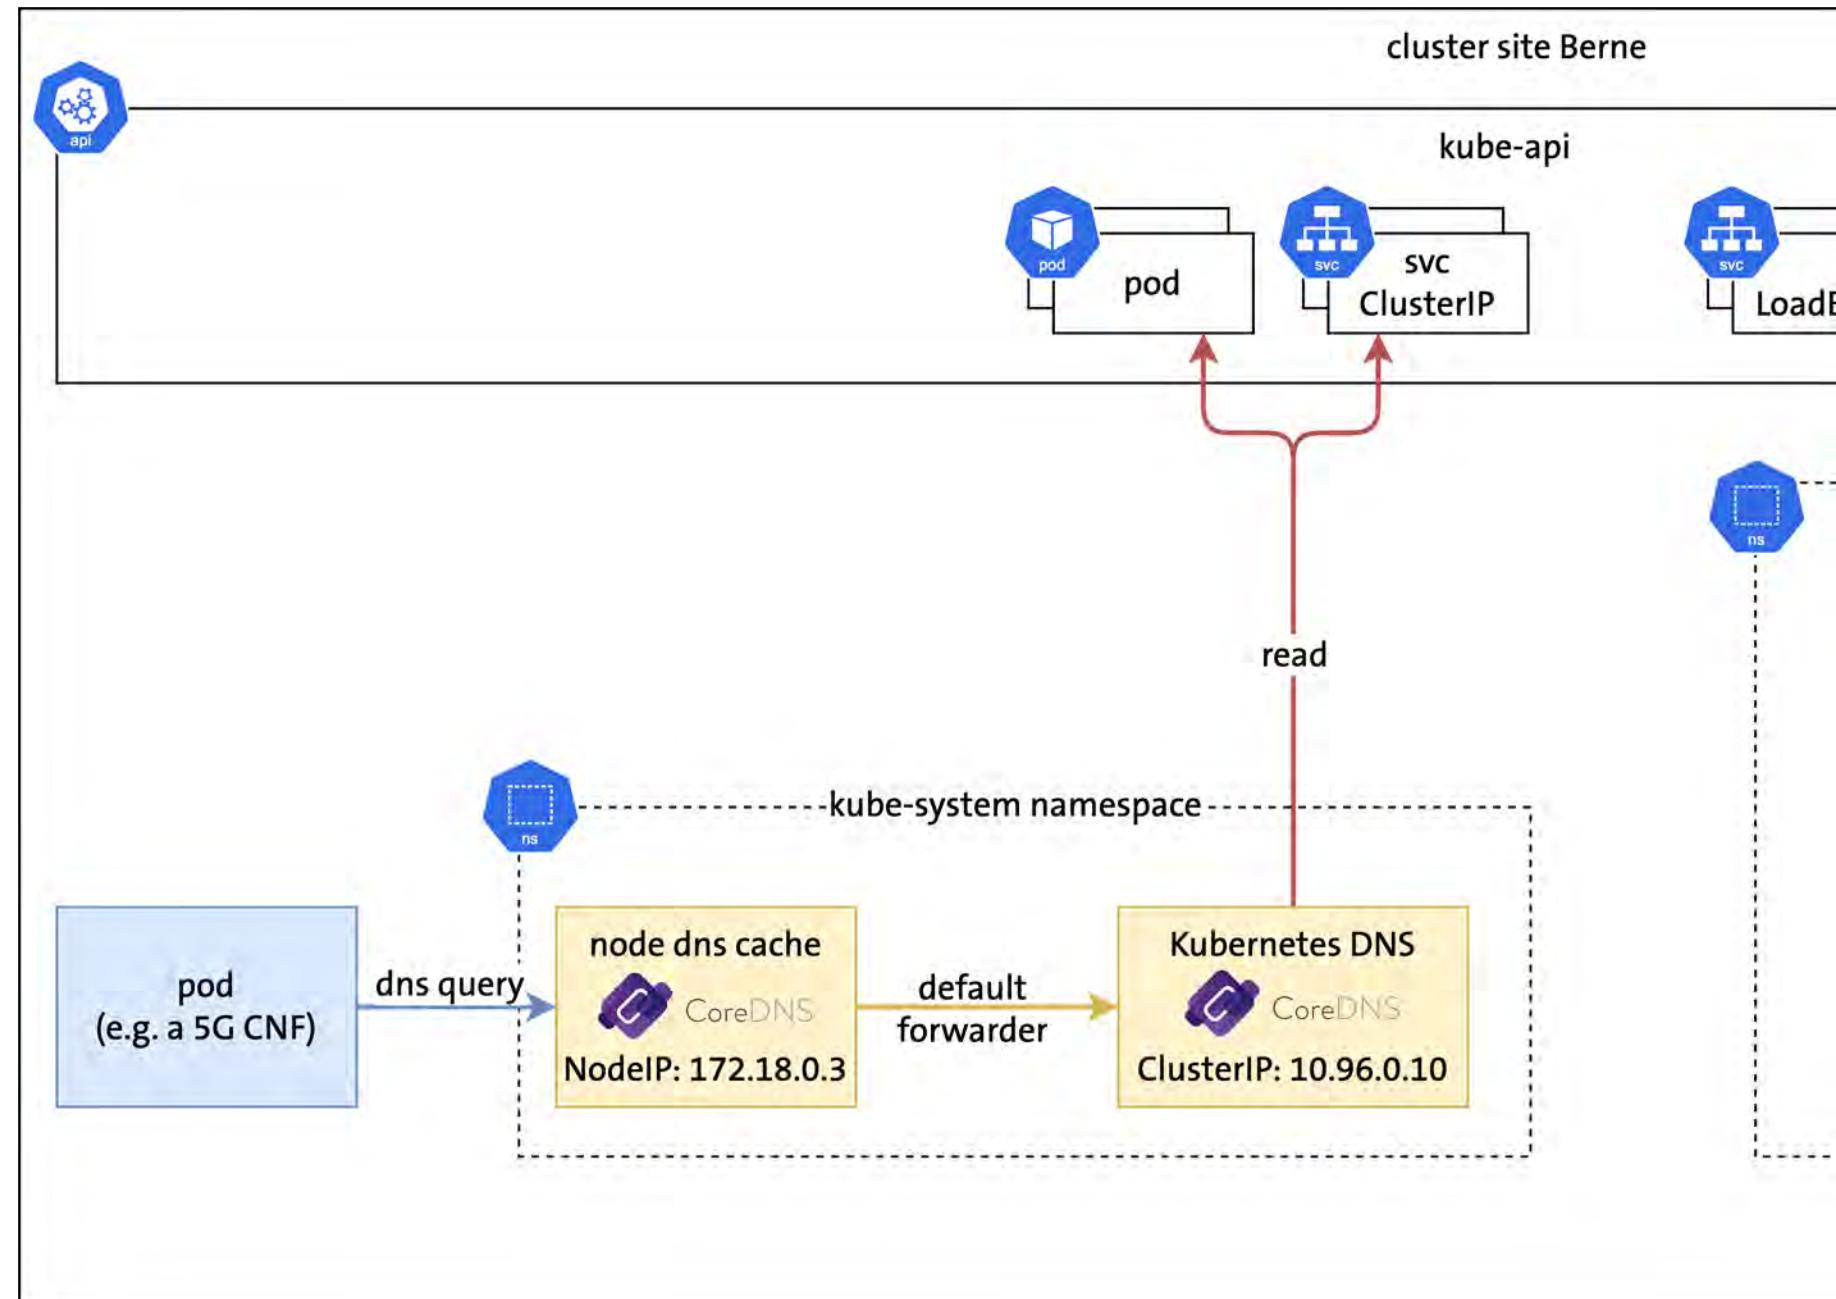


SaaS



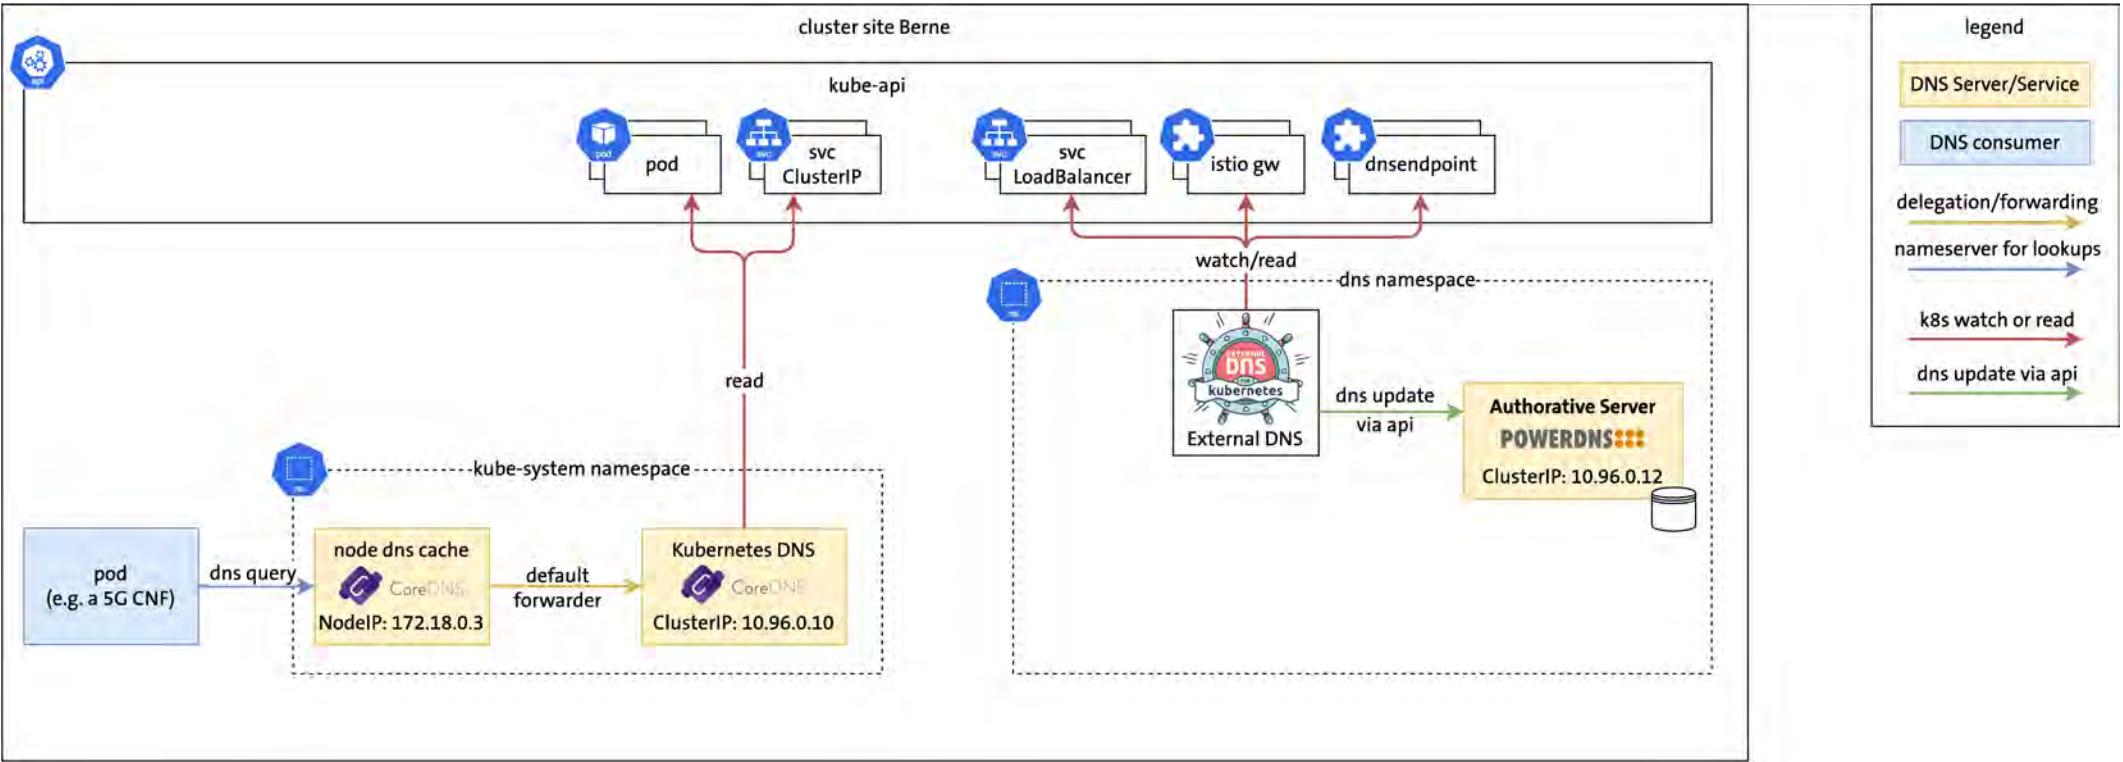
* <https://github.com/kubernetes-sigs/external-dns>

** After a fix in external-dns <https://github.com/kubernetes-sigs/external-dns/pull/4212>





Overview

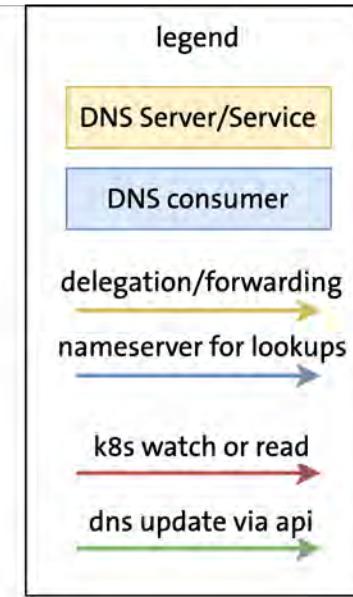
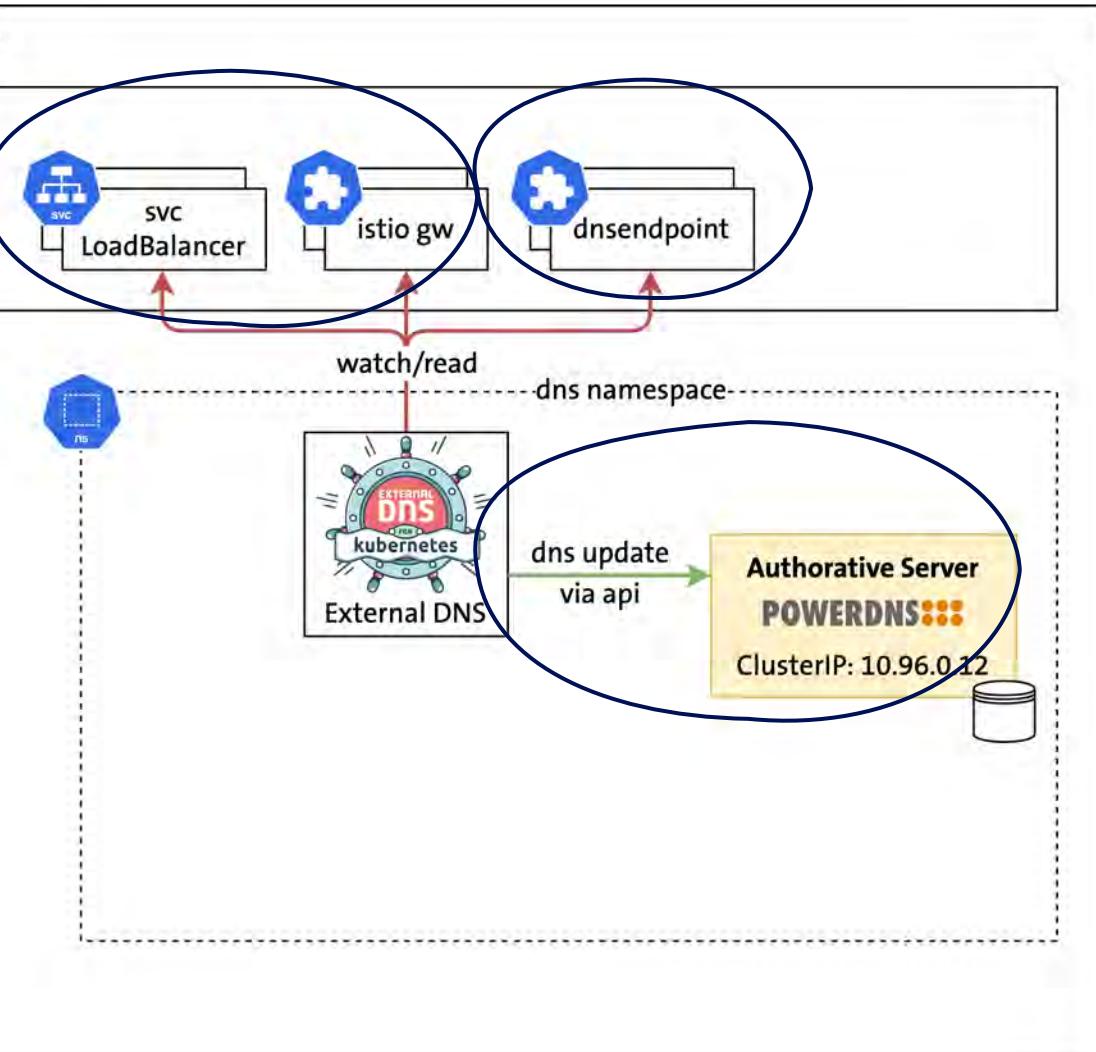




Automation of Authoritative Server Using ExternalDNS

Berne

i



ExternalDNS reads from kube-api:

- Static Resource Records as DNST Endpoint Custom Resources
- Dynamic Type A Records using Annotations
 - Name definition via Annotation
 - IP fetched from Service / Ingress / Istiogw status field

ExternalDNS writes to PowerDNS API



ExternalDNS State Management for Dynamic IP Assignment: GitOps + Kubernetes



```
apiVersion: v1
kind: Service
metadata:
  annotations:
    external-dns.alpha.kubernetes.io/hostname: my-app.example.com
  name: my-app
spec:
  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: 80
  selector:
    name: my-app
  type: LoadBalancer
```

DNS Name
defined in git



```
apiVersion: v1
kind: Service
metadata:
  annotations:
    external-dns.alpha.kubernetes.io/hostname: my-app.example.com
  name: my-app
spec:
  ports:
    - name: http
      port: 80
      protocol: TCP
      targetPort: 80
  selector:
    name: my-app
  type: LoadBalancer
status:
  loadBalancer:
    ingress:
      - ip: 192.168.0.35
```

IP read by
ExternalDNS

ExternalDNS creates

DNS Backend:
my-app.example.com. 3600 IN A 192.168.0.35



ExternalDNS State Management for Static Assignment: GitOps



```
apiVersion: externaldns.k8s.io/v1alpha1
kind: DNSEndpoint
metadata:
  name: my-mx-record
spec:
  endpoints:
    - dnsName: my-app.example.com
      recordTTL: 3600
      recordType: A
      targets:
        - 192.168.0.35
    - dnsName: example.com
      recordTTL: 3600
      recordType: MX
      targets:
        - 10 mailhost1.example.com
        - 20 mailhost2.example.com
```

Resource
Records
defined in git



```
apiVersion: externaldns.k8s.io/v1alpha1
kind: DNSEndpoint
metadata:
  name: my-mx-record
spec:
  endpoints:
    - dnsName: my-app.example.com
      recordTTL: 1
      recordType: A
      targets:
        - 192.168.0.35
    - dnsName: example.com
      recordTTL: 180
      recordType: MX
      targets:
        - 10 mailhost1.example.com
        - 20 mailhost2.example.com
```

ExternalDNS creates

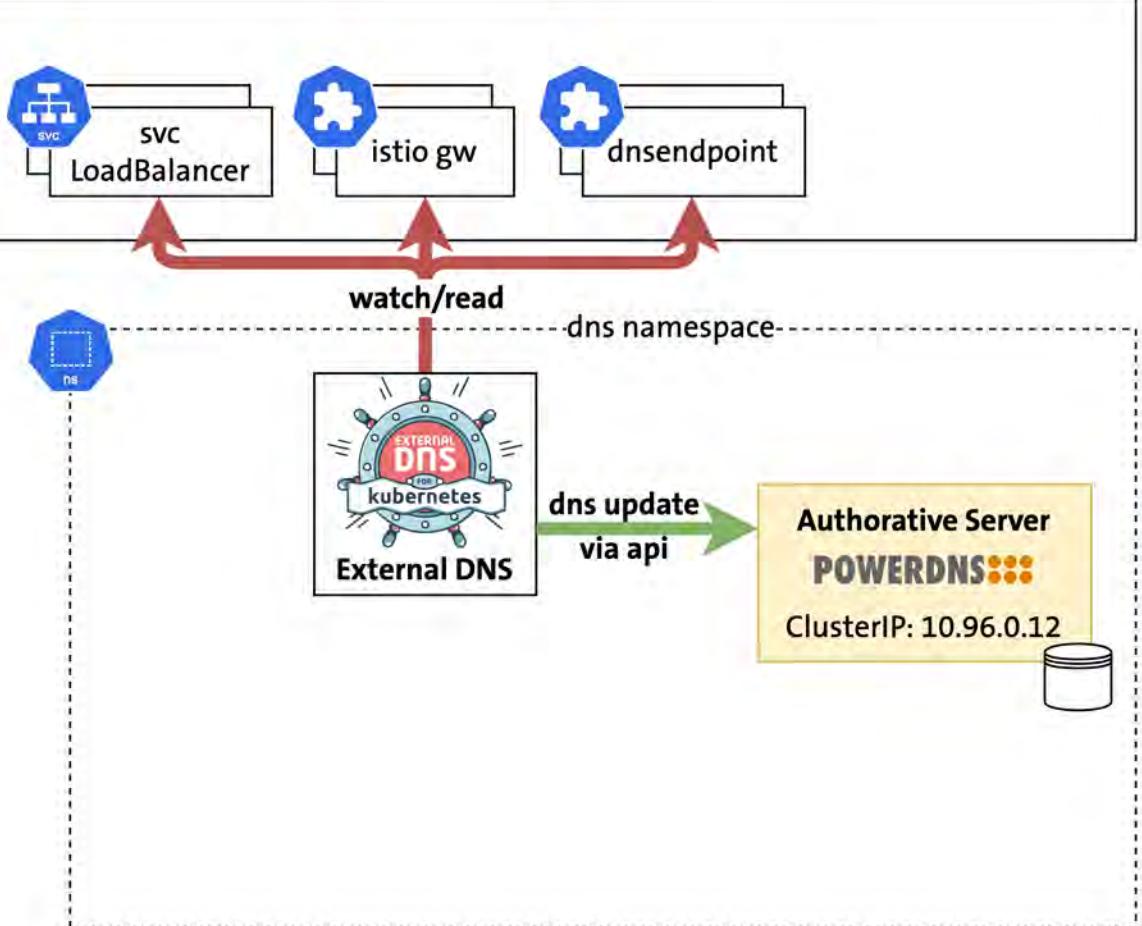
DNS Backend:

my-app.example.com.	3600	IN	A	192.168.0.35
example.com.	3600	IN	MX	10 mailhost1.example.com
example.com.	3600	IN	MX	20 mailhost2.example.com



Berne

i



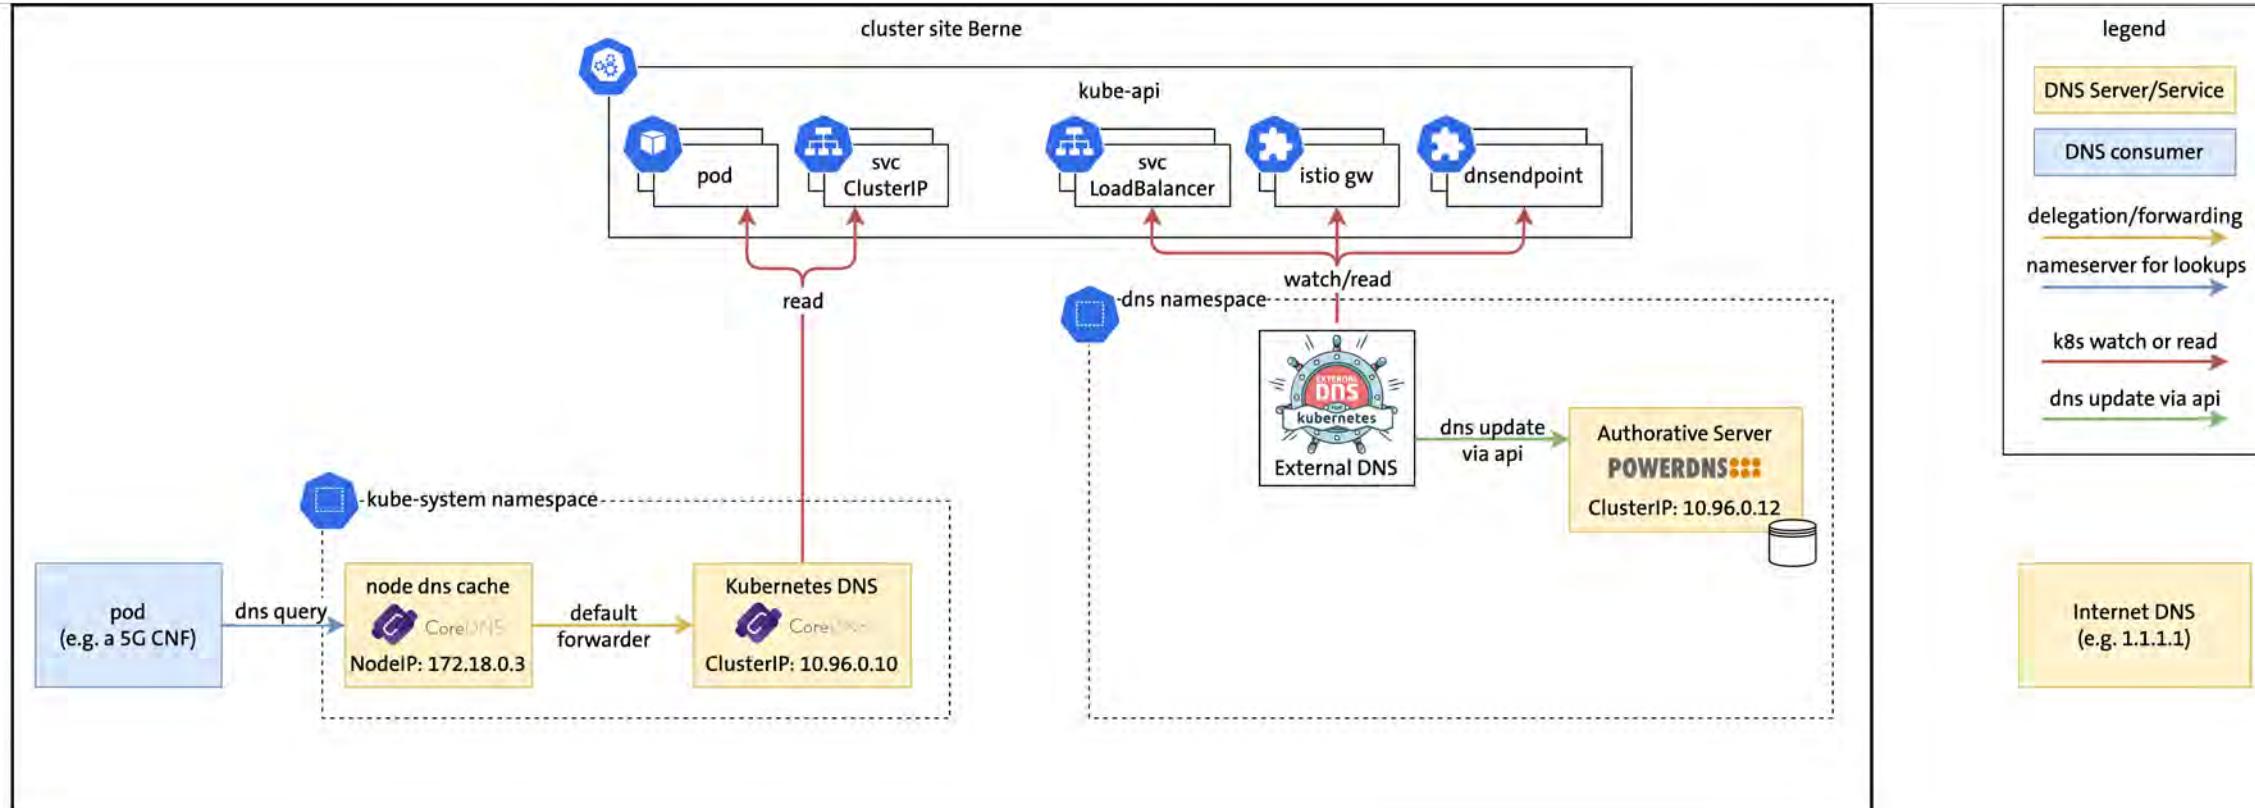
Demo ExternalDNS + PowerDNS Single Cluster



<https://github.com/swisscom/cloud-native-telco/tree/main/prototypes/dns/1-demo-external-dns>

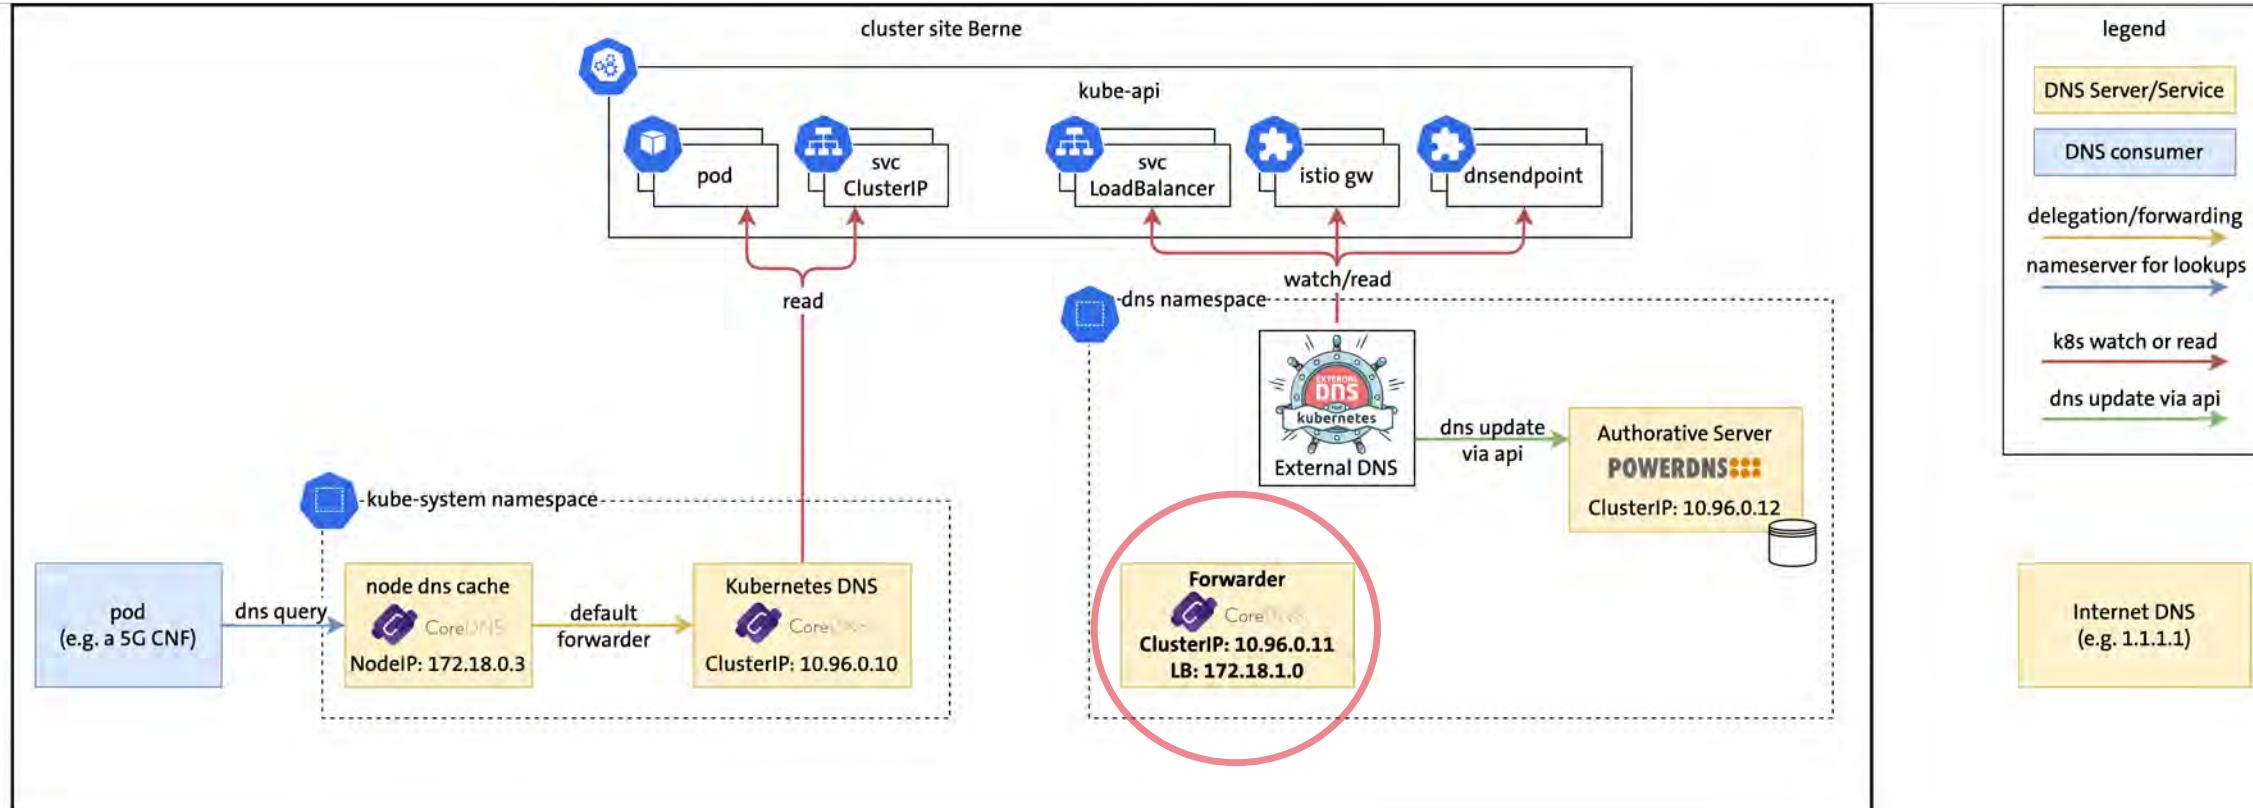


Forwarding to Authoritative Server



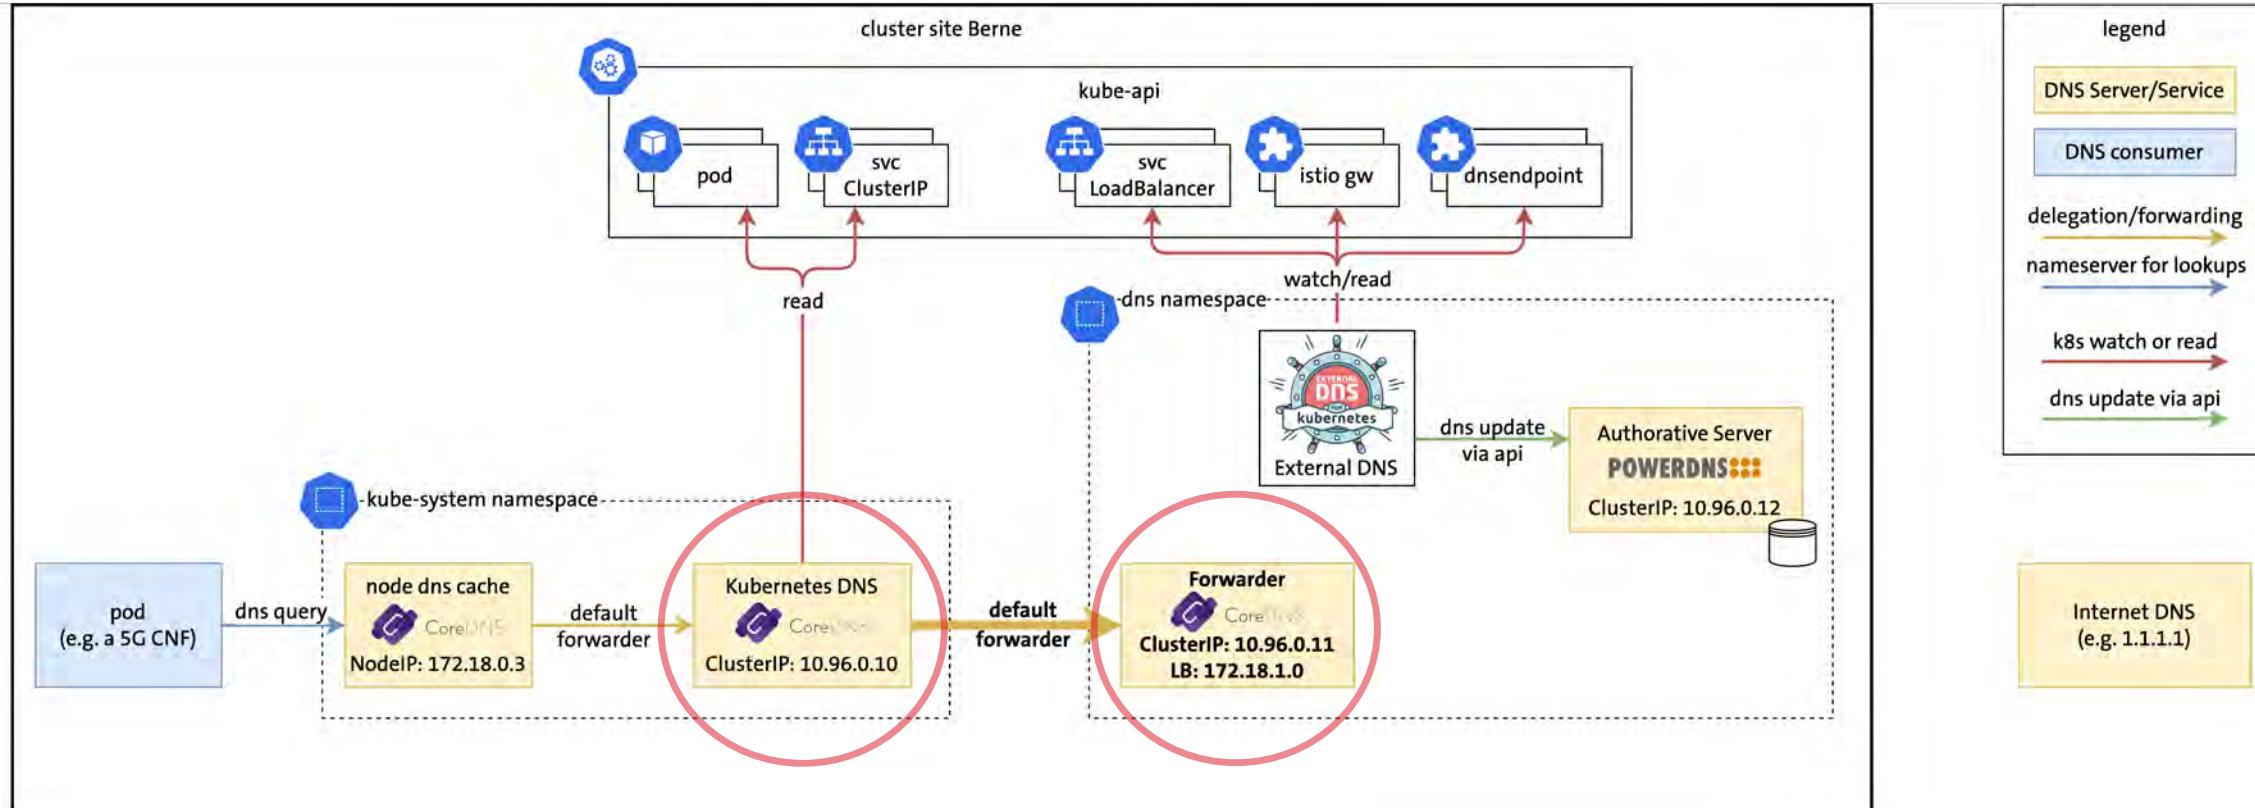


Forwarding to Authoritative Server



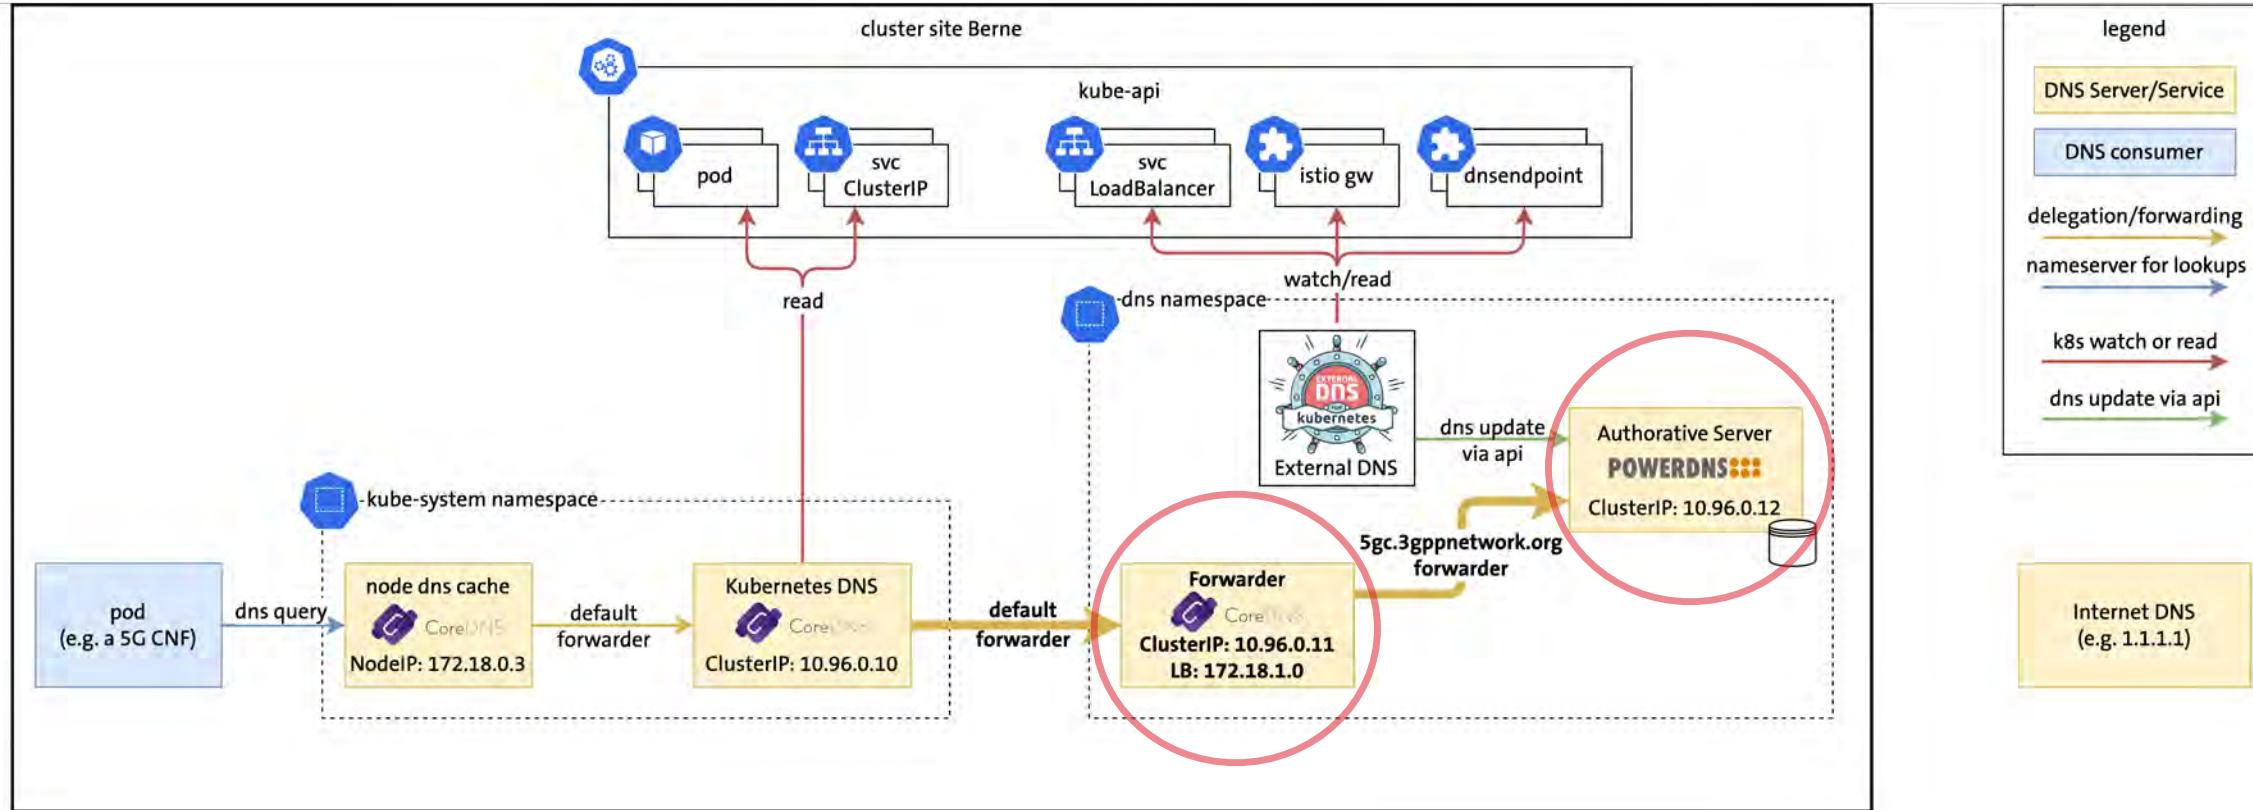


Forwarding to Authoritative Server



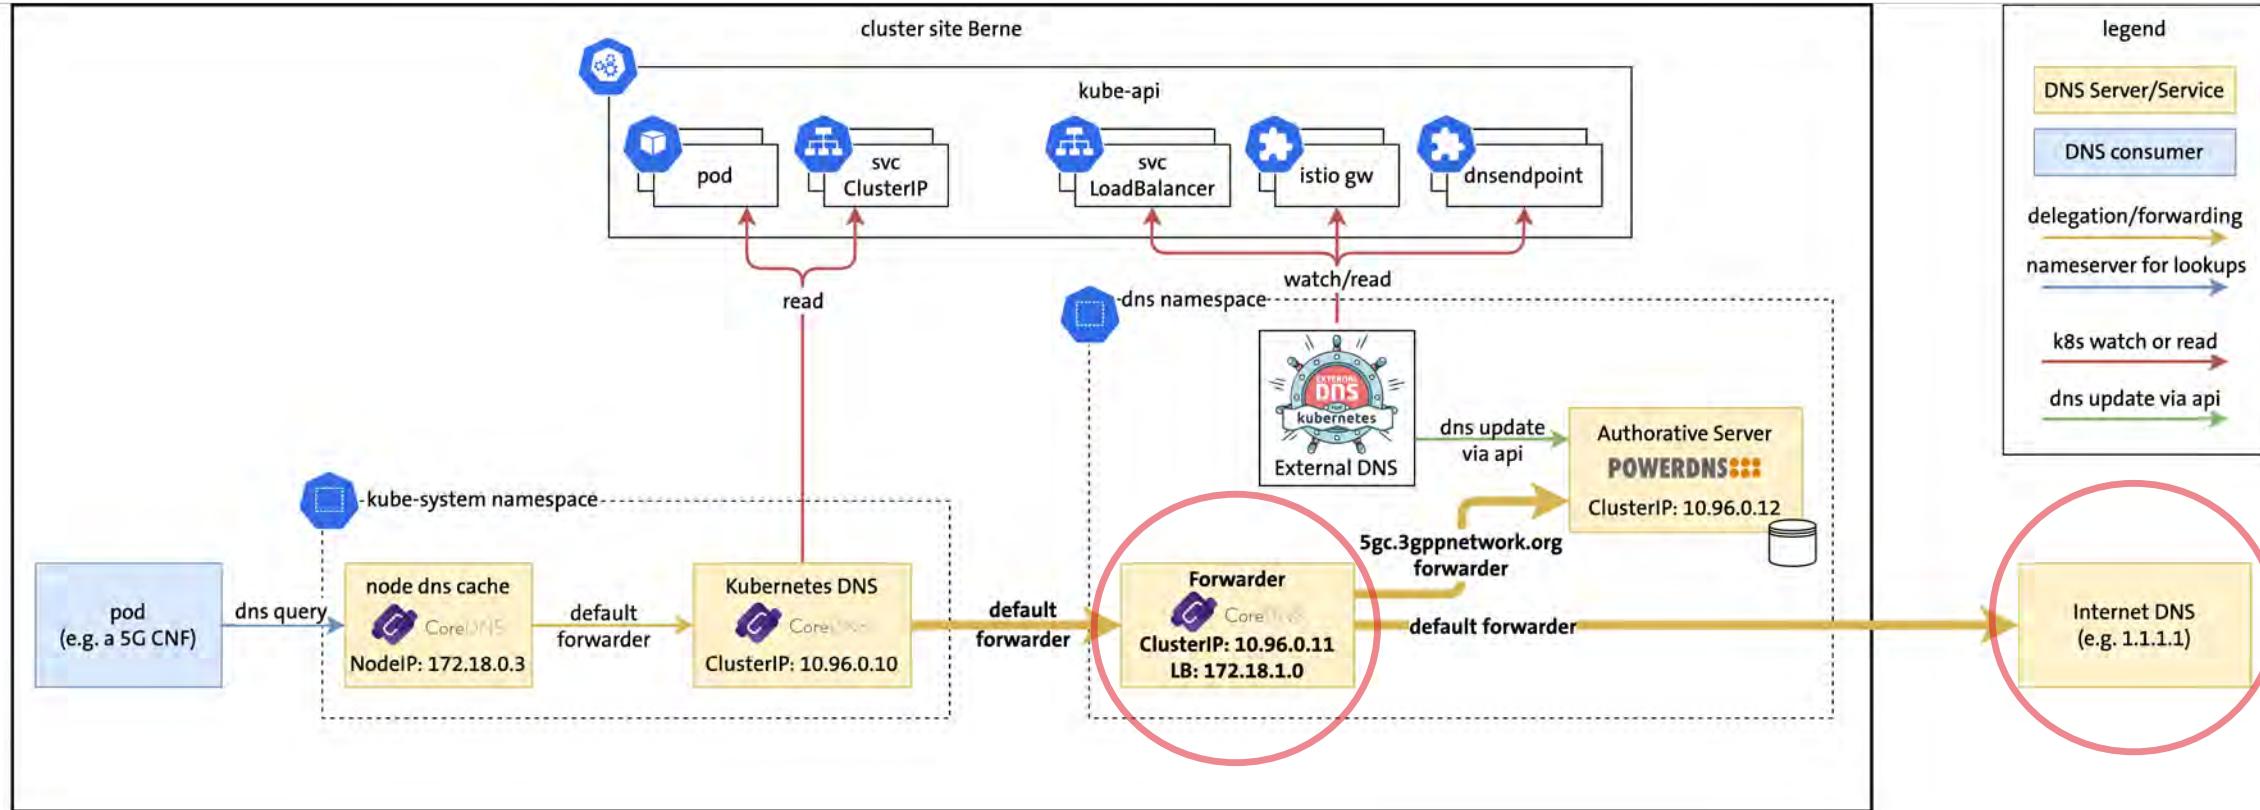


Forwarding to Authoritative Server



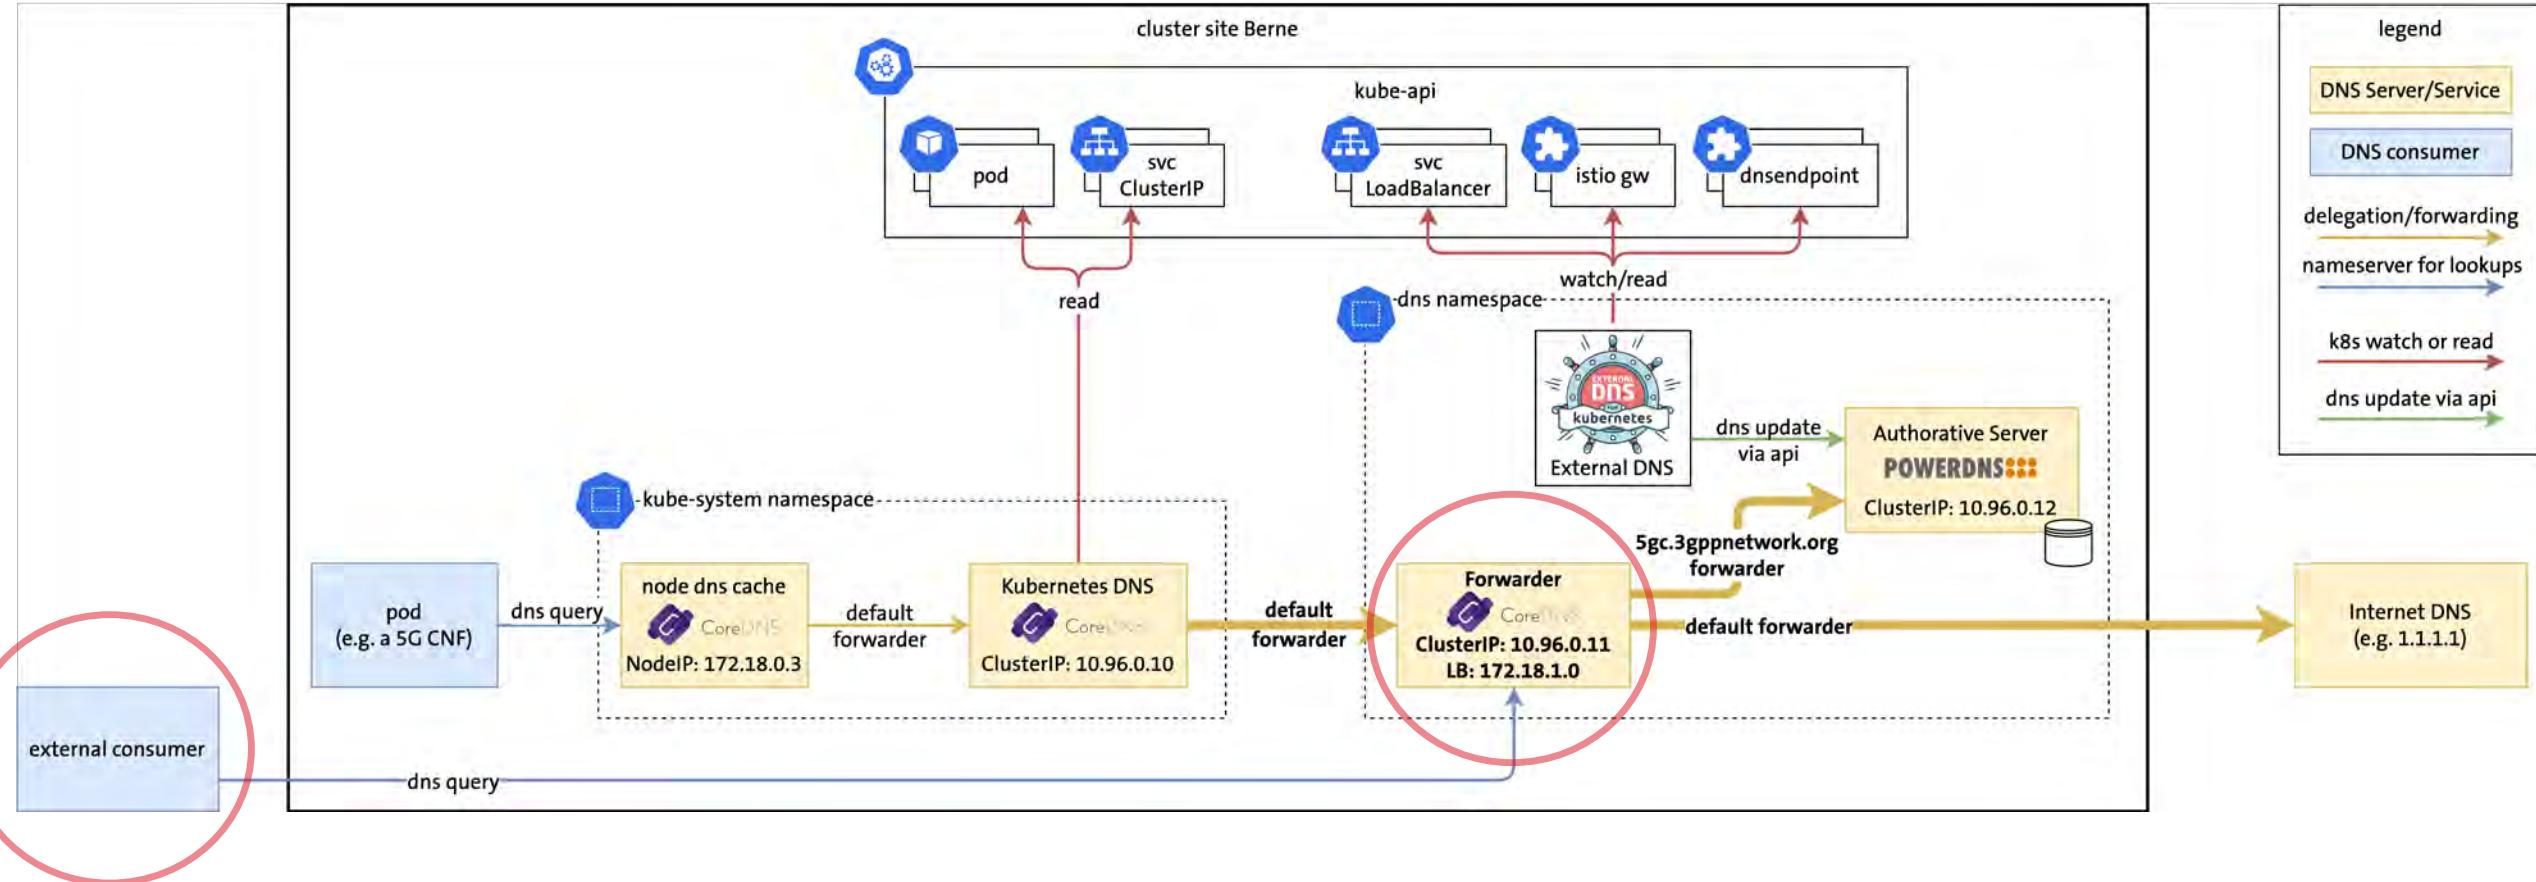


Forwarding to Authoritative Server





Forwarding to Authoritative Server

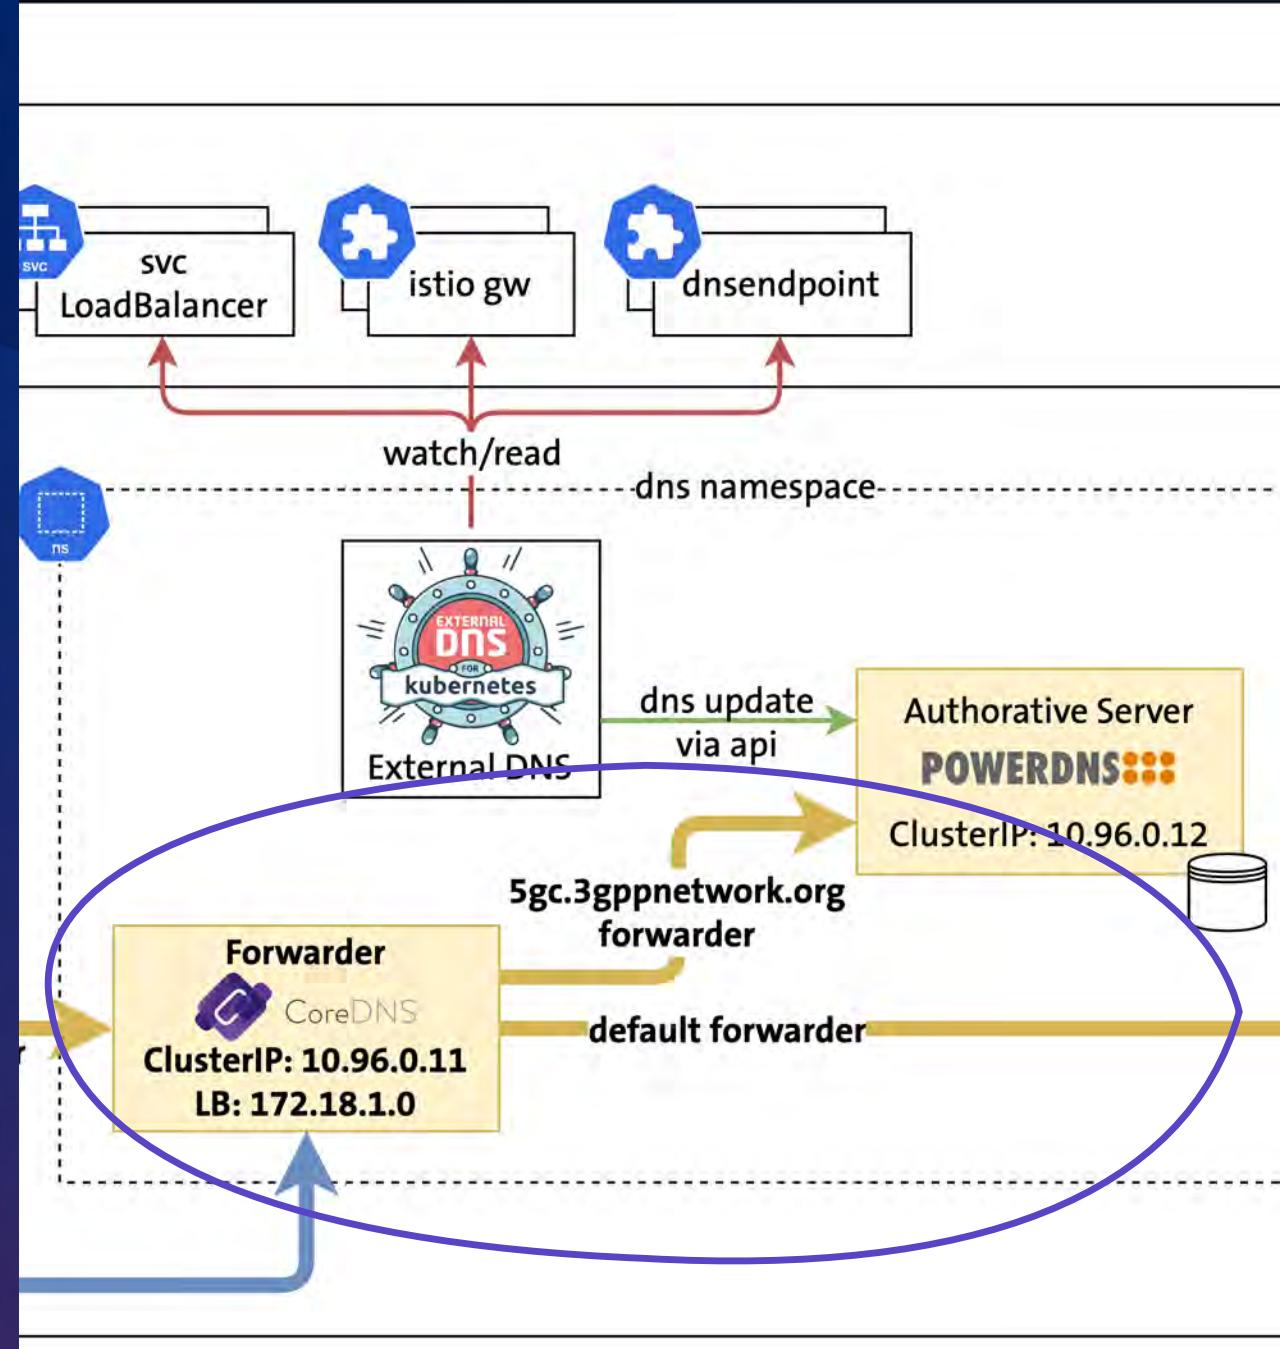




Demo Forwarding

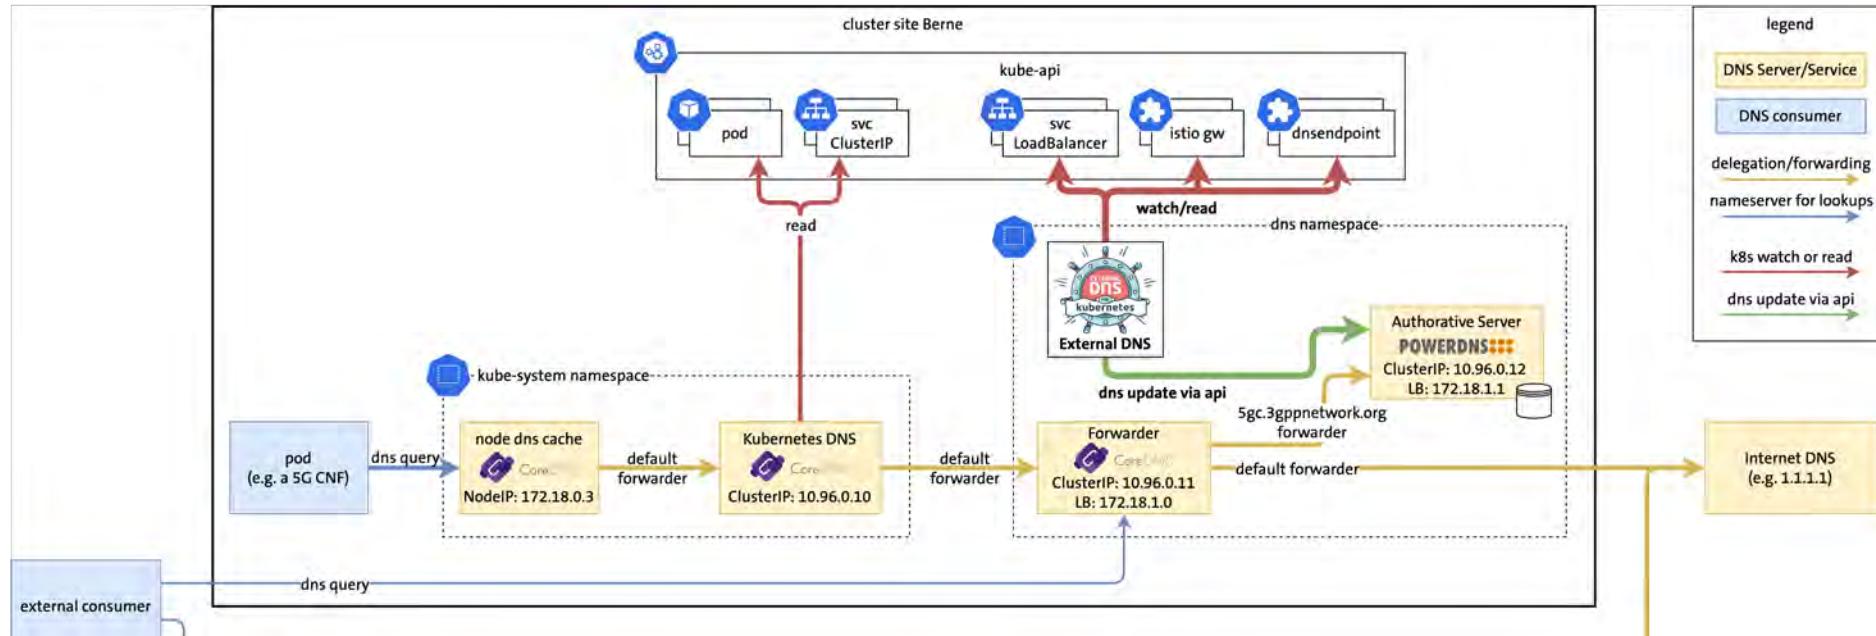


<https://github.com/swisscom/cloud-native-telco/tree/main/prototypes/dns/2-demo-forwarding>



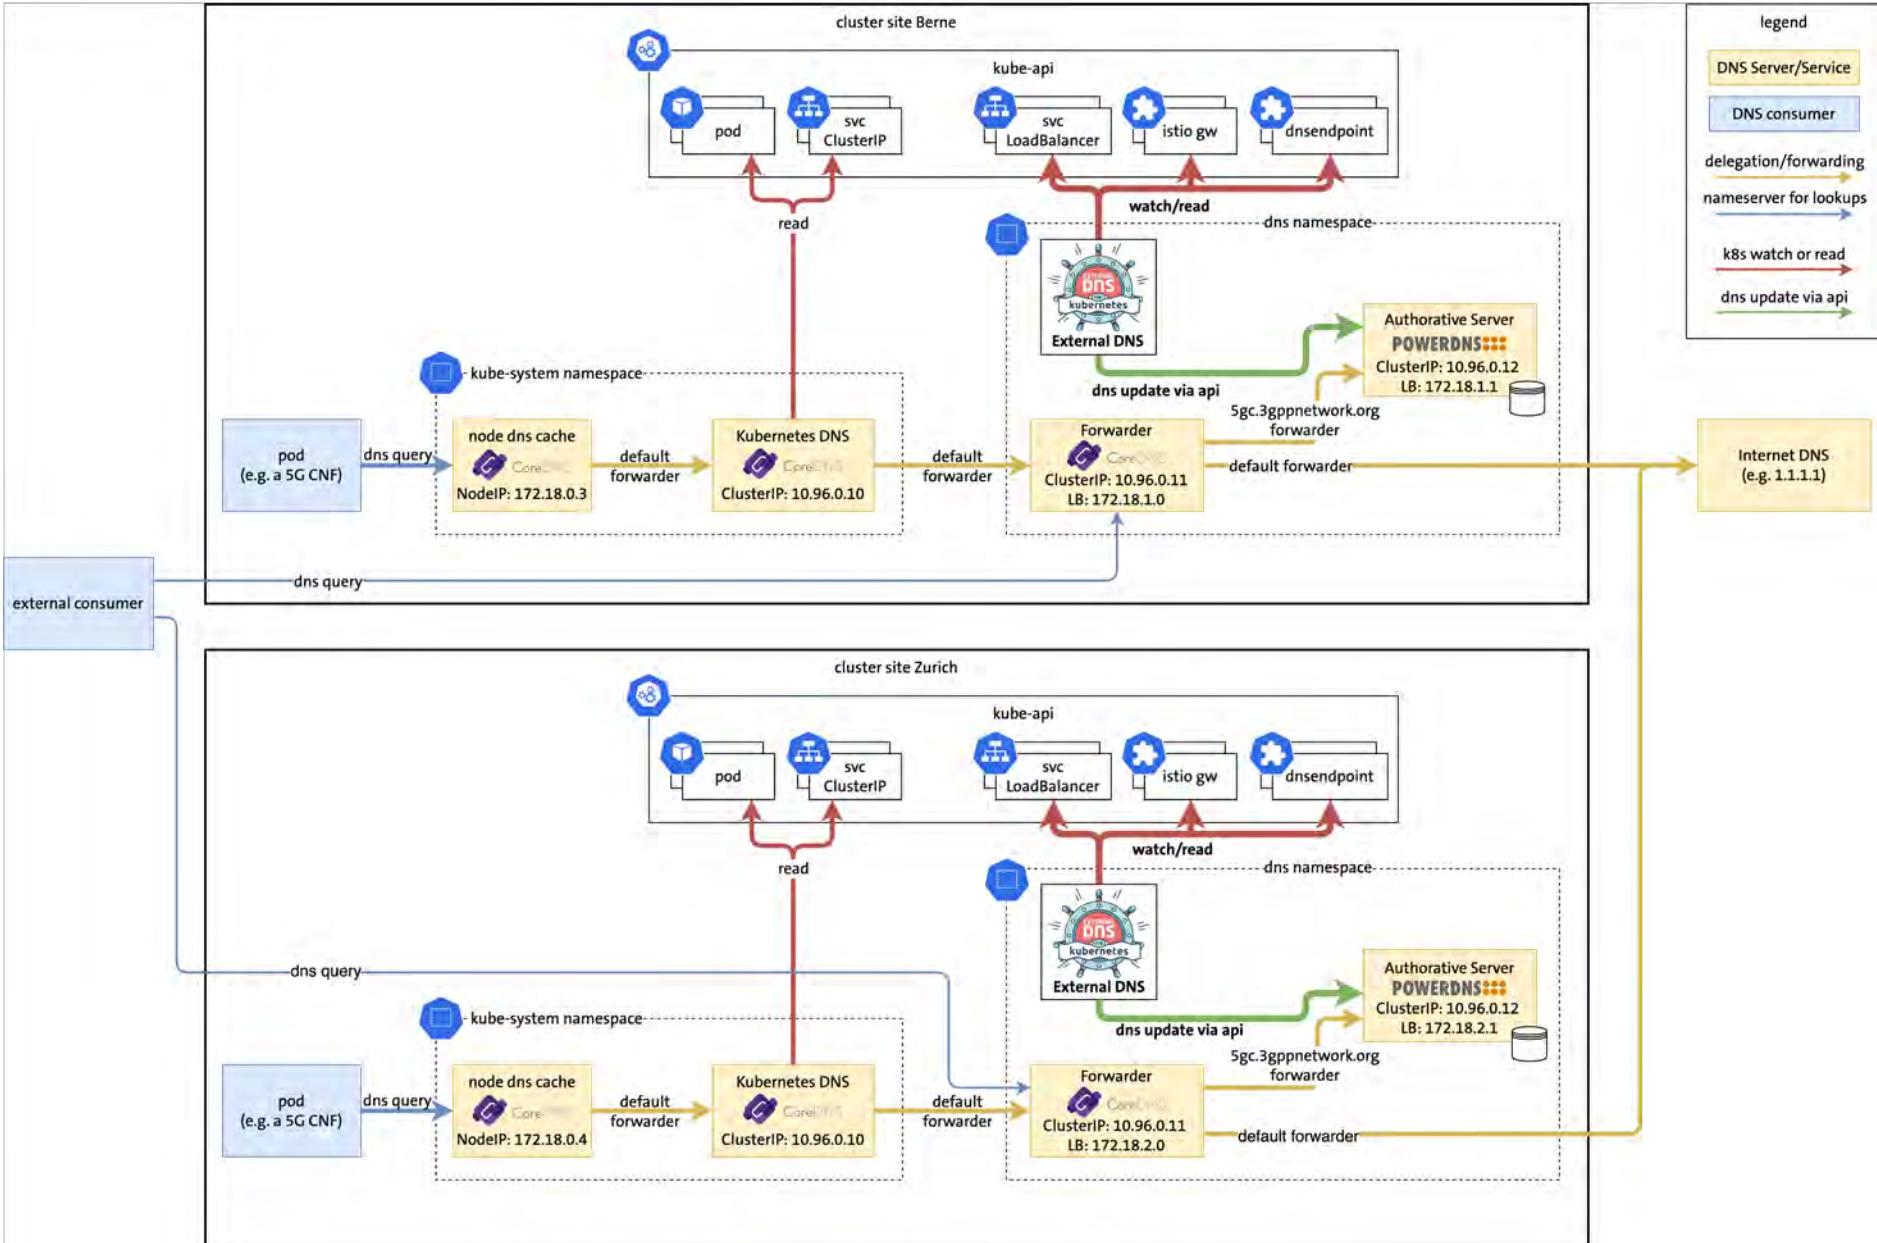


Dual-Cluster Using ExternalDNS



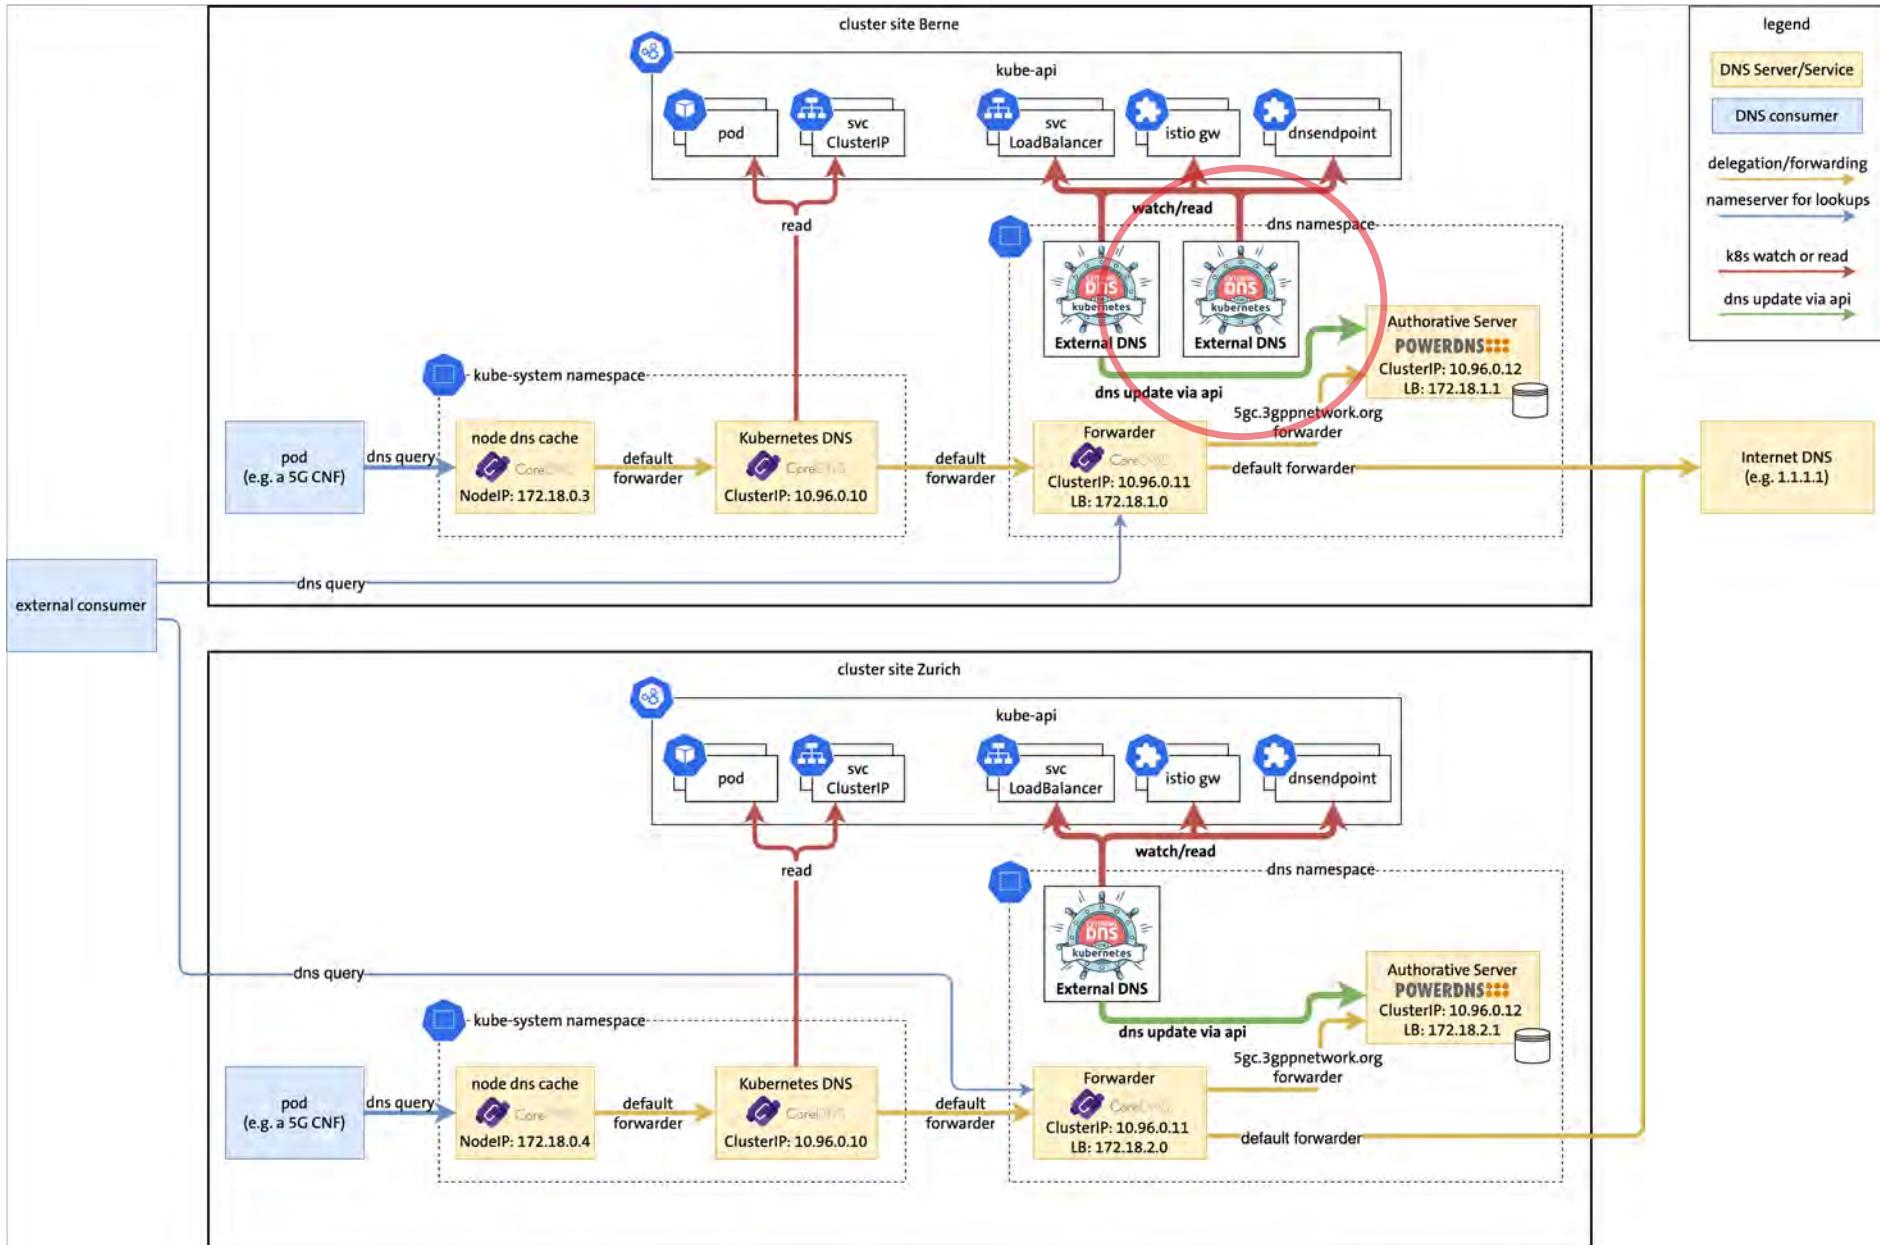


Dual-Cluster Using ExternalDNS



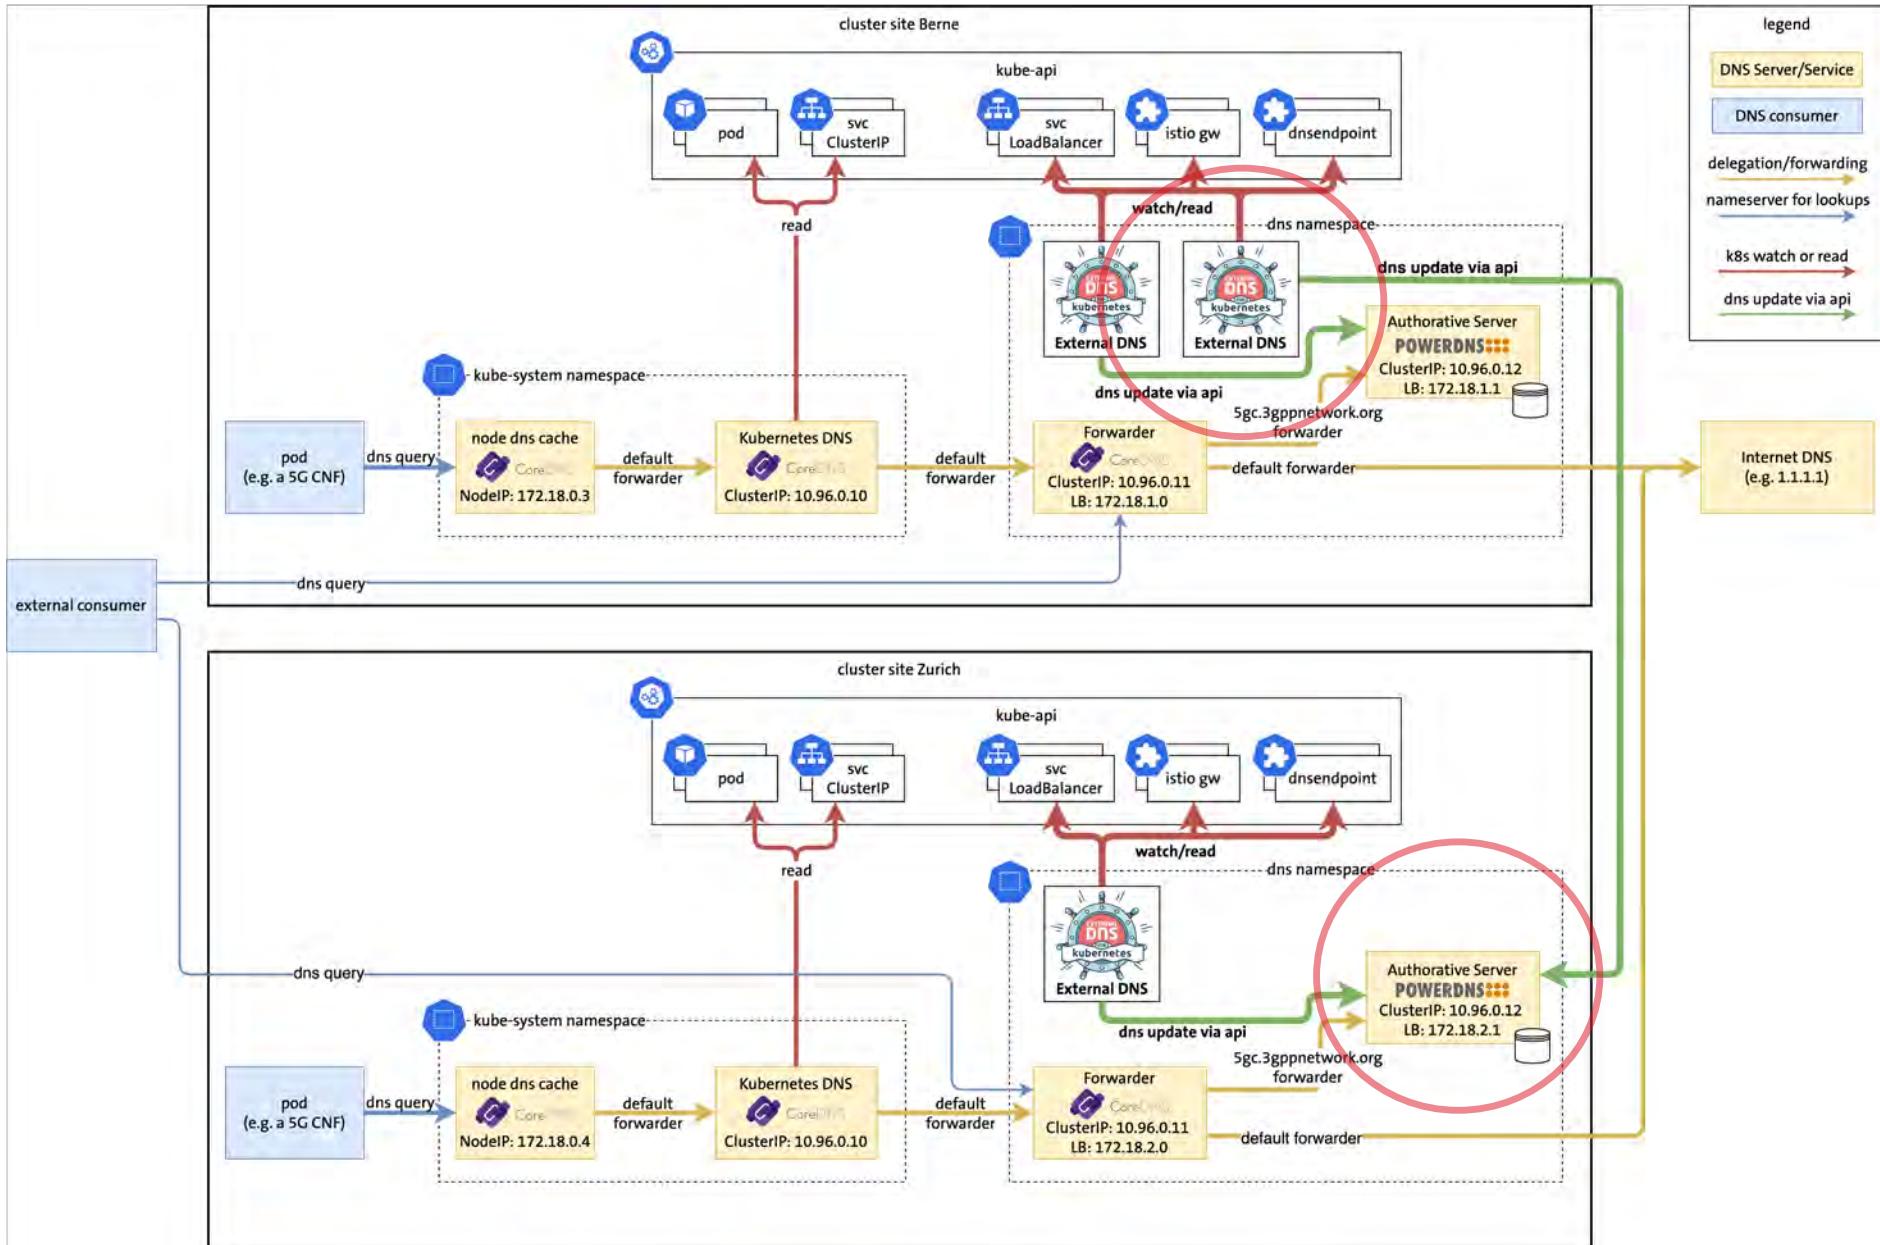


Dual-Cluster Using ExternalDNS



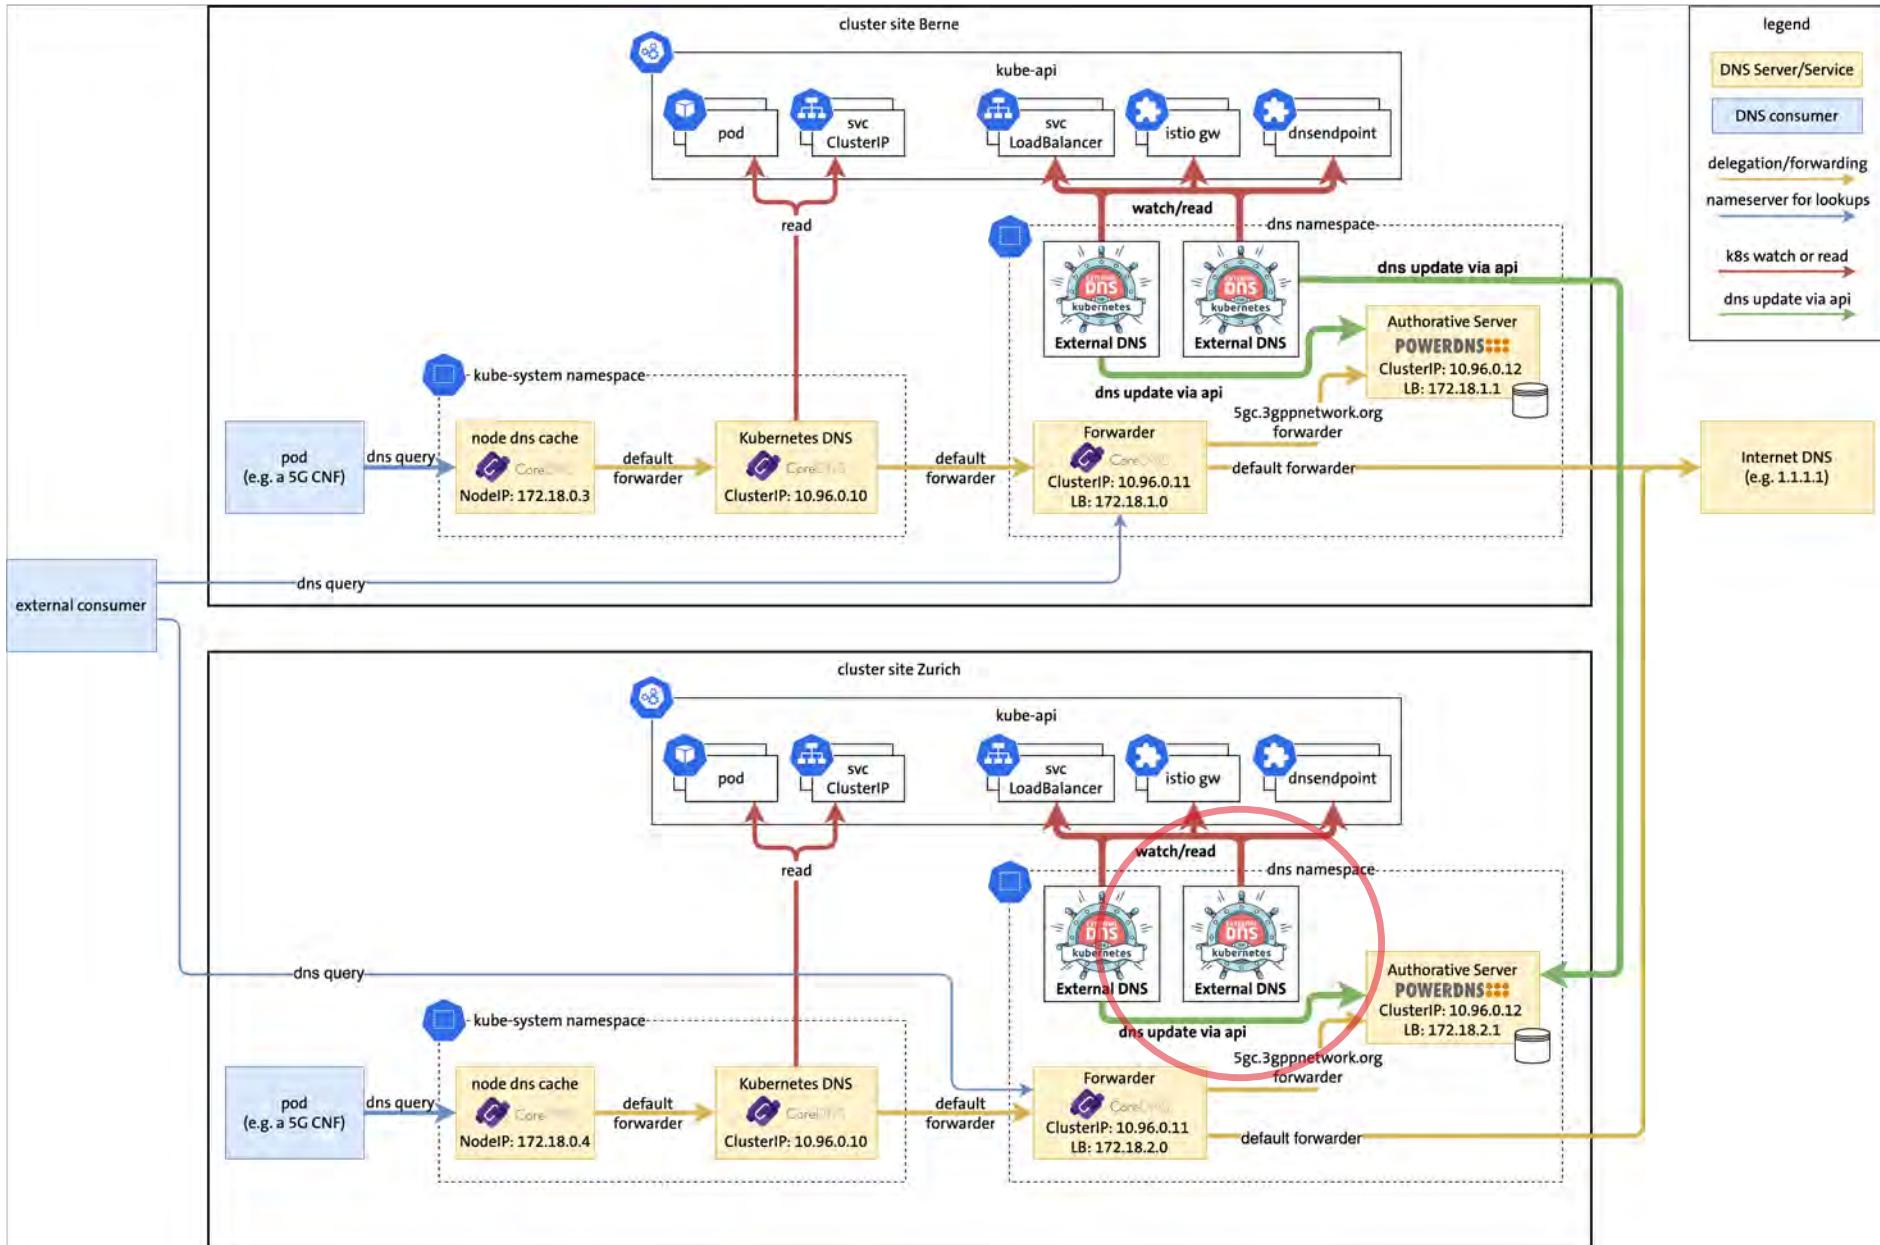


Dual-Cluster Using ExternalDNS



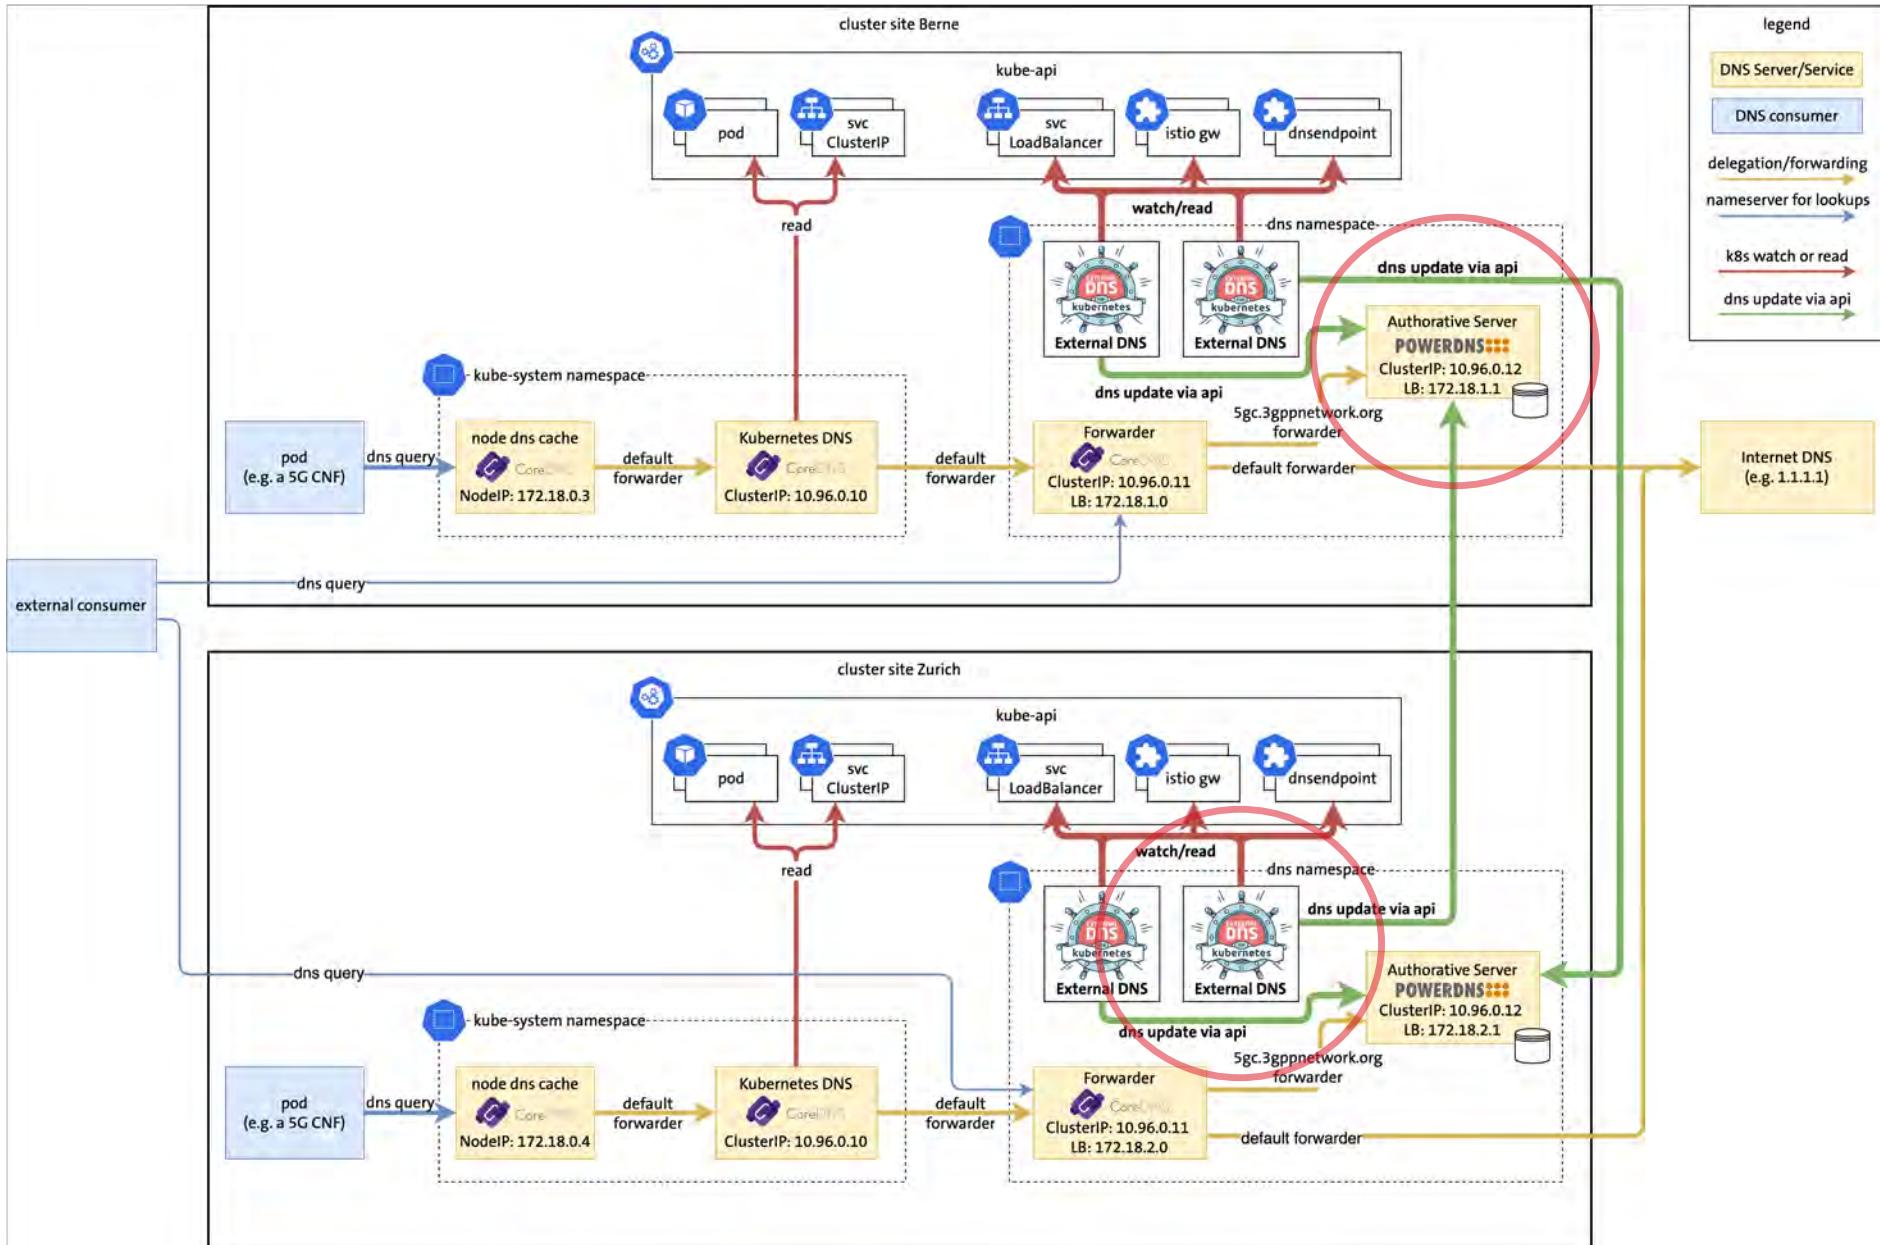


Dual-Cluster Using ExternalDNS



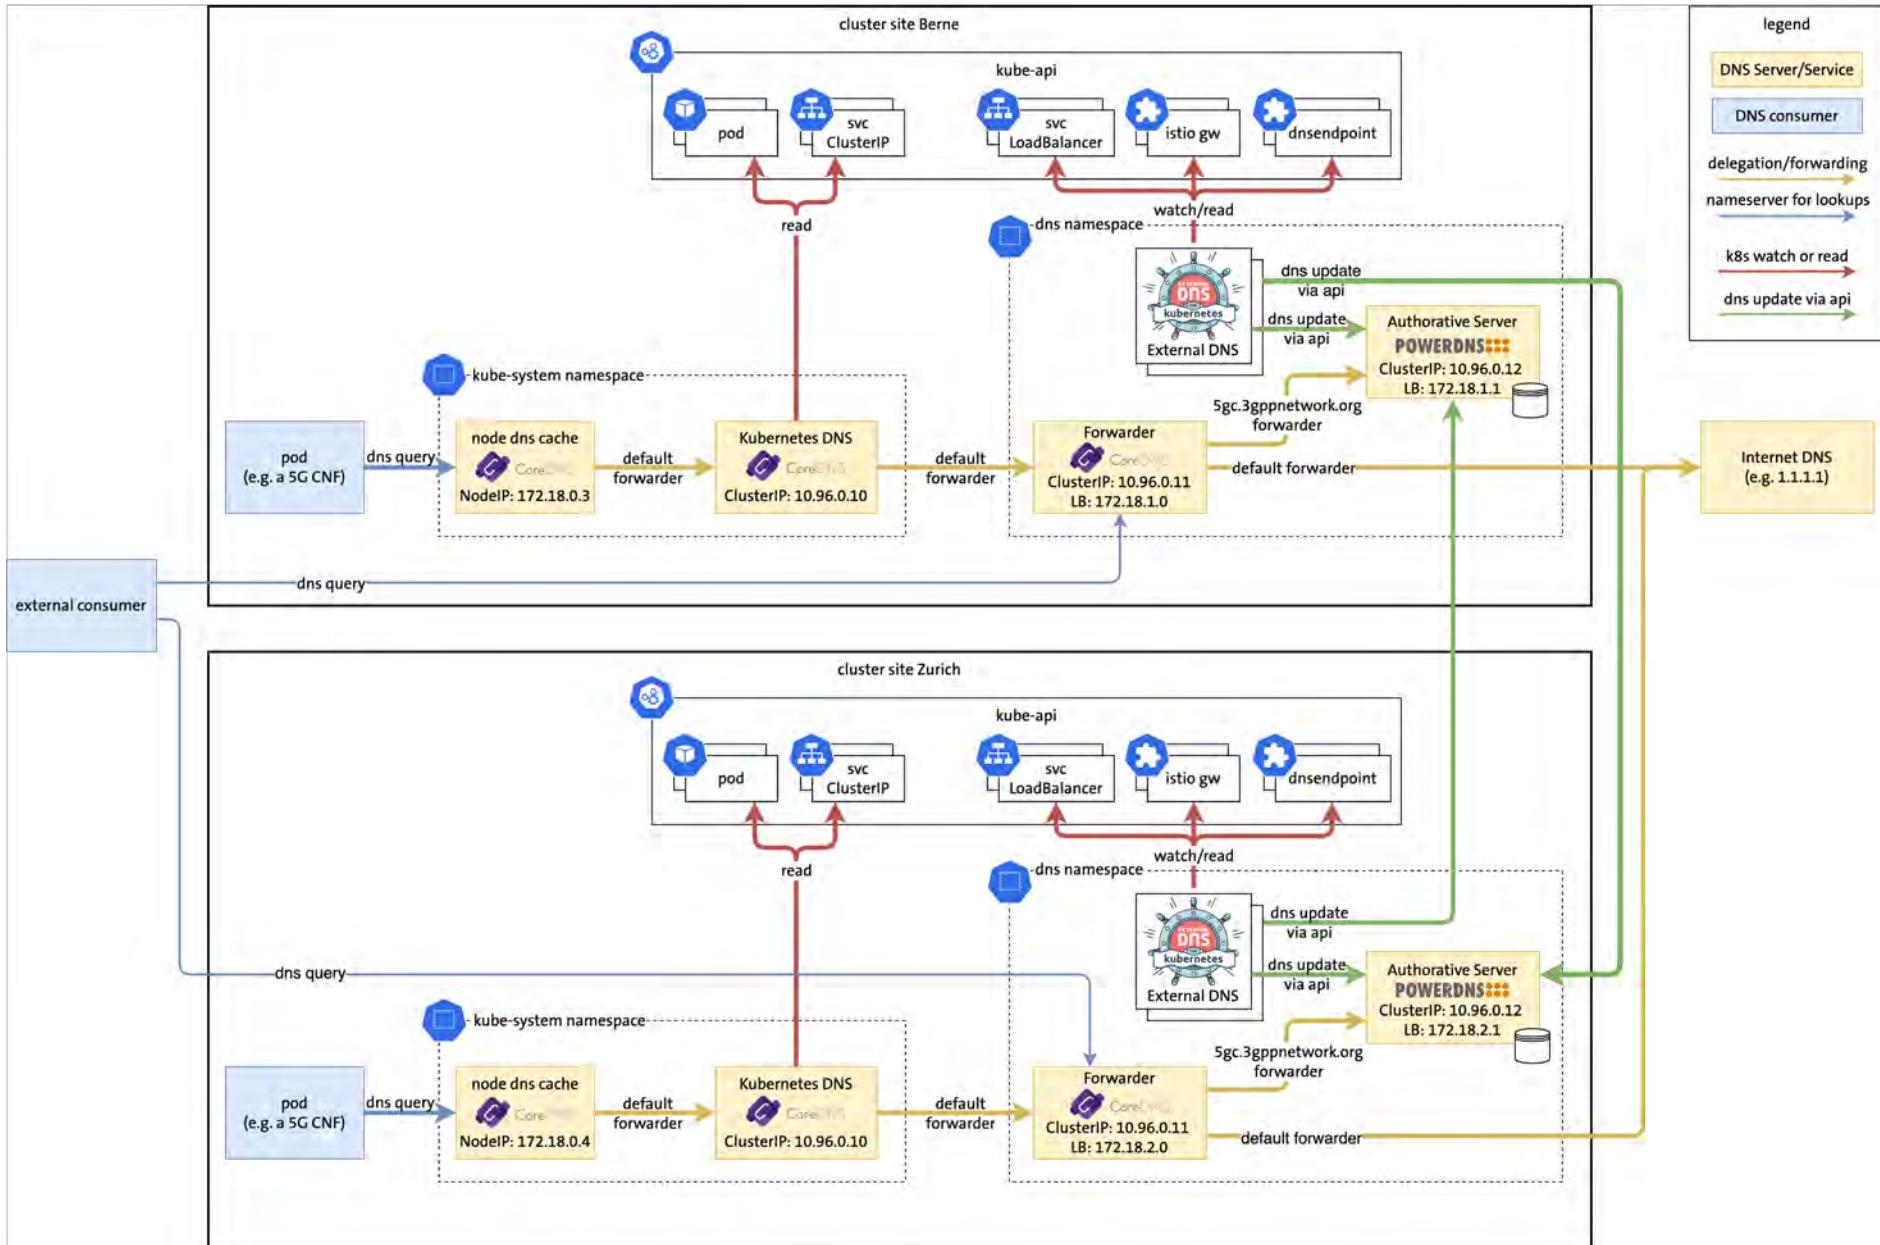


Dual-Cluster Using ExternalDNS



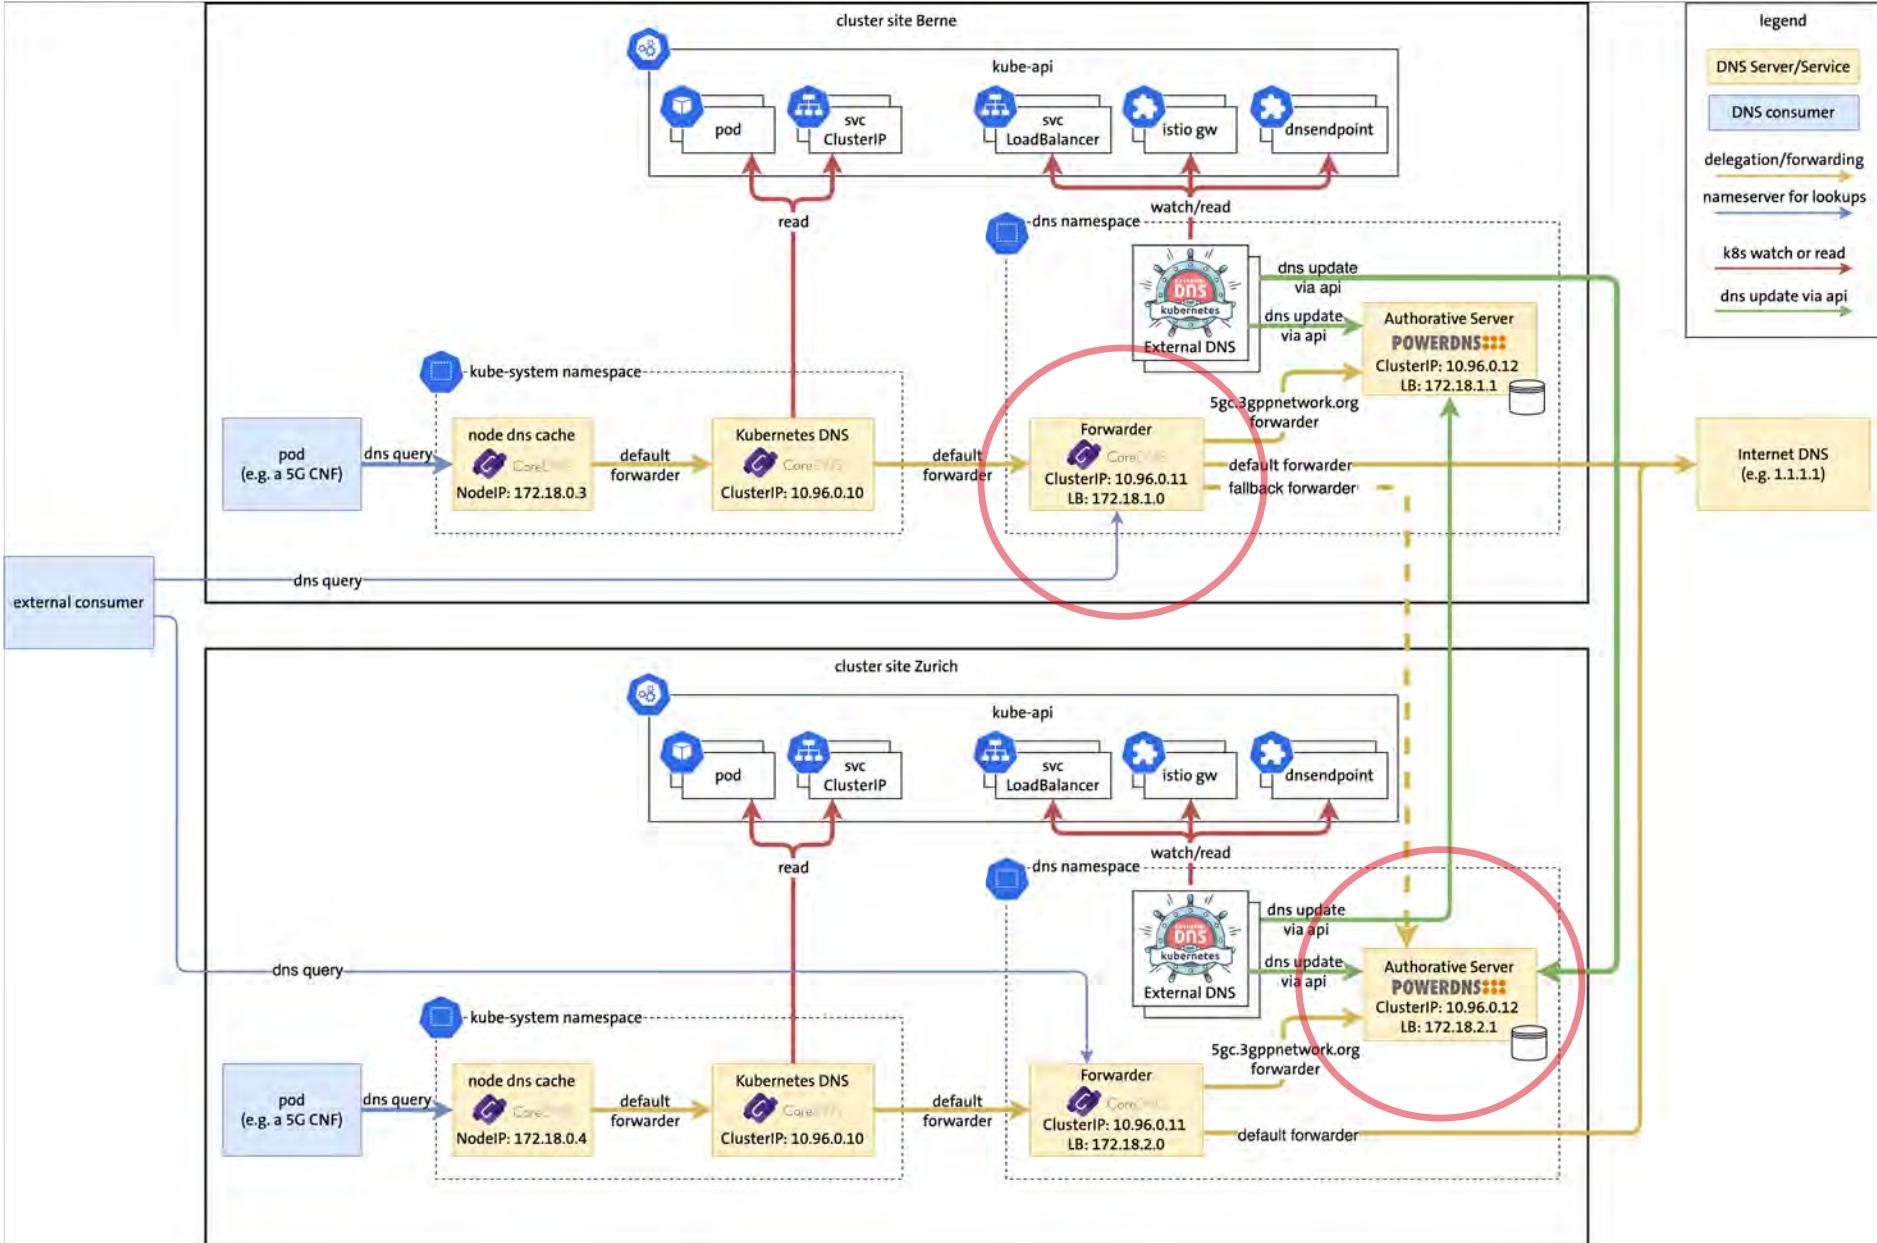


Dual-Cluster Using ExternalDNS



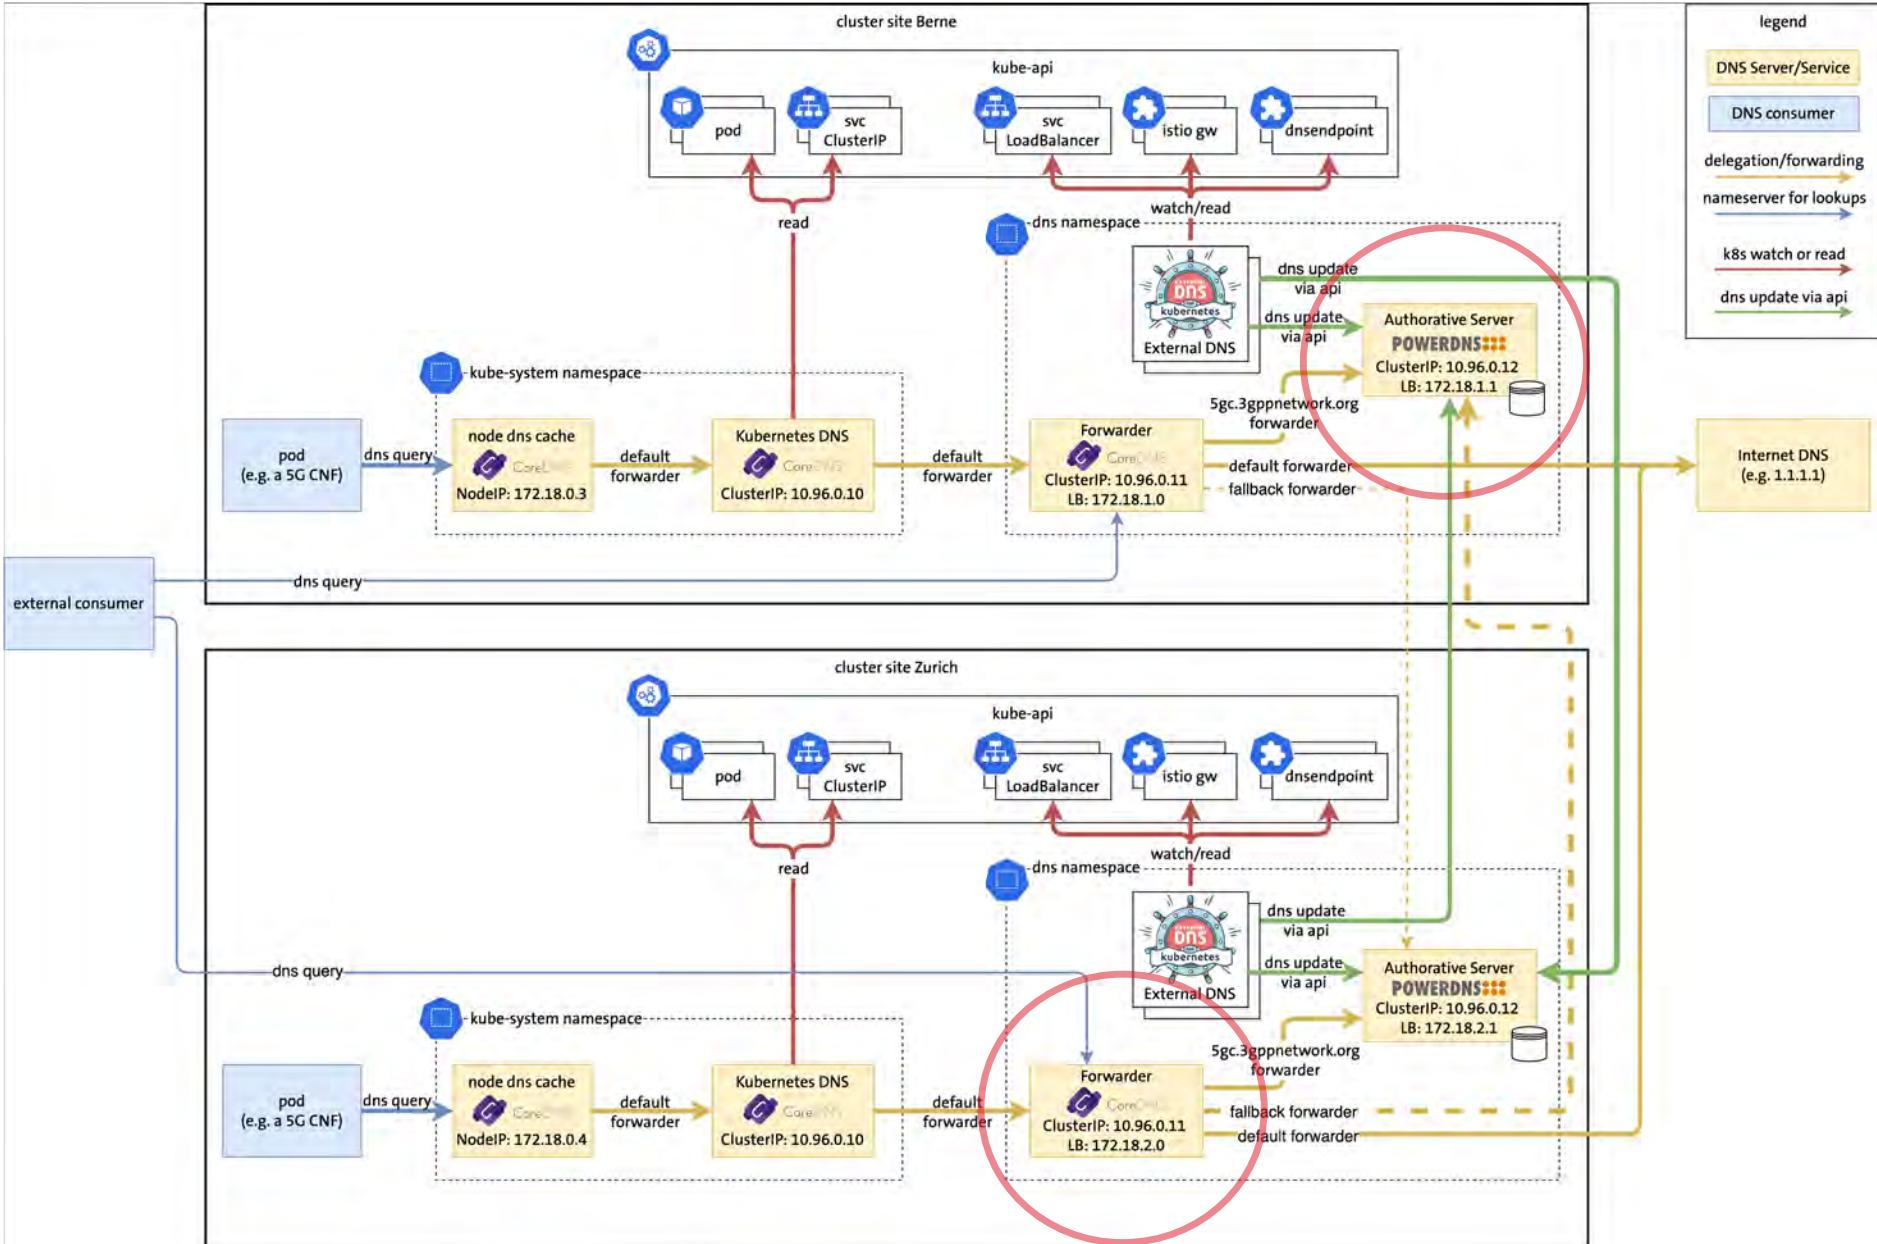


Eliminating Single Point of Failure



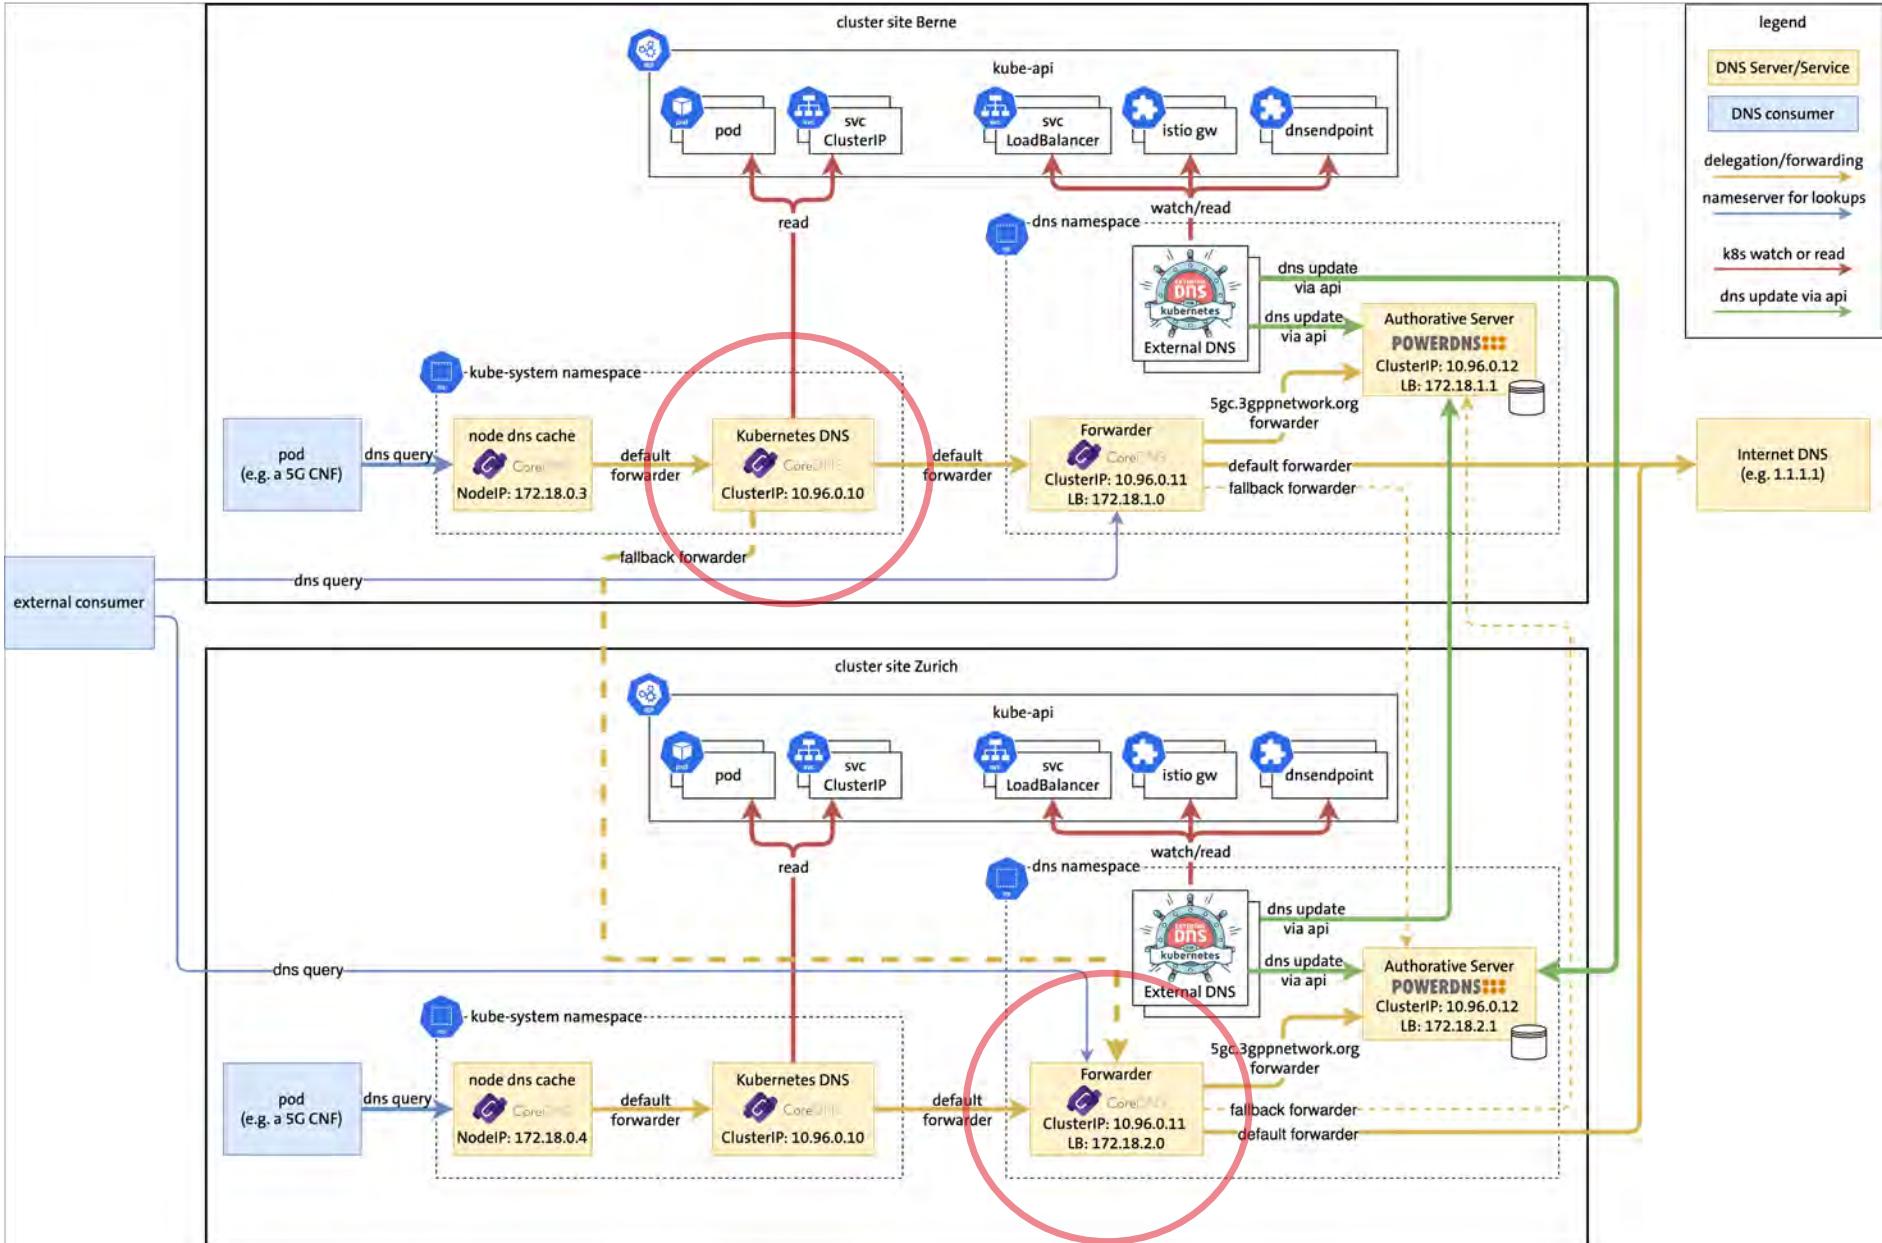


Eliminating Single Point of Failure



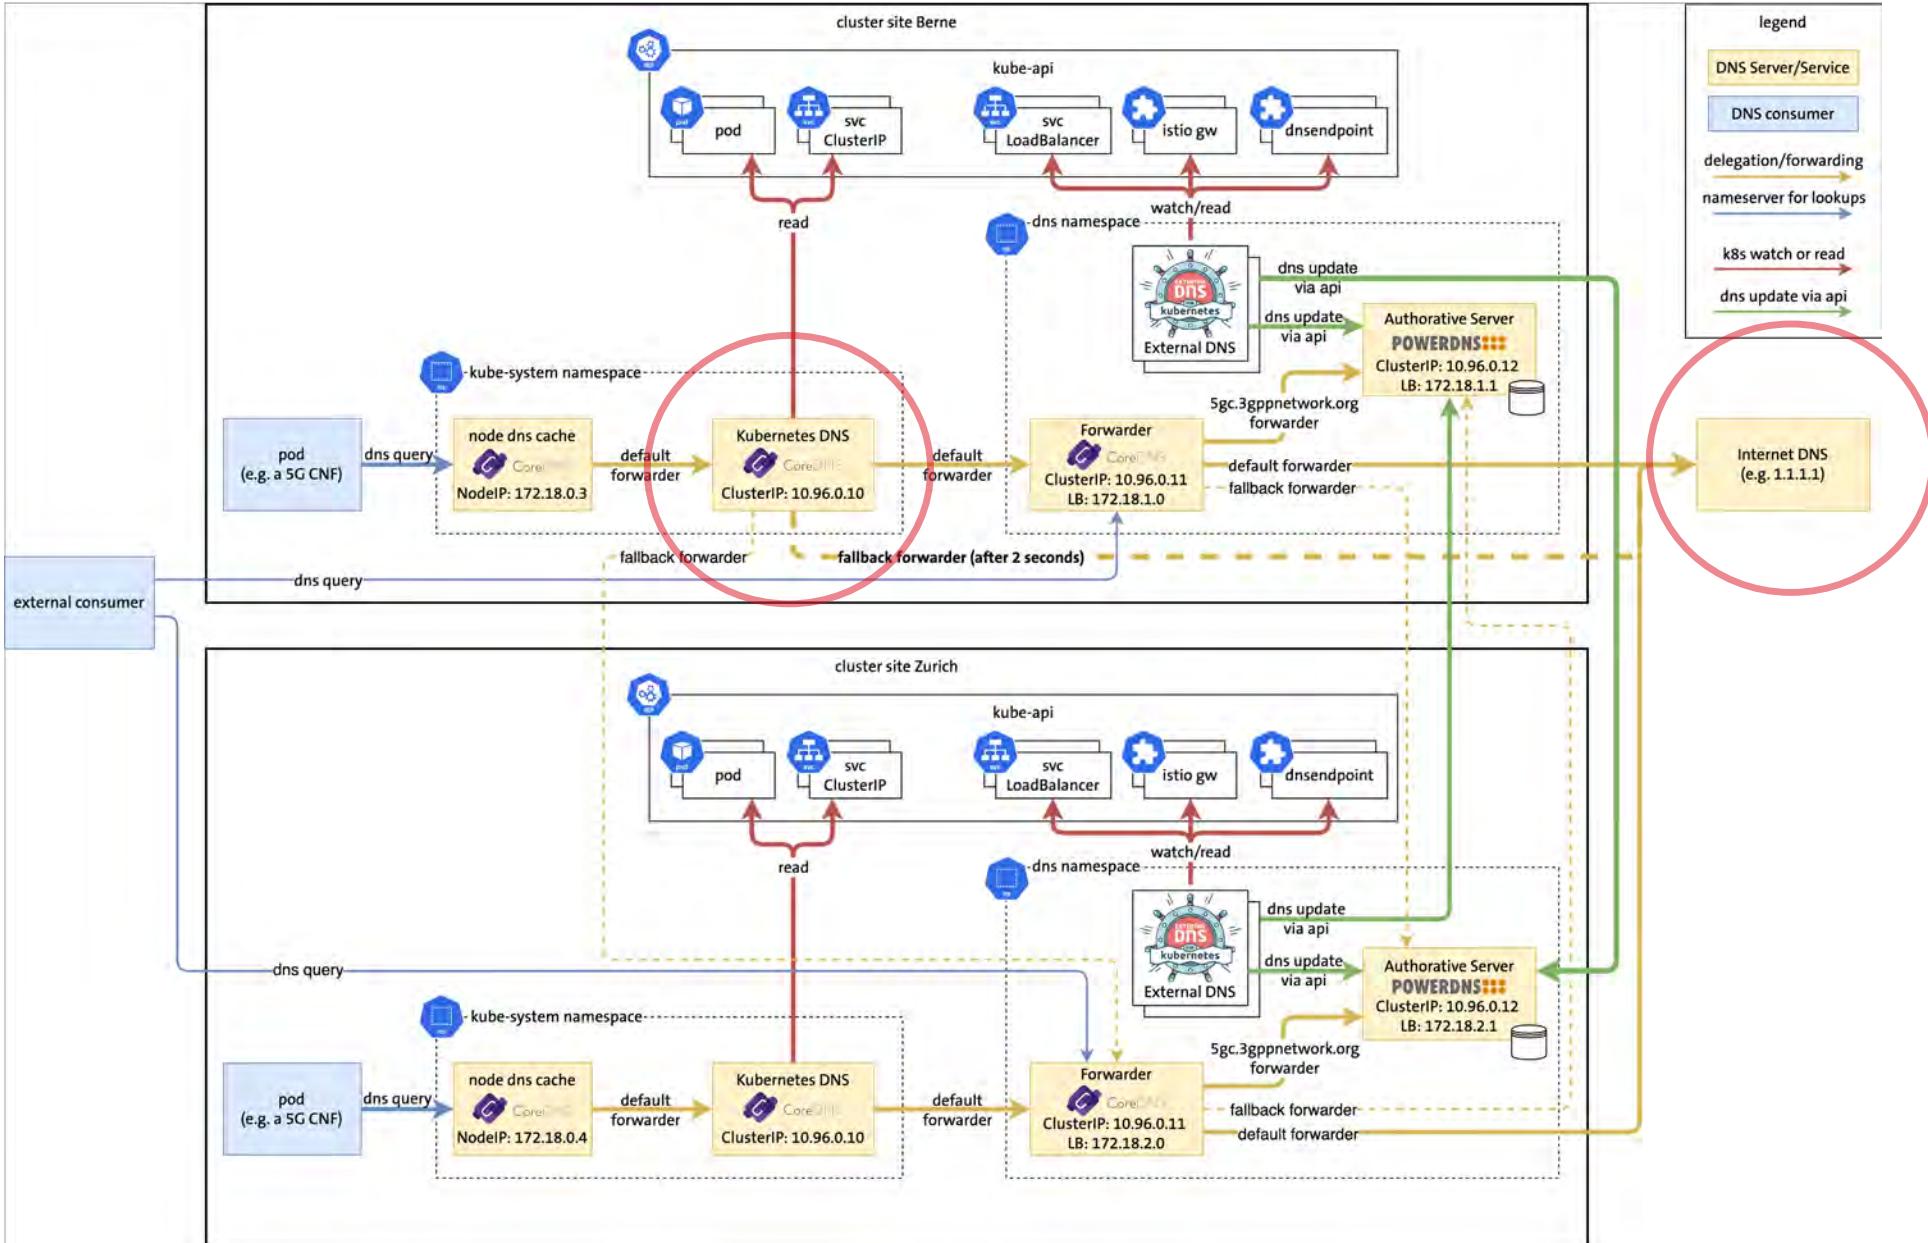


Eliminating Single Point of Failure



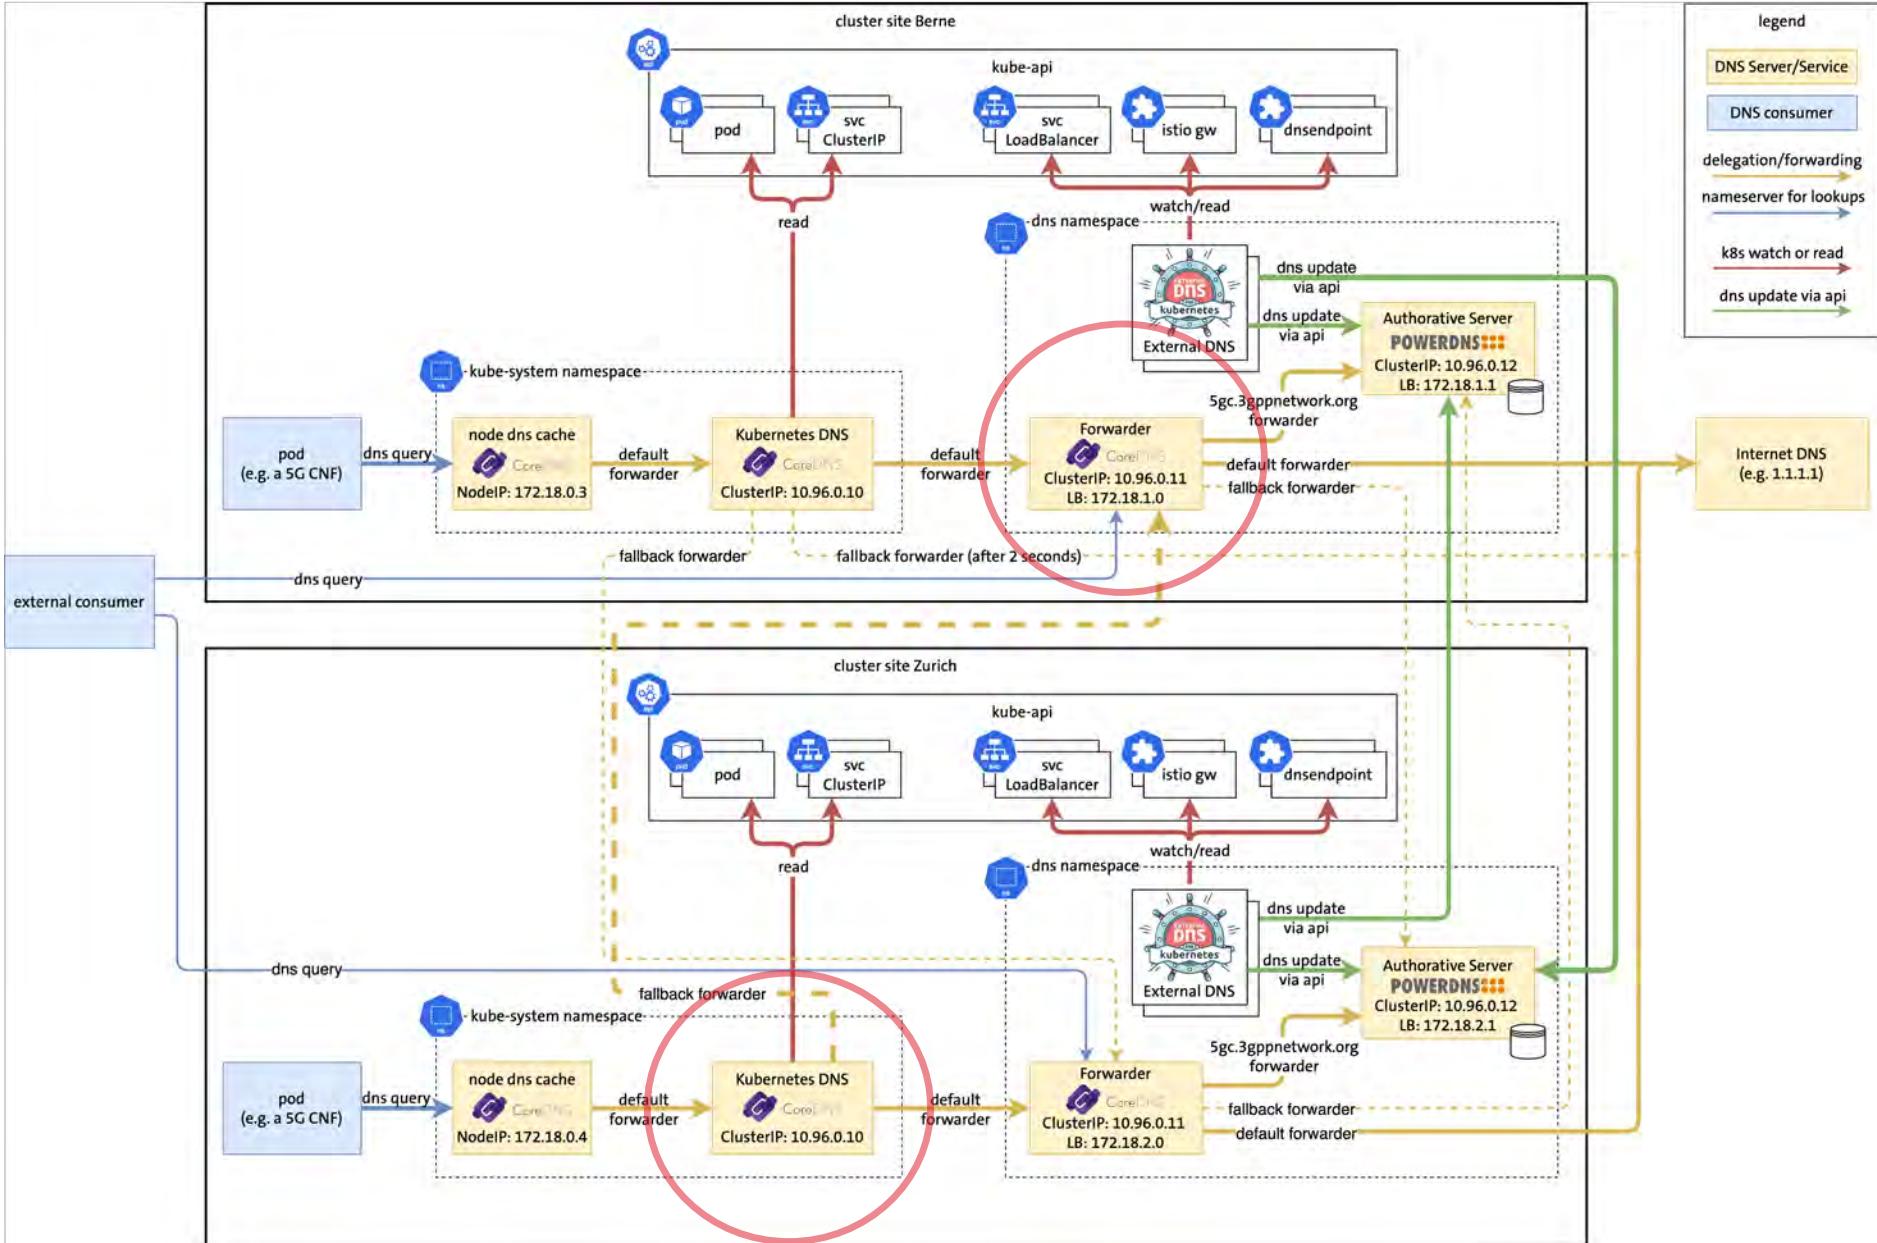


Eliminating Single Point of Failure



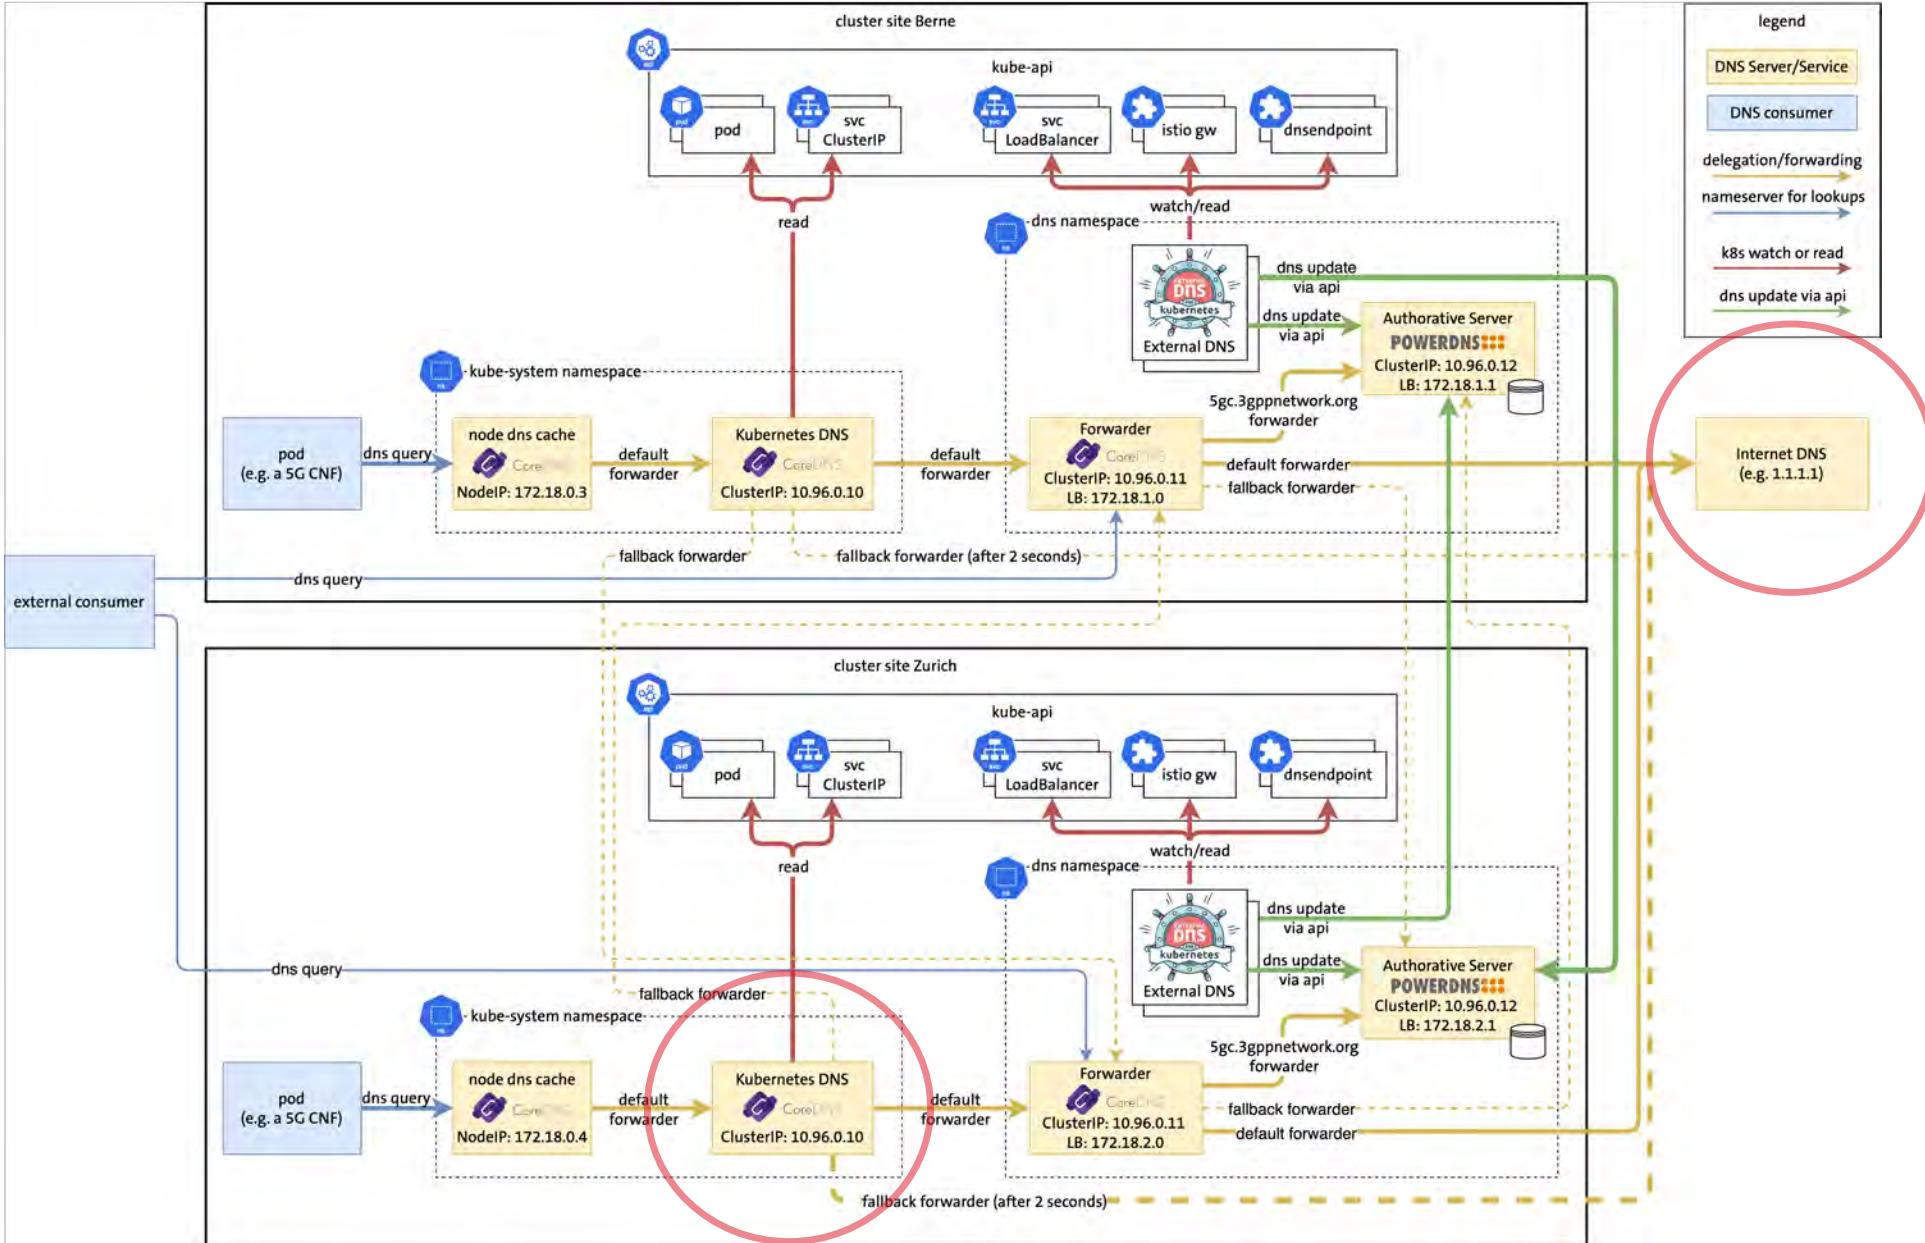


Eliminating Single Point of Failure



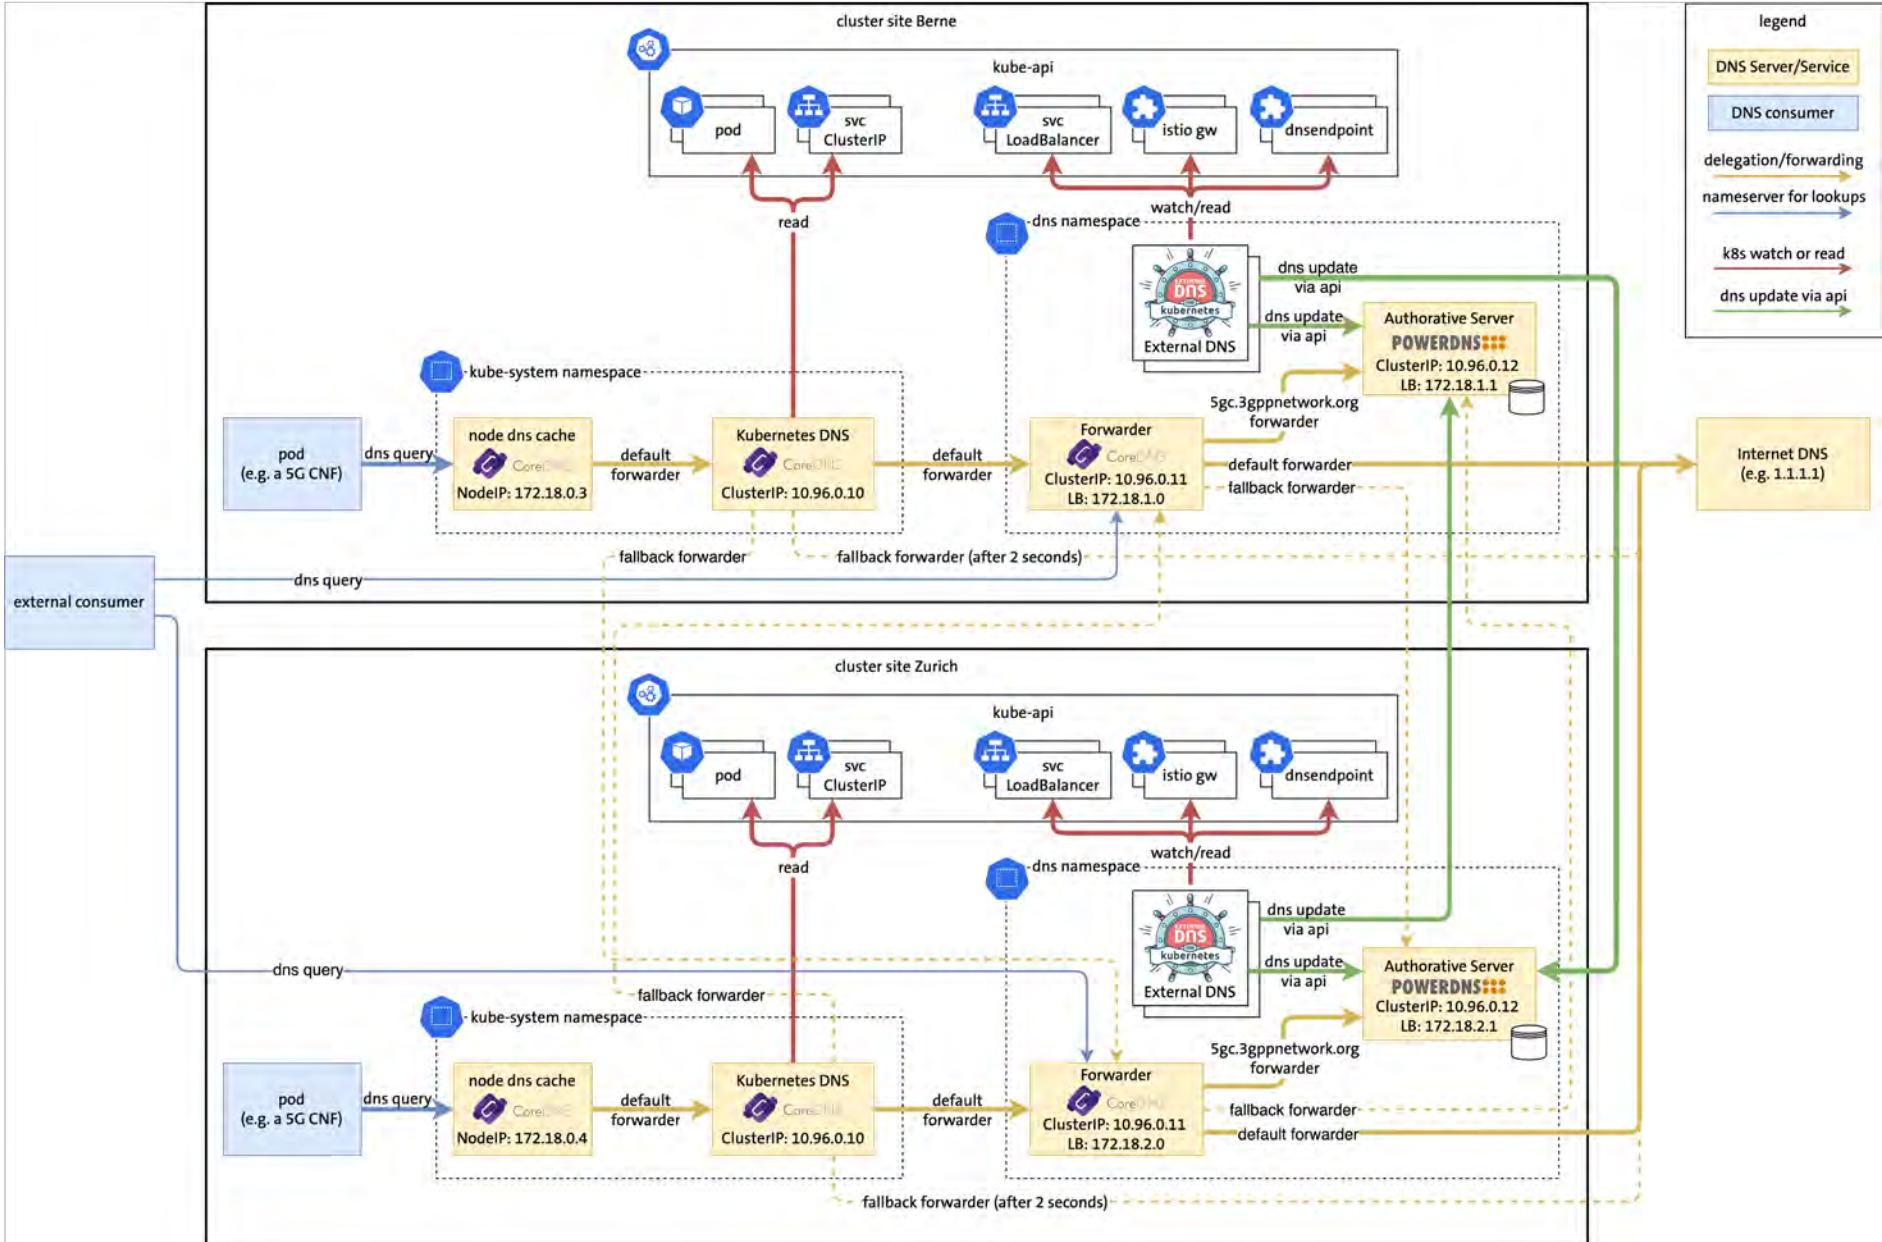


Eliminating Single Point of Failure





Final Dual Cluster Setup

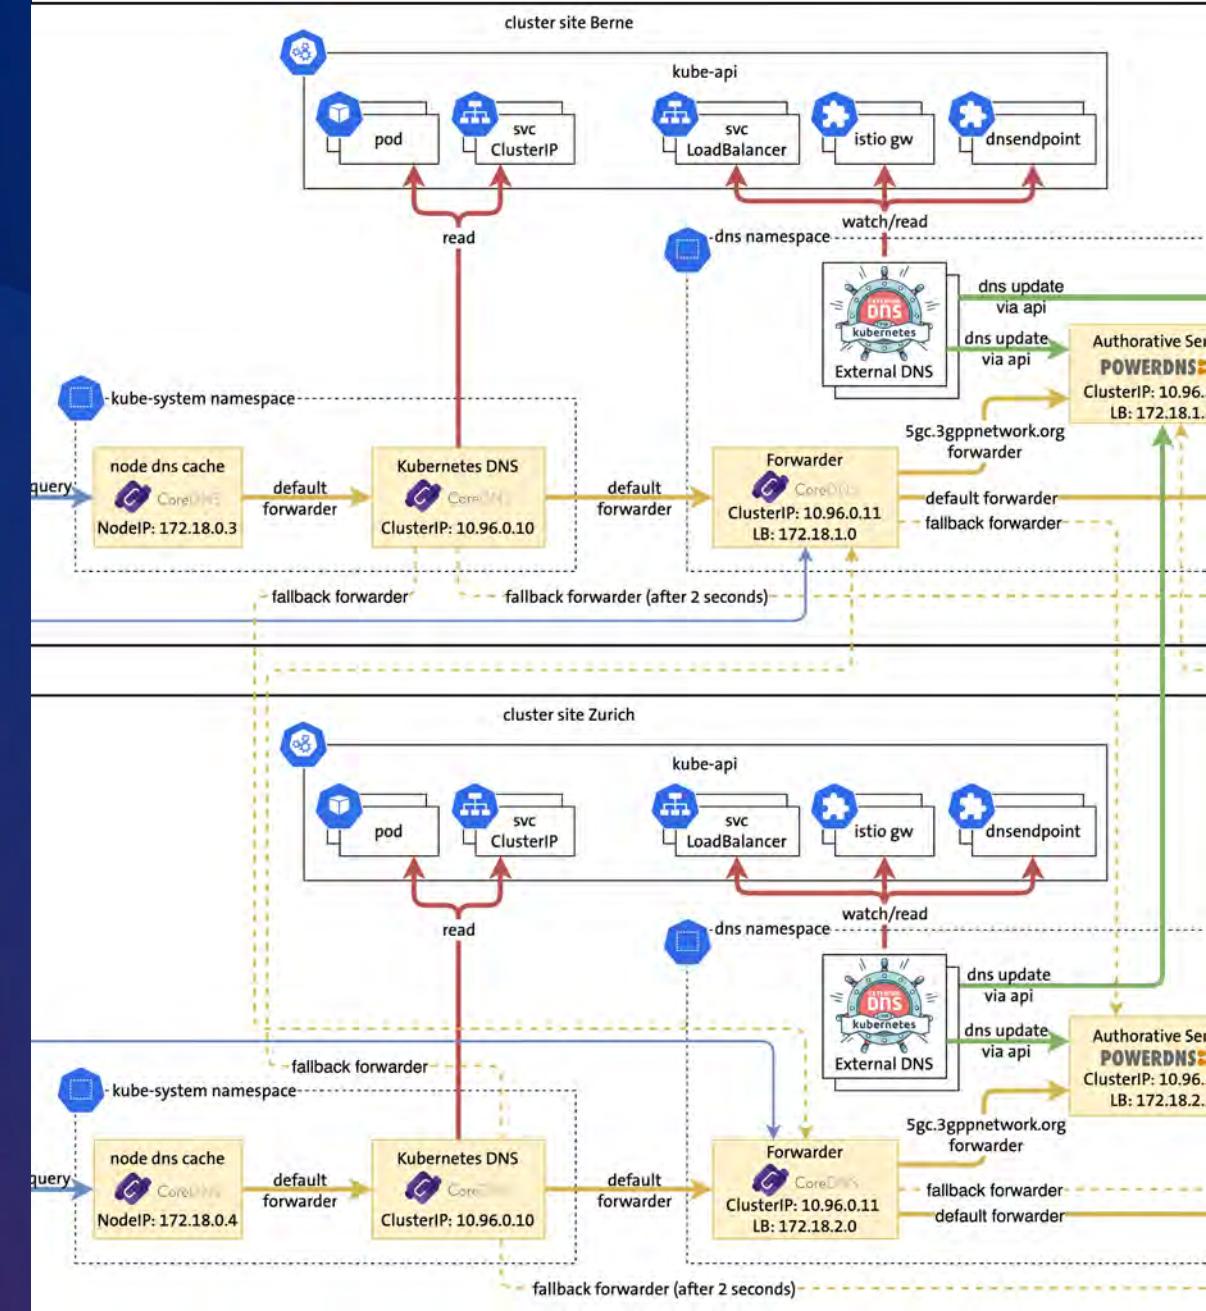




Demo Multi Cluster



<https://github.com/swisscom/cloud-native-telco/tree/main/prototypes/dns/3-demo-multi-cluster-dns>





Limitations of Our DNS Service



Self-dependence

Complexity increases when consuming the Service from within the same clusters.



Kubernetes Resources only

Limited to Kubernetes Resources and GitOps



Service Discovery

ExternalDNS not suited for service discovery



Limitations of ExternalDNS: Service Discovery

Interval-Based Syncing due to architectural decisions

- ⚠️ Delayed Resource Record creation

No Health Checks (e.g. integration into [Kubernetes Services/EndpointSlices](#))

- ⚠️ Cannot rely on ExternalDNS for app readiness

No Multi Cluster Round Robin for A records: one record cannot be shared by multiple ExternalDNS

- ⚠️ Cannot use DNS records created by ExternalDNS for routing across multiple clusters

Full cluster outage will not revoke DNS records

- ⚠️ Tight monitoring and additional automation needed to avoid outages



What Did We Achieve?



Proximity to Consumer

Minimal amount of hops between 5G Core and DNS

✓ On-prem deployment



Fully Automated

GitOps driven and automated provisioning of DNS records

✓ GitOps + ExternalDNS



Geo Redundant & HA

Spread across multiple K8s clusters and geo regions to increase reliability

✓ Spread across multiple K8s Clusters



Support of Advanced DNS features

Resource Records such as NAPTR and SRV supported for e.g. SIP Phone Calls

✓ Advanced RRs supported



K8s integration with ExternalDNS

The System leverages Kubernetes Patterns such as CRs and Operators

✓ 100% Kubernetes Resources



Minimal Amount of SPOFs

Remove single points of failure from the System

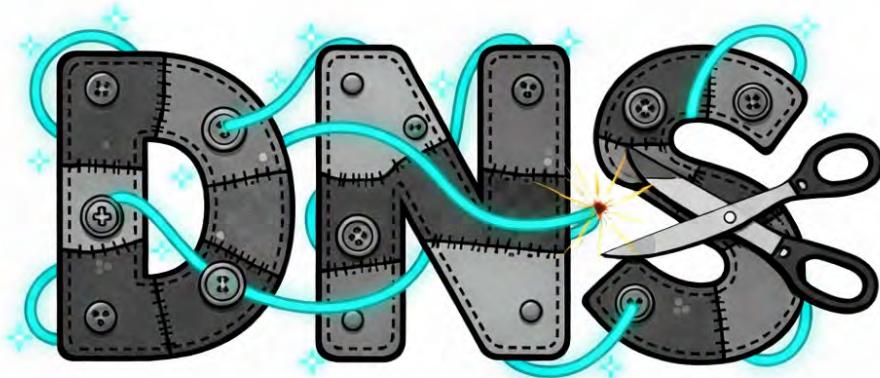
✓ Distributed control plane

✓ Cross-Cluster Forwarding



Thanks!





RESILIENCE TESTING

Resilience Testing on Cloud Native DNS
Georgios Daskalopoulos

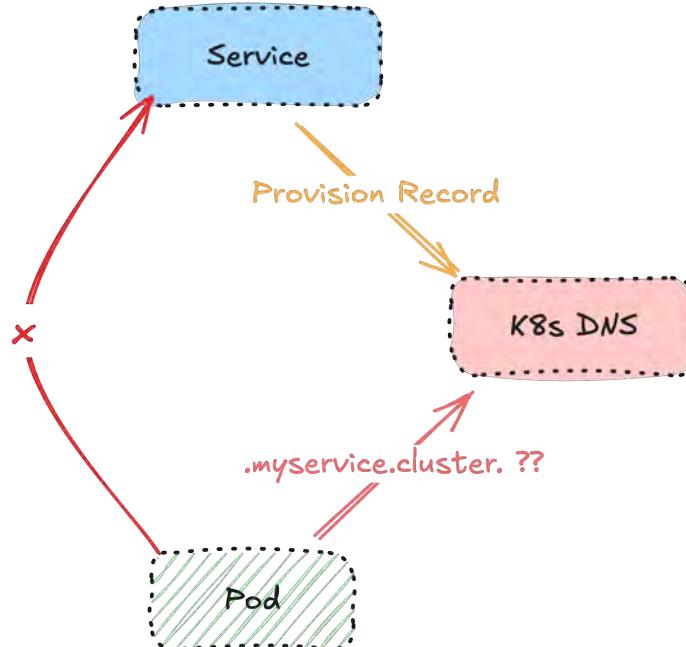
Swisscom

- Fabian Schulz
- Joel Studler

Communications Systems Lab @ UZH

- Dr. Burkhard Stiller
- Thomas Grübl

Cloud Native DNS -- Resilience Testing

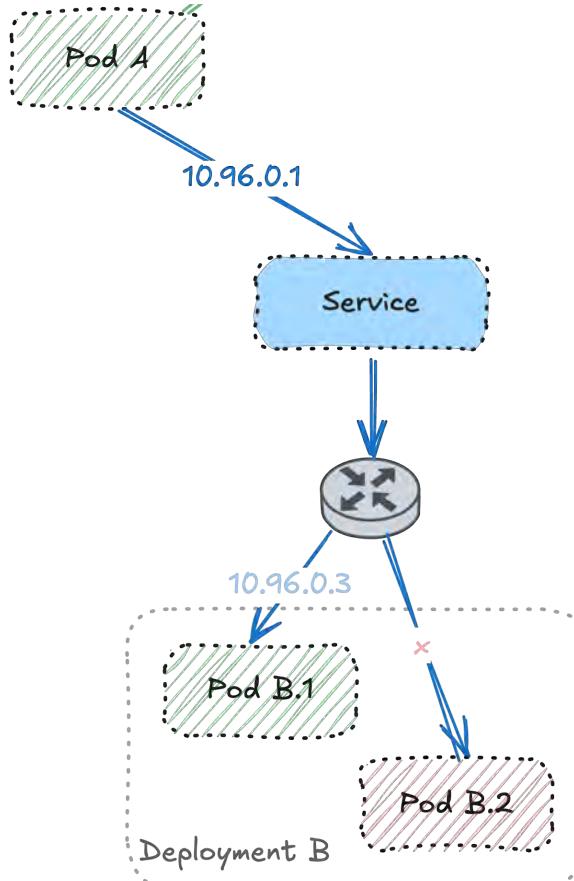


Cloud Native DNS differs from Public DNS

- DNS records are updated way more frequently
- DNS outages or slow responses have cascading effects
- Fast & Accurate DNS is crucial for cloud application reliability
- Private DNS setups are not as tested as Public DNS !



Cloud Native DNS -- Resilience Testing



- Cloud Native DNS can be tested with Resilience Testing
- Resilience
 - Ability of an application to withstand and recover from failures in its components.
 - Achieved with multiple instances of the same Process
- Resilience Testing / Chaos Engineering
 - Controlled experiments which assess a system's Resilience.
 - Define Stady State based on KPIs
 - Introduce chaos, process crashes, network disruptions
 - Measure the impact on the system behavior



Real World Incidents

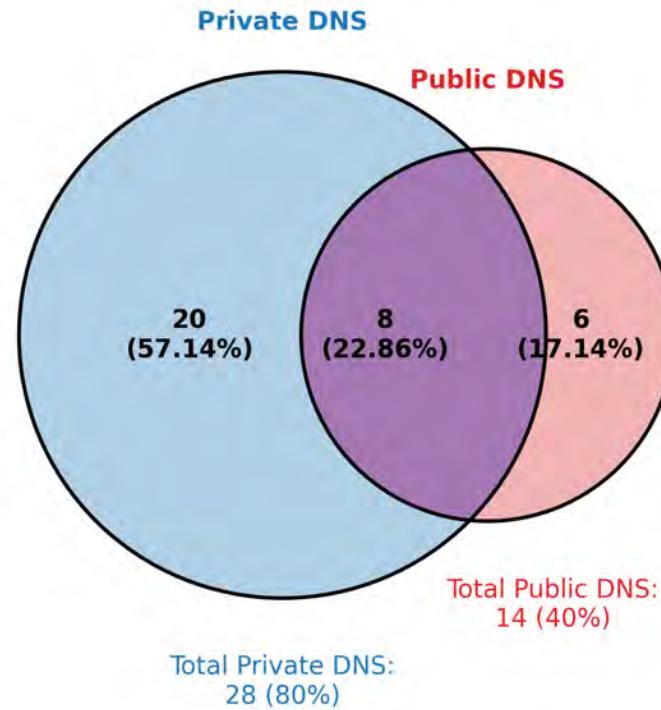
35 Post - Mortems analysed from cloud providers

Source	Incidents
Google Cloud	19
AWS	7
Azure	4
Post-Mortem List	5
Total	35

- . Affected DNS Infrastructure
- . DNS being the Root Cause
- . Name Resolution
- . Record Provisioning affected
- . Public or Private Zone affected



Real World Incidents



Incident Impact on DNS Zones

80% of incidents related to private DNS zones

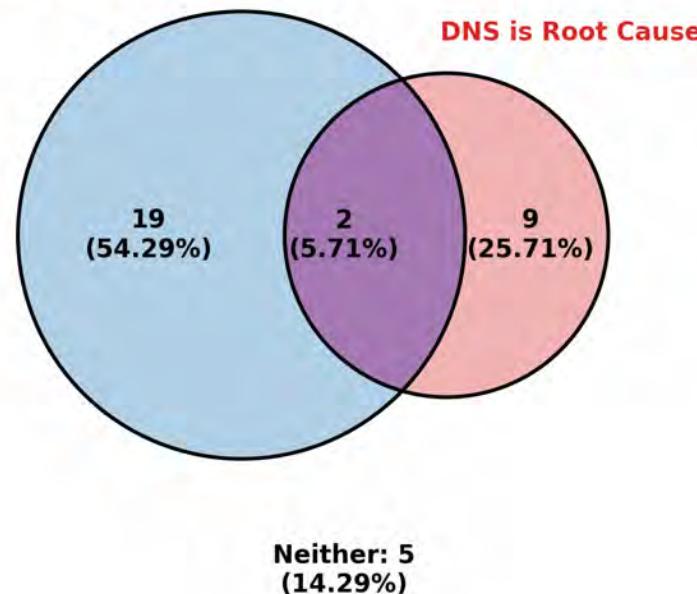
57% of incidents related **only** to private DNS zones

Private DNS is involved in the majority of incidents



Real World Incidents

Infrastructure affected



DNS Infrastructure Impact & DNS Root Cause Classification

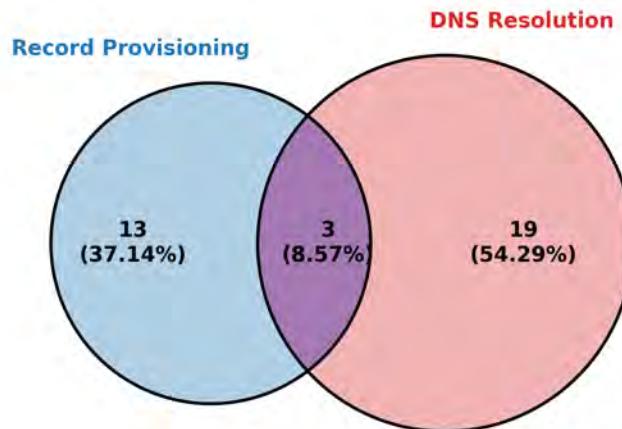
54% of the incidents, DNS is impacted by infrastructure problems

25% of the incidents DNS is the root cause

DNS should withstand failure in its underlying components



Real World Incidents



Incident Impact on DNS Functionality

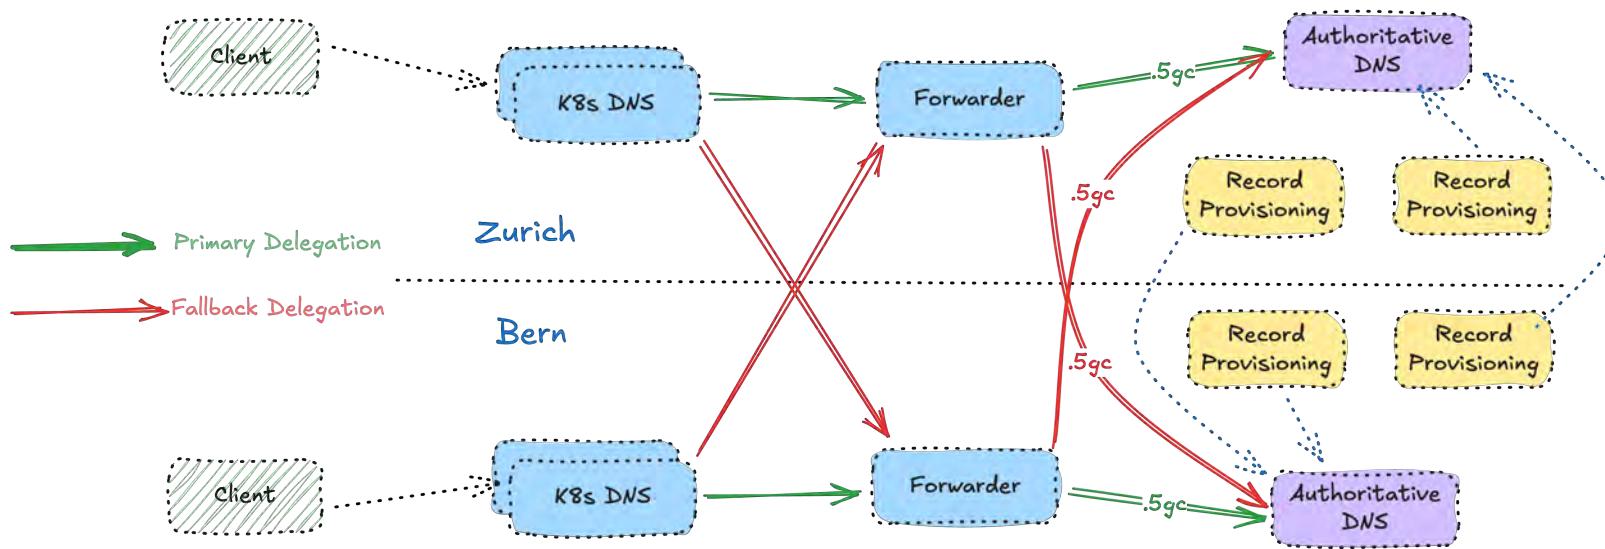
54% of incidents, DNS Resolution is affected

37% of incidents, Record Provisioning is affected

No function is immune to disruption and need to be tested.

Resilience Testing on Cloud Native DNS

How → DNS Setup -- Testing Strategy – Metrics

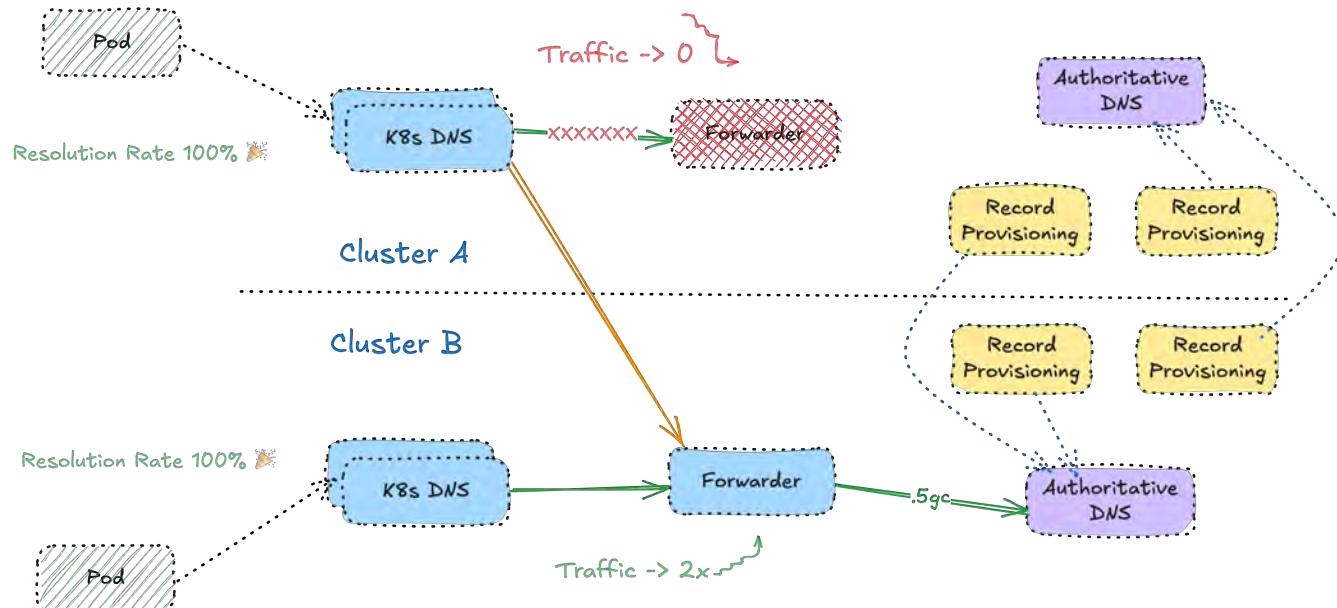


- Swisscom DNS Setup
- Multi-cluster DNS Resolution
- 5G Core Zone
- Primary & Fallback delegation
- [Github Project](#)



Resilience Testing on Cloud Native DNS

How → DNS Setup -- Testing Strategy -- Metrics



Chaos Injection, Pod Disruption:

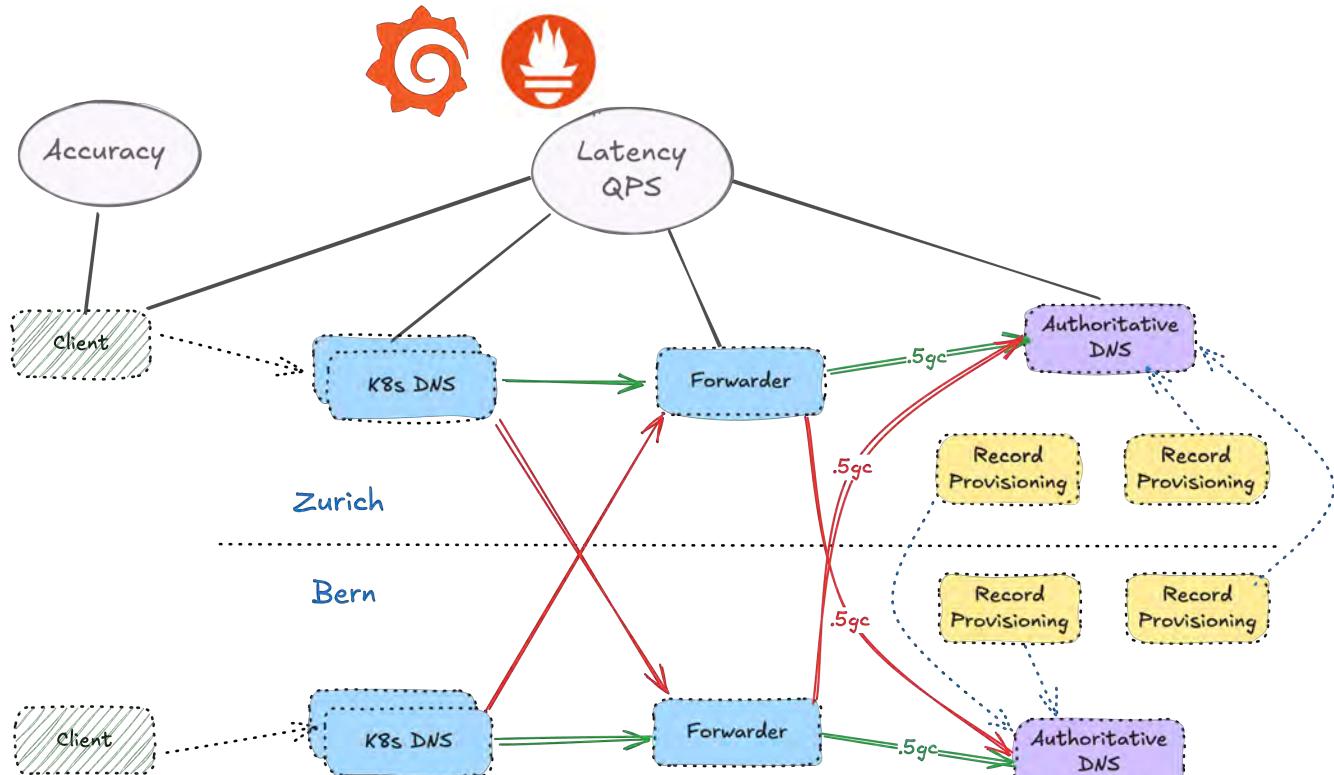
- Components on Resolution Path TC1, TC2, TC3
- Record Provisioning component TC4

DNS Traffic per Cluster

- 50 QPS
- Record Update per 10 seconds

Resilience Testing on Cloud Native DNS

How → DNS Setup -- Testing Strategy -- Metrics



Metrics Collected:

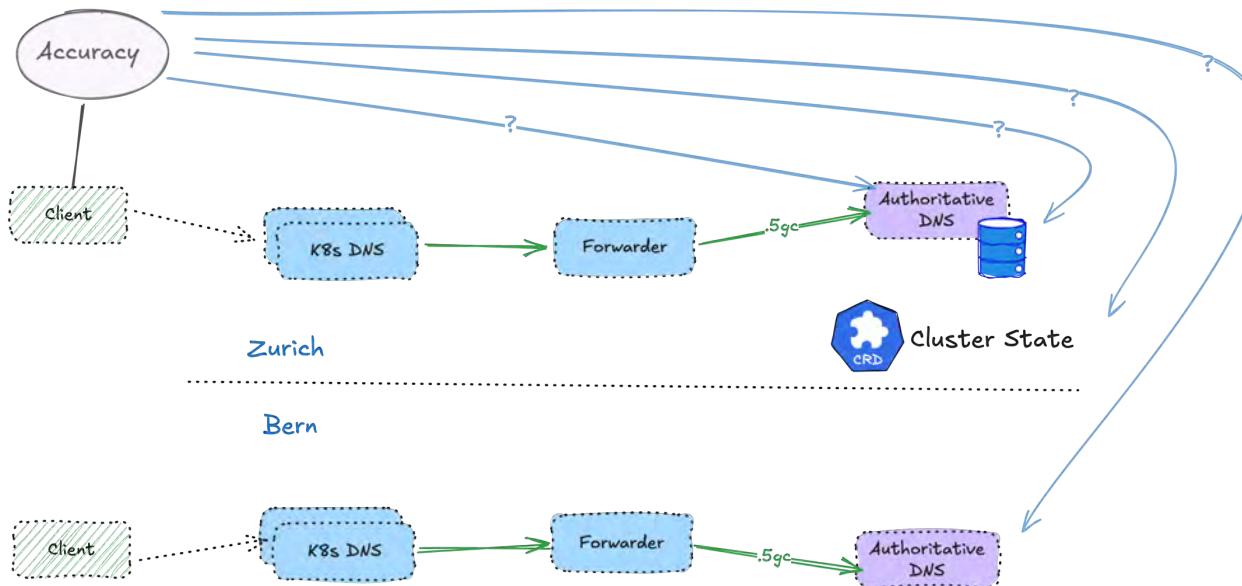
- Latency
- Queries Per Second
- Accuracy of Resolutions %

Observability Stack:

- Separate Kind Cluster
- Prometheus
- Grafana

Resilience Testing on Cloud Native DNS

How → Assessing Accuracy



What is the source of truth?

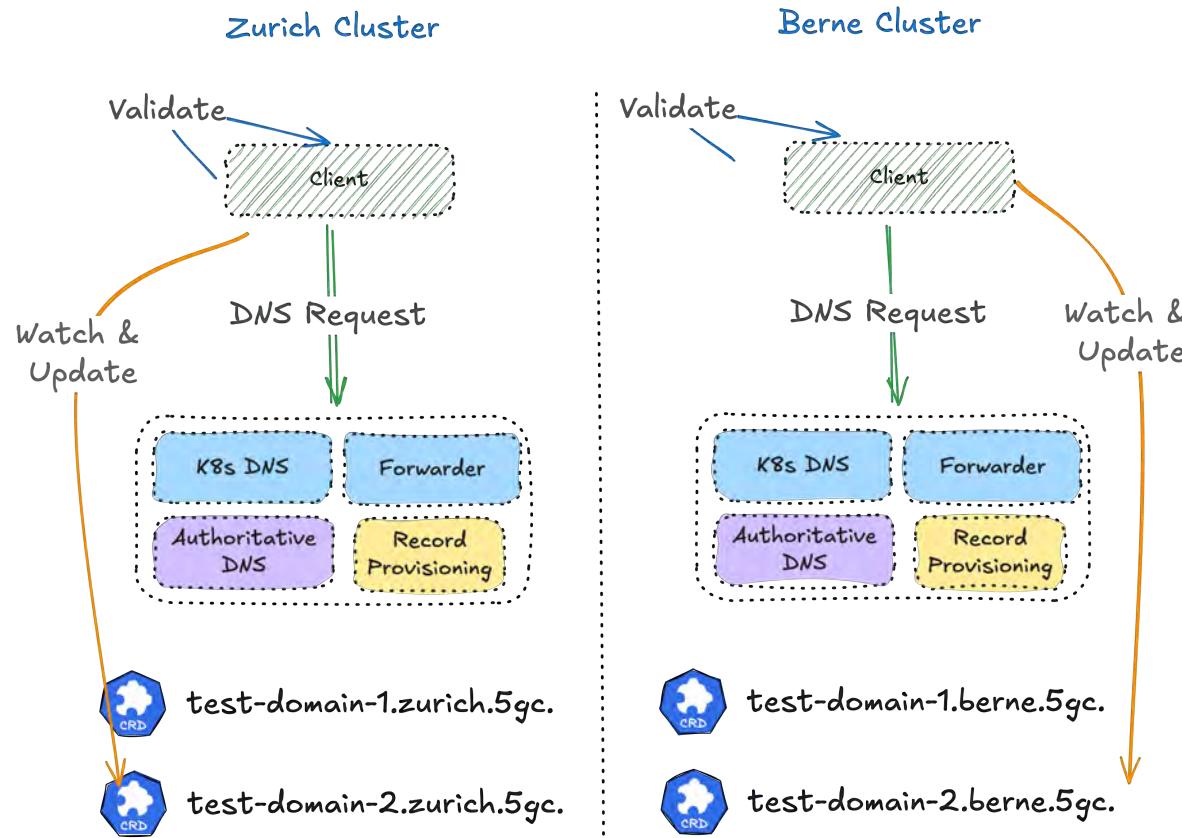
- Authoritative Server?
- Authoritative Server Database?
- Quorum of Auth Servers?
- Cluster K8s API State?

Challenge of temporal coupling

⚠ More checks, more uncertainty.

Resilience Testing on Cloud Native DNS

How → DNS Client -- Traffic Patterns



Client becomes the source of truth.

- Maintains cluster state internally
- Issues updates
- Compares resolution against internal state
- Minimised temporal coupling

Traffic Patterns:

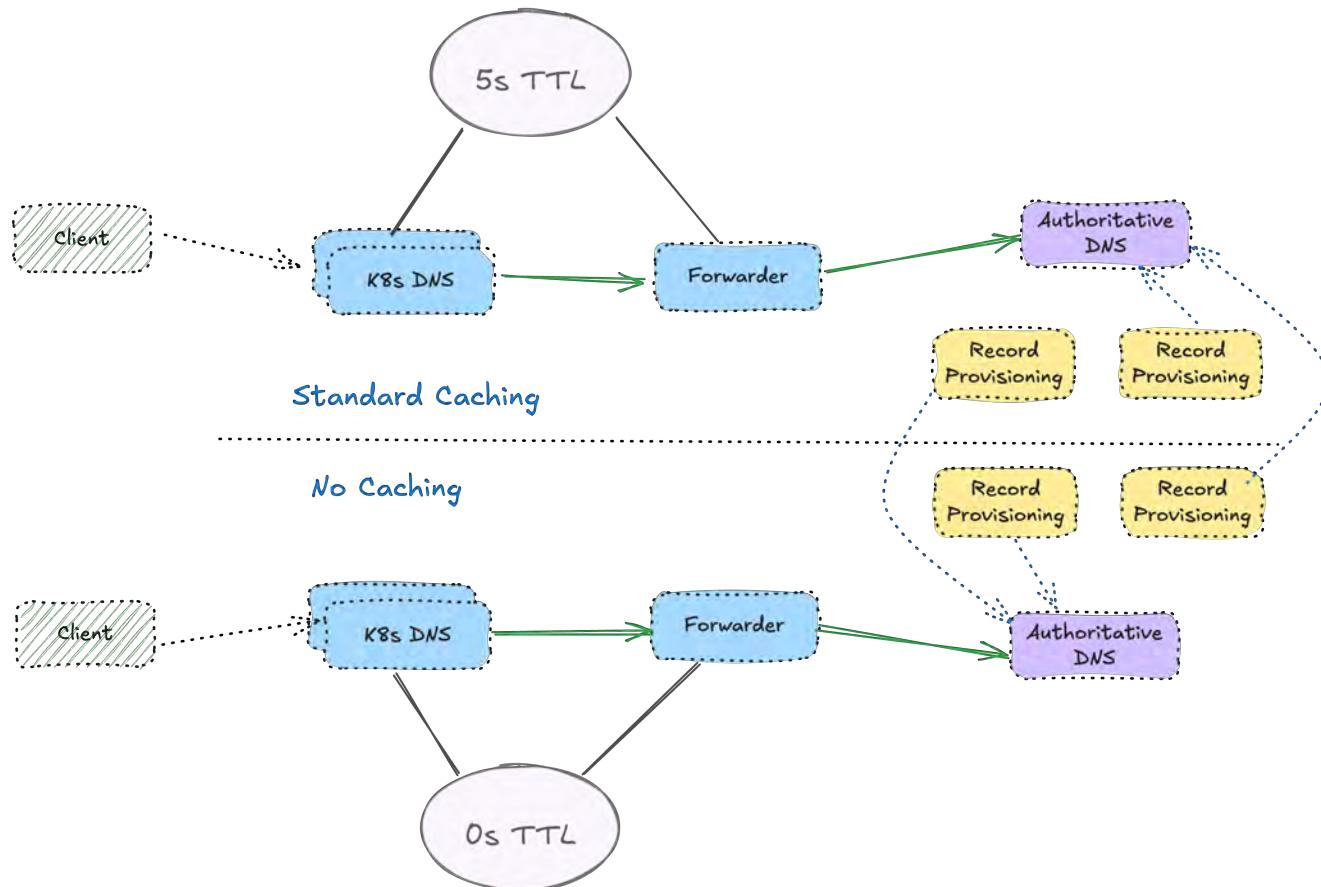
- Disjoined domain sets in each cluster
- Client only validates records in its cluster

Limitation over cross cluster validation.

Implementation based on [Hostlookuper](#)

Resilience Testing on Cloud Native DNS

How → Caching Configurations -- Steady State

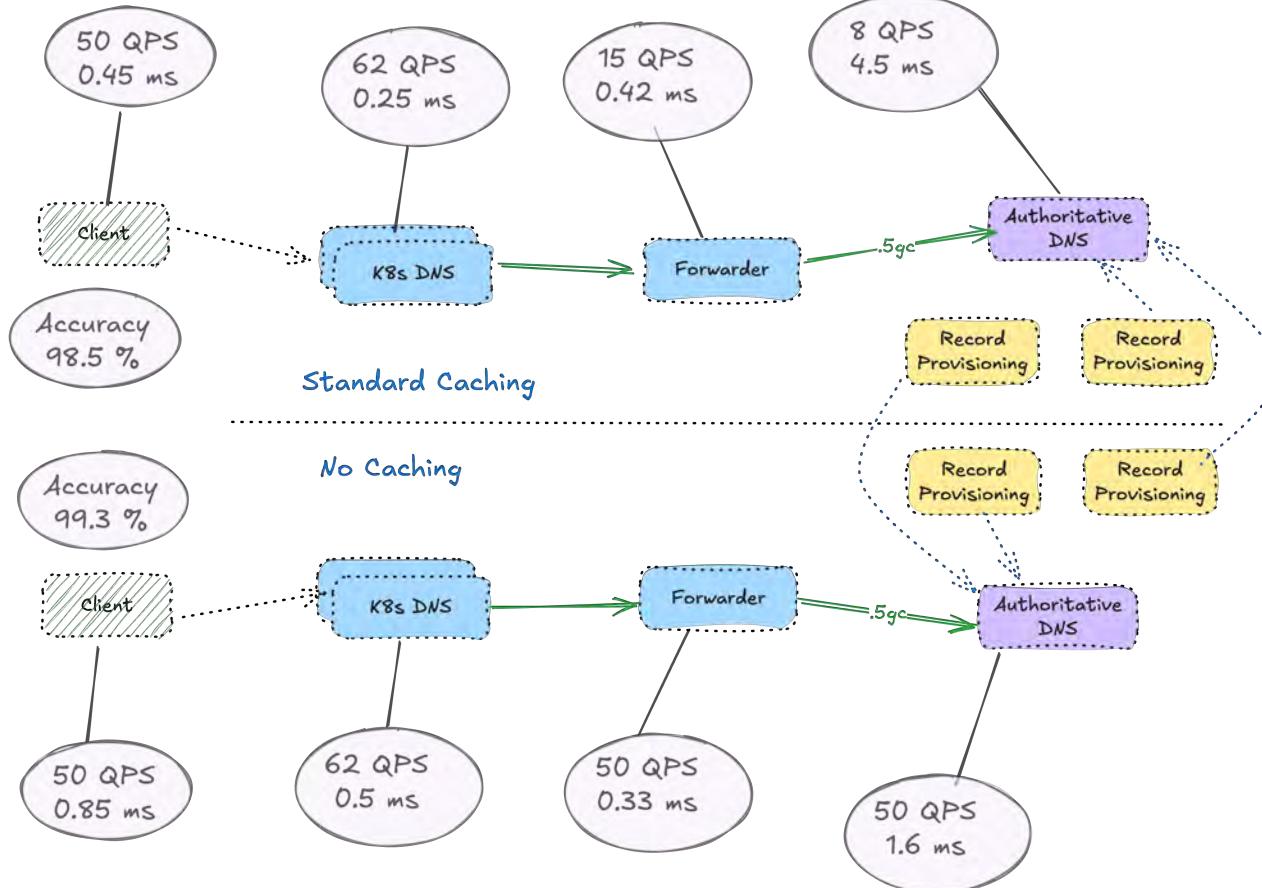


2 Caching Configurations:

- 5s TTL – Standard Caching
- No Caching

Resilience Testing on Cloud Native DNS

How → Caching Configurations -- Steady State



Steady State Measurements:

- One Client per Cluster
- 50 Unique Domains per Second



How → Chainsaw -- Datadog Chaos Injector

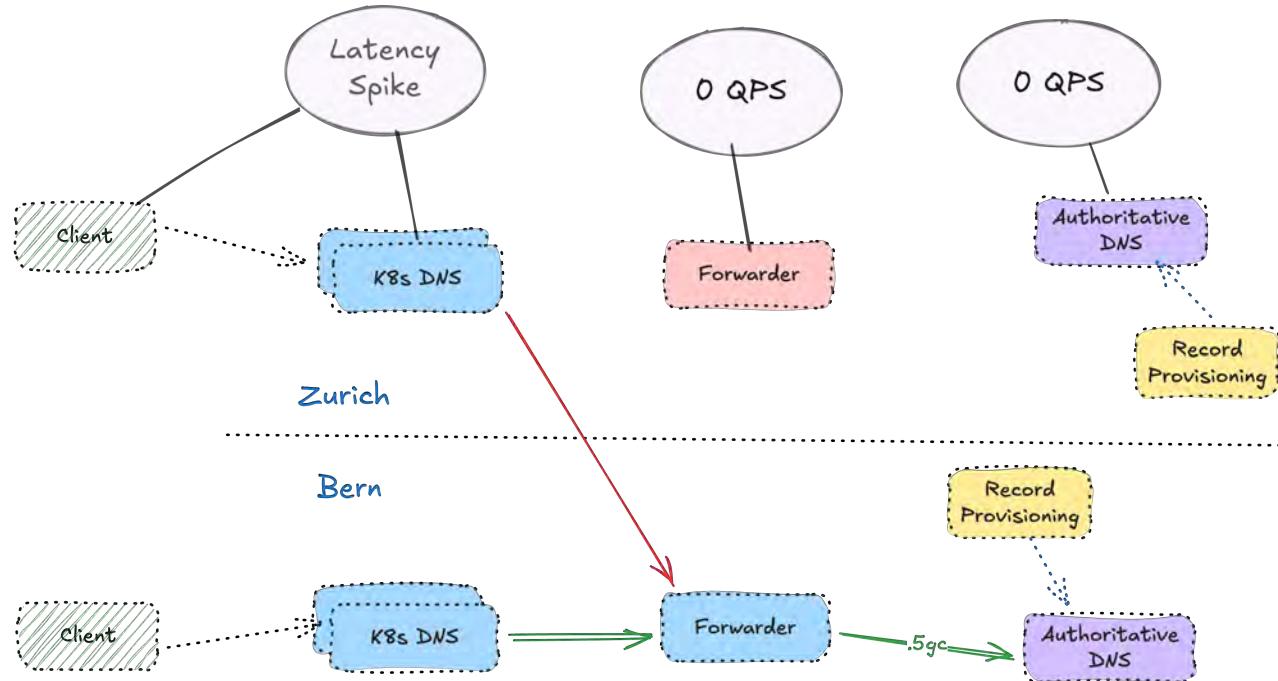
```
== CONN Chainsaw
== RUN chainsaw/k8s-coredns-chaos-test
| 14:46:03 | k8s-coredns-chaos-test | @chainsaw | CREATE | OK
| 14:46:03 | k8s-coredns-chaos-test | Post Test Start | TRY | BEGIN
| 14:46:03 | k8s-coredns-chaos-test | Post Test Start | CMD | RUN
    == COMMAND
    /usr/bin/sh -c ../utils/annotate-grafana.sh test-k8s-dns-start
| 14:46:03 | k8s-coredns-chaos-test | Post Test Start | SCRIPT | LOG
    == STDOUT
    {"id":124,"message":"Annotation added"}
| 14:46:03 | k8s-coredns-chaos-test | Post Test Start | SCRIPT | DONE
| 14:46:03 | k8s-coredns-chaos-test | Post Test Start | TRY | END
| 14:46:03 | k8s-coredns-chaos-test | Inject Chaos | TRY | BEGIN
| 14:46:03 | k8s-coredns-chaos-test | Inject Chaos | APPLY | RUN
| 14:46:03 | k8s-coredns-chaos-test | Inject Chaos | CREATE | OK
| 14:46:03 | k8s-coredns-chaos-test | Inject Chaos | APPLY | DONE
| 14:46:03 | k8s-coredns-chaos-test | Inject Chaos | ASSERT | RUN
| 14:46:03 | k8s-coredns-chaos-test | Inject Chaos | ASSERT | DONE
| 14:48:04 | k8s-coredns-chaos-test | Inject Chaos | TRY | END
| 14:48:04 | k8s-coredns-chaos-test | Inject Chaos | TRY | BEGIN
| 14:48:04 | k8s-coredns-chaos-test | Post Test Stop | TRY | RUN
| 14:48:04 | k8s-coredns-chaos-test | Post Test Stop | CMD | RUN
    == COMMAND
    /usr/bin/sh -c ../utils/annotate-grafana.sh test-k8s-dns-end
| 14:48:04 | k8s-coredns-chaos-test | Post Test Stop | SCRIPT | LOG
    == STDOUT
    {"id":125,"message":"Annotation added"}
| 14:48:04 | k8s-coredns-chaos-test | Post Test Stop | SCRIPT | DONE
| 14:48:04 | k8s-coredns-chaos-test | Post Test Stop | TRY | END
| 14:48:04 | k8s-coredns-chaos-test | Validate Revival | TRY | BEGIN
| 14:48:04 | k8s-coredns-chaos-test | Validate Revival | ASSERT | RUN
| 14:49:13 | k8s-coredns-chaos-test | Validate Revival | ASSERT | DONE
| 14:49:13 | k8s-coredns-chaos-test | Validate Revival | TRY | END
| 14:49:13 | k8s-coredns-chaos-test | Annotate System Restored | TRY | BEGIN
| 14:49:13 | k8s-coredns-chaos-test | Annotate System Restored | CMD | RUN
    == COMMAND
    /usr/bin/sh -c ../utils/annotate-grafana.sh test-k8s-dns-system-restored
| 14:49:13 | k8s-coredns-chaos-test | Annotate System Restored | SCRIPT | LOG
    == STDOUT
    {"id":126,"message":"Annotation added"}
| 14:49:13 | k8s-coredns-chaos-test | Annotate System Restored | SCRIPT | DONE
| 14:49:13 | k8s-coredns-chaos-test | Annotate System Restored | TRY | END
| 14:49:13 | k8s-coredns-chaos-test | Inject Chaos | CLEANUP | BEGIN
| 14:49:13 | k8s-coredns-chaos-test | Inject Chaos | DELETE | OK
| 14:49:44 | k8s-coredns-chaos-test | Inject Chaos | CLEANUP | END
| 14:49:44 | k8s-coredns-chaos-test | @chainsaw | CLEANUP | BEGIN
| 14:49:44 | k8s-coredns-chaos-test | @chainsaw | DELETE | OK
| 14:49:49 | k8s-coredns-chaos-test | @chainsaw | CLEANUP | END
—— PASS: chainsaw (226.74s)
—— PASS: chainsaw/k8s-coredns-chaos-test (226.74s)
PASS
```

Test Steps via Chainsaw:

1. Annotate chaos injection start
2. Inject Chaos for 2 minutes
3. Annotate chaos injection finish
4. Wait for the component to recover
5. Annotate recovery

Resilience Testing on Cloud Native DNS

Evaluation → TC2 -- Forwarder Disruption

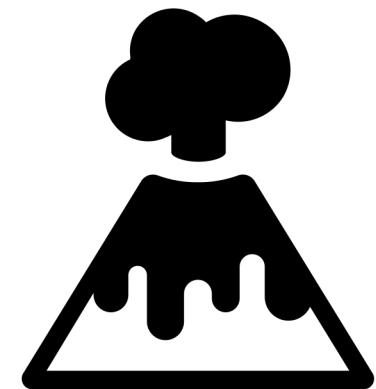
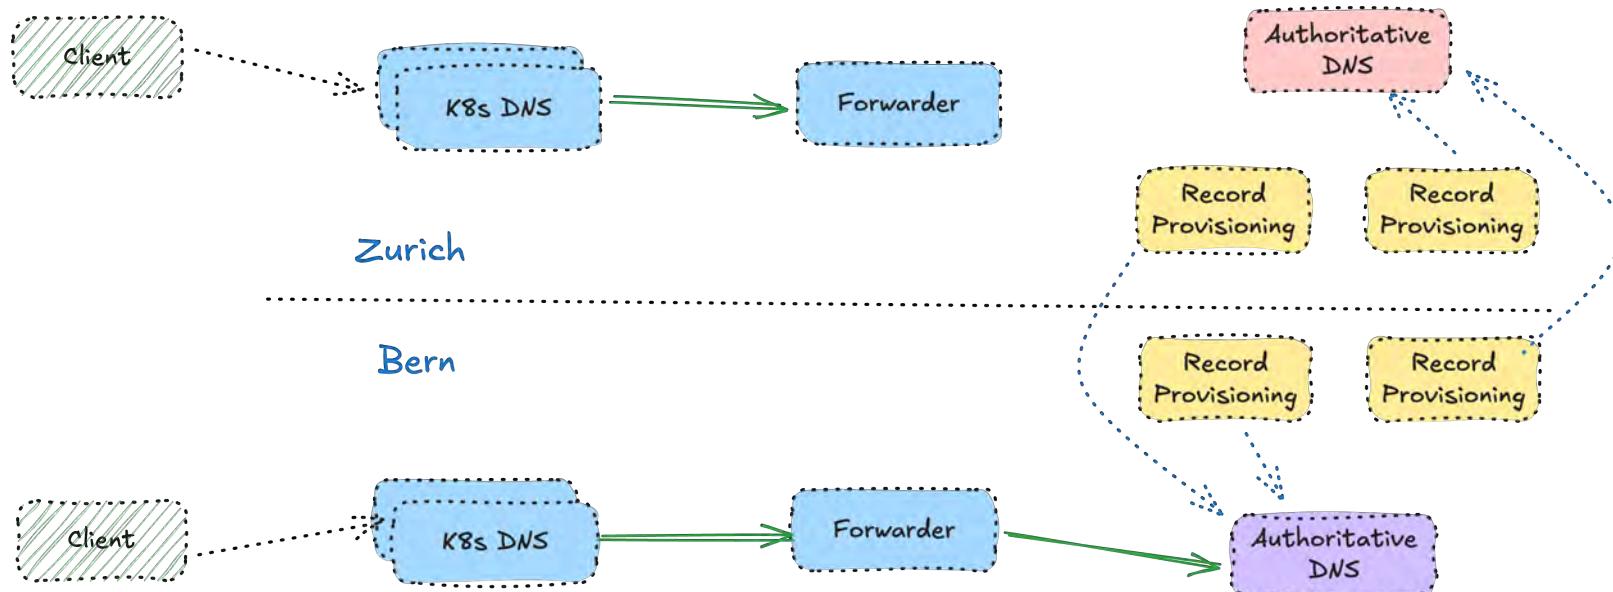


Fallback route is used

Latency spike in the beginning

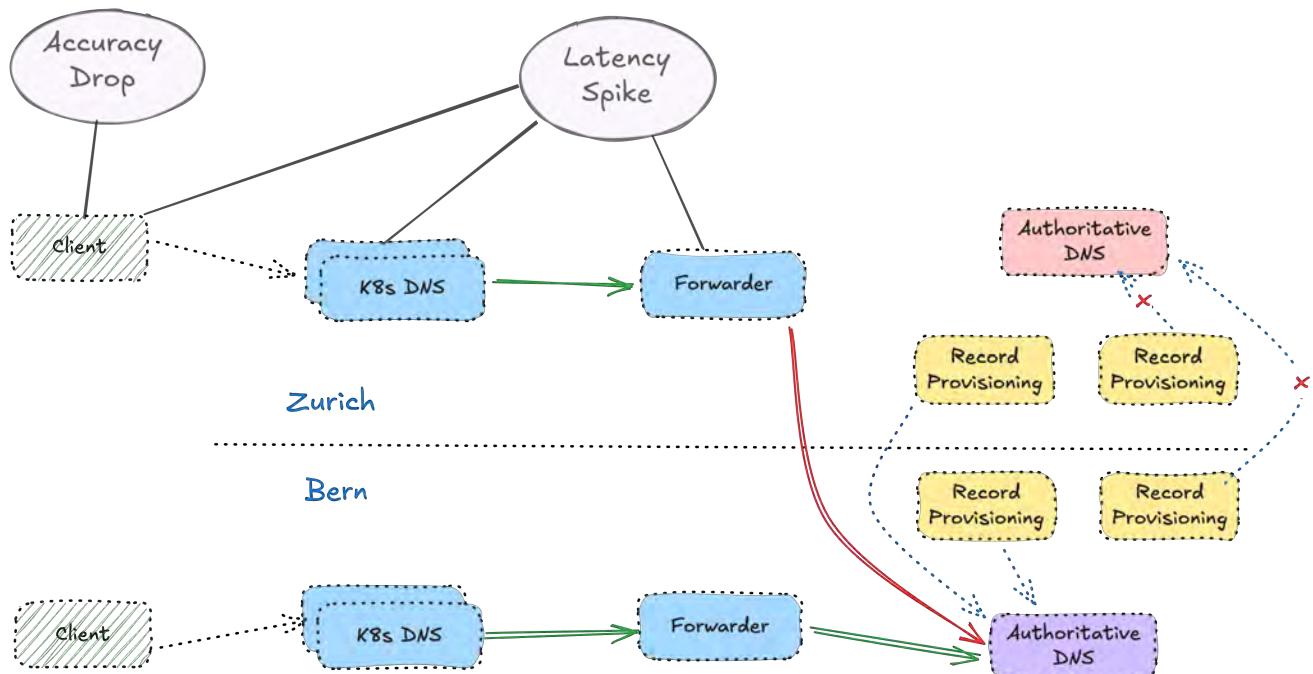
Higher latency spike for no-caching

Demo → TC3, Authoritative DNS Disruption





Evaluation → TC3 -- Authoritative DNS Disruption



Fallback route is used

Latency spike in the beginning

Higher latency spike for no-caching

Accuracy dips during restoration



Evaluation → Automation & Resilience



Fully Automated - GitOps driven
DNS record provisioning

✓ *Updates Propagate without
manual intervention*



Geo-Redundant & HA
System remains operational under
all tested disruptions

✓ *Operational under disruption*

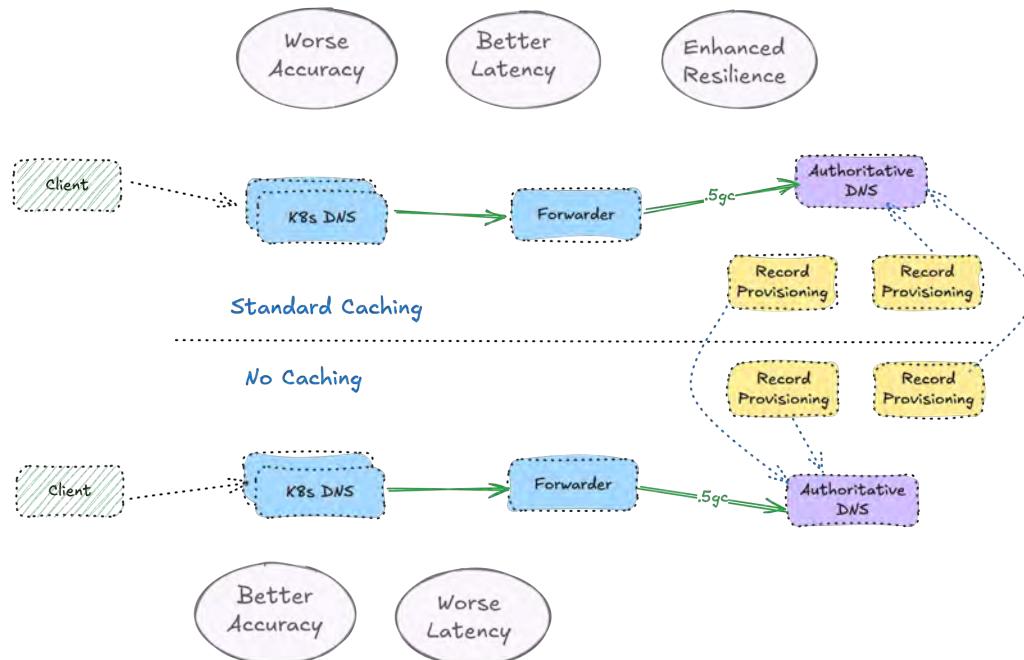


Minimal Amount of SPOFs
Remove single points of failure from
the System, Recover from Failure

✓ *Cross-Cluster Forwarding*
✓ *Predictable Recovery*



Evaluation → Caching

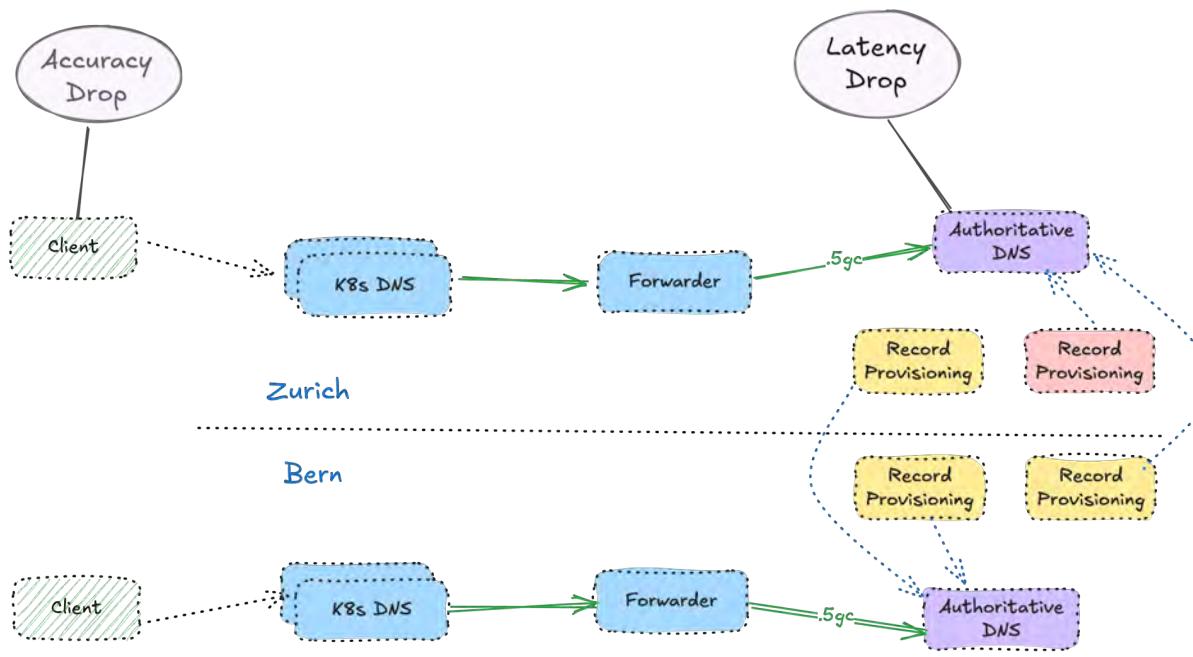


Caching improves resilience,
especially for failures occurred far from the client

No-caching, higher accuracy, worse latency
With TTL of 5s, Accuracy drop is estimated to ~0.5%

Resilience Testing on Cloud Native DNS

Evaluation → Testing Strategy



Testing strategy expose weaknesses

Client-side metrics → validate resilience

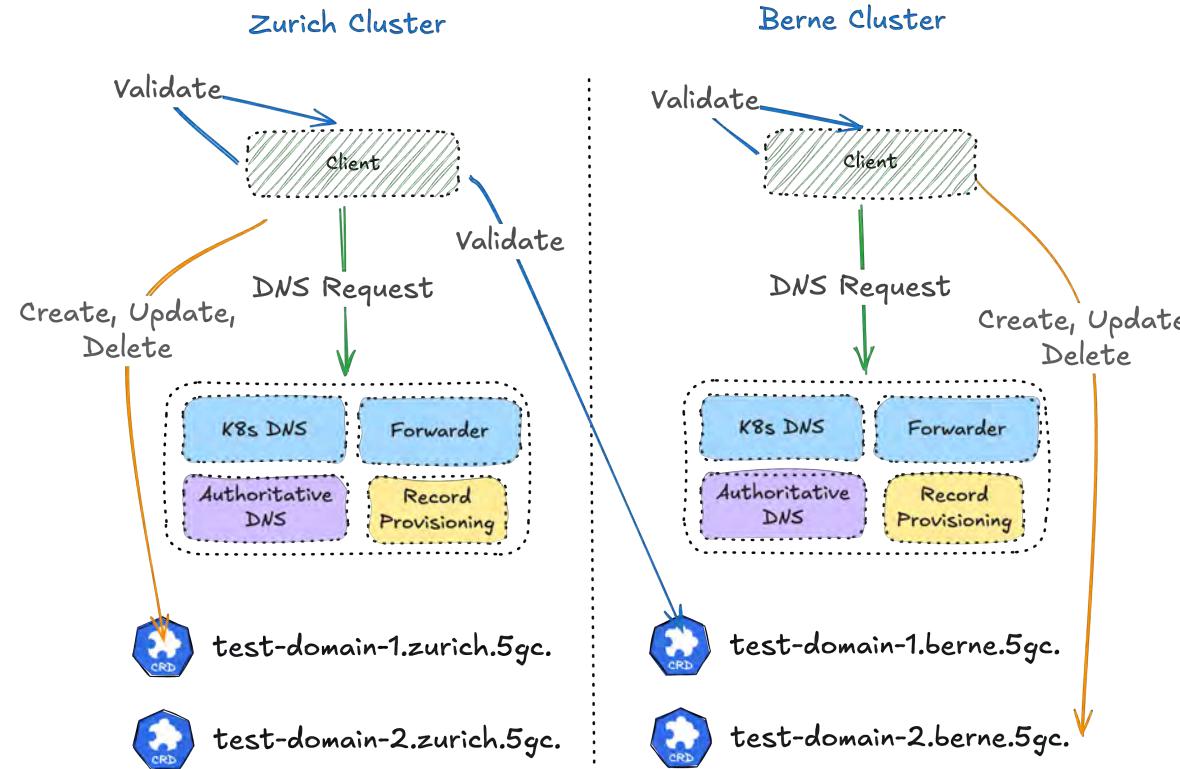
Internal system metrics → explain behaviour and root causes

Methodology easily extendable:

- Additional Clusters
- Alternative Caching Configurations

Resilience Testing on Cloud Native DNS

Future Work → Client Improvements



Limitations:

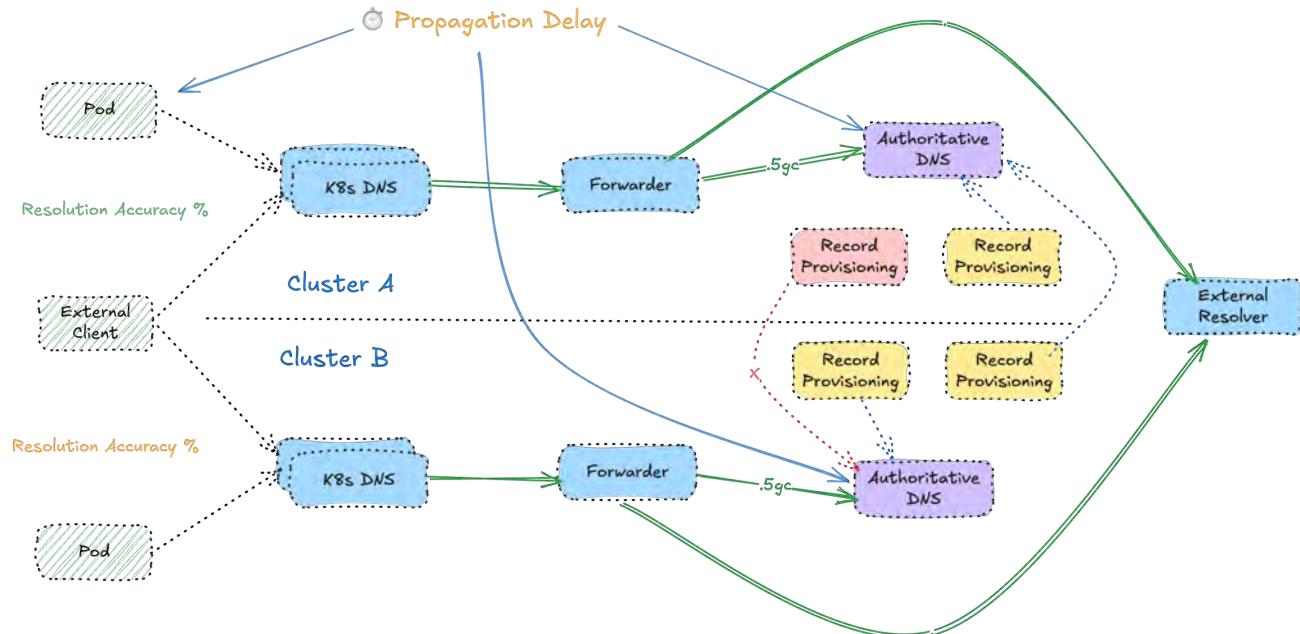
- Support only for updates
- Cross cluster validation

Improvements:

- Support Creation, Update & Deletion
- Cross Cluster lookups
- Fail fast delayed requests.



Future Work → Client Improvements



Estimate Propagation Delay

Measure time from record update to first correct client resolution

Error bounded by DNS query frequency



Thank you 



Q&A